

# Deep Learning for NLP

**Week 12: CNNs**

Sharid Loáiciga – July 7th, 2020

**Slides by Miikka Silfverberg & Hande Celikkanat**



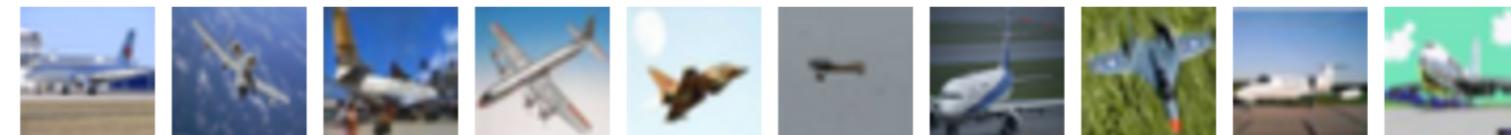
HELSINGIN YLIOPISTO  
HELSINGFORS UNIVERSITET  
UNIVERSITY OF HELSINKI

# Convolutional Neural Networks



# Image Classification

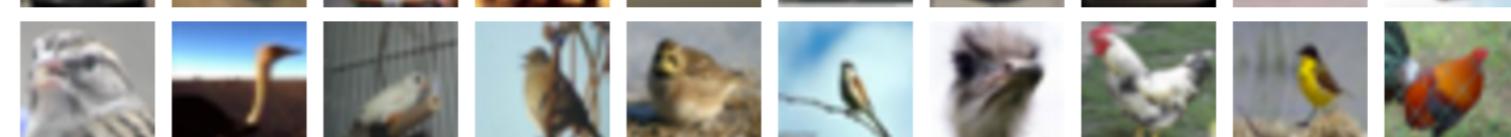
**airplane**



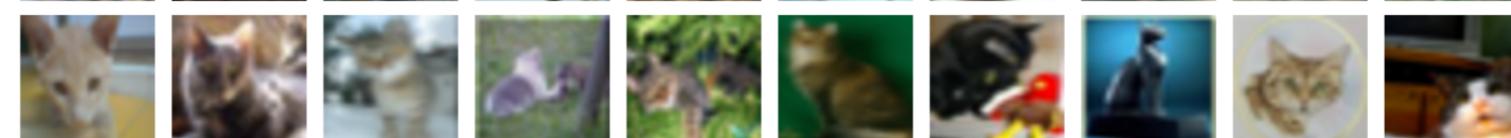
**automobile**



**bird**



**cat**



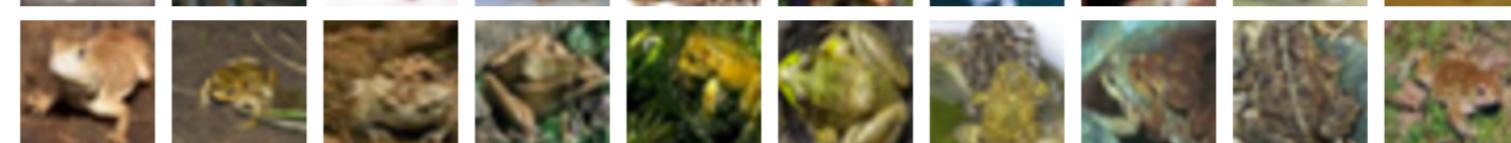
**deer**



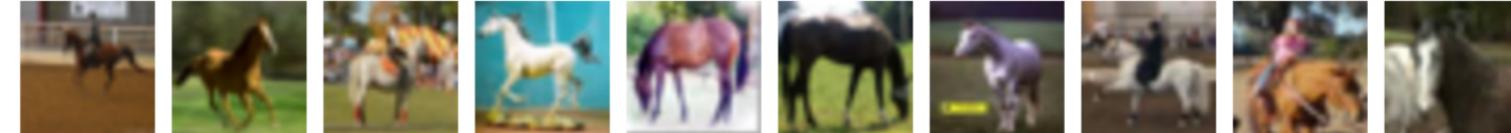
**dog**



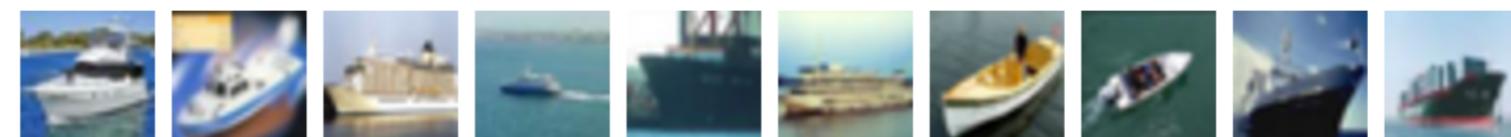
**frog**



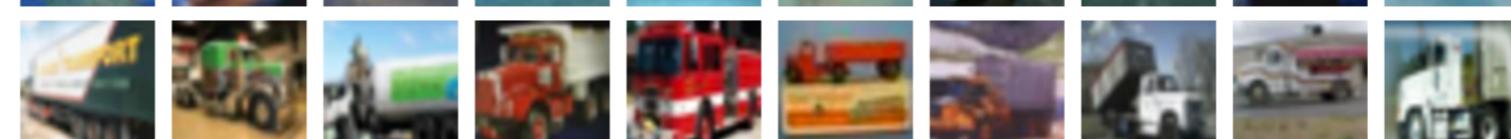
**horse**



**ship**



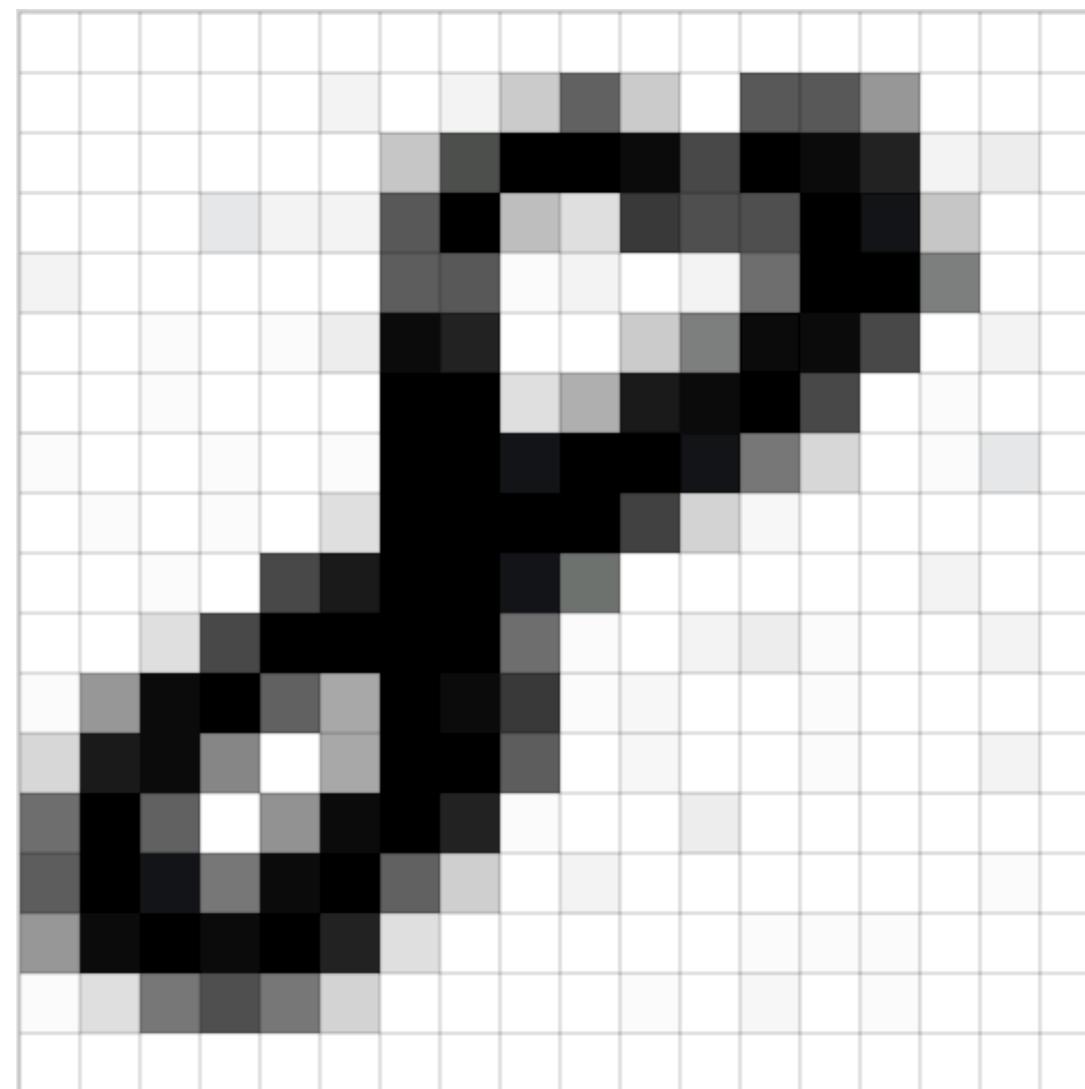
**truck**



CIFAR-10 dataset



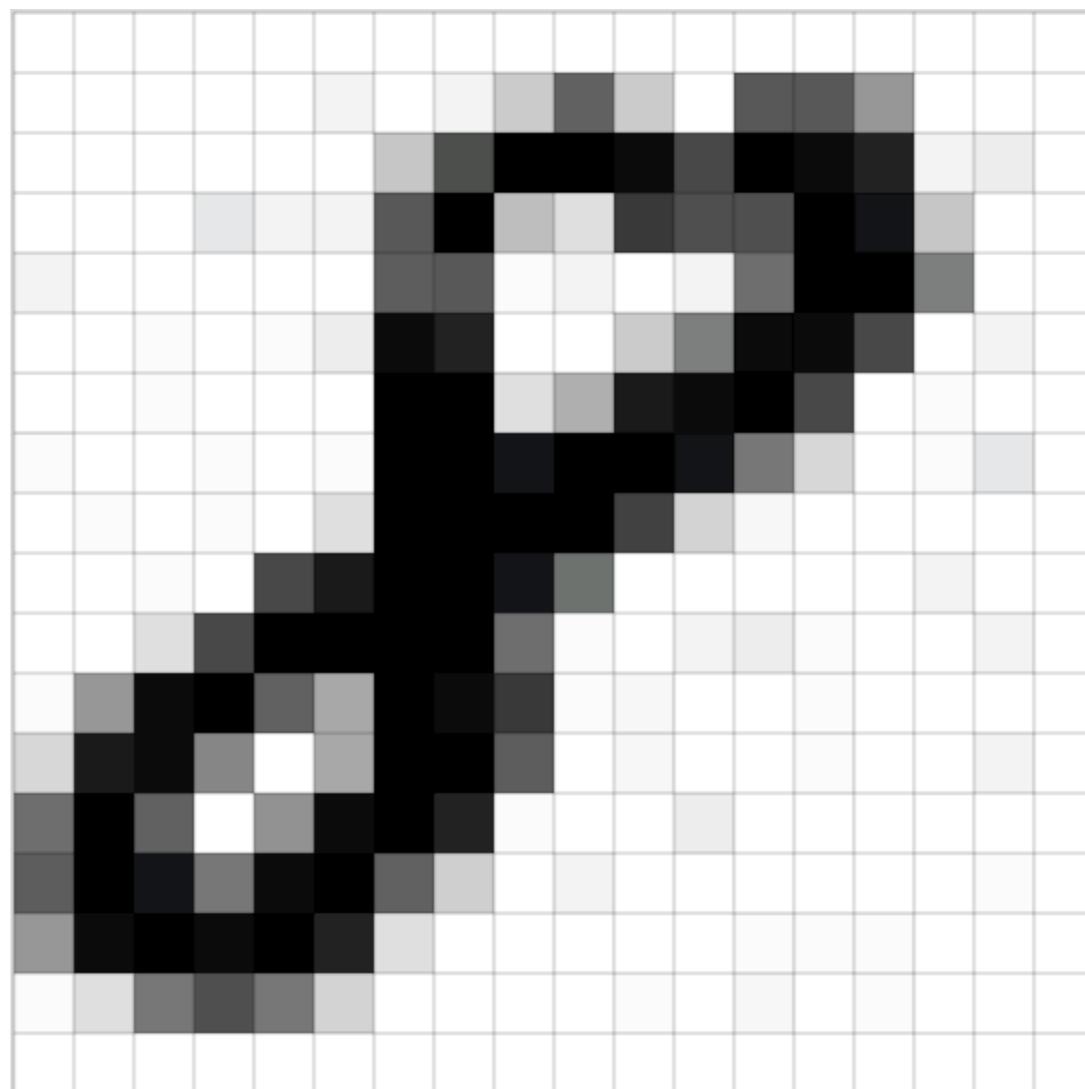
# Representing Images



The easiest way to represent images as matrices is to interpret each pixel as a real number value



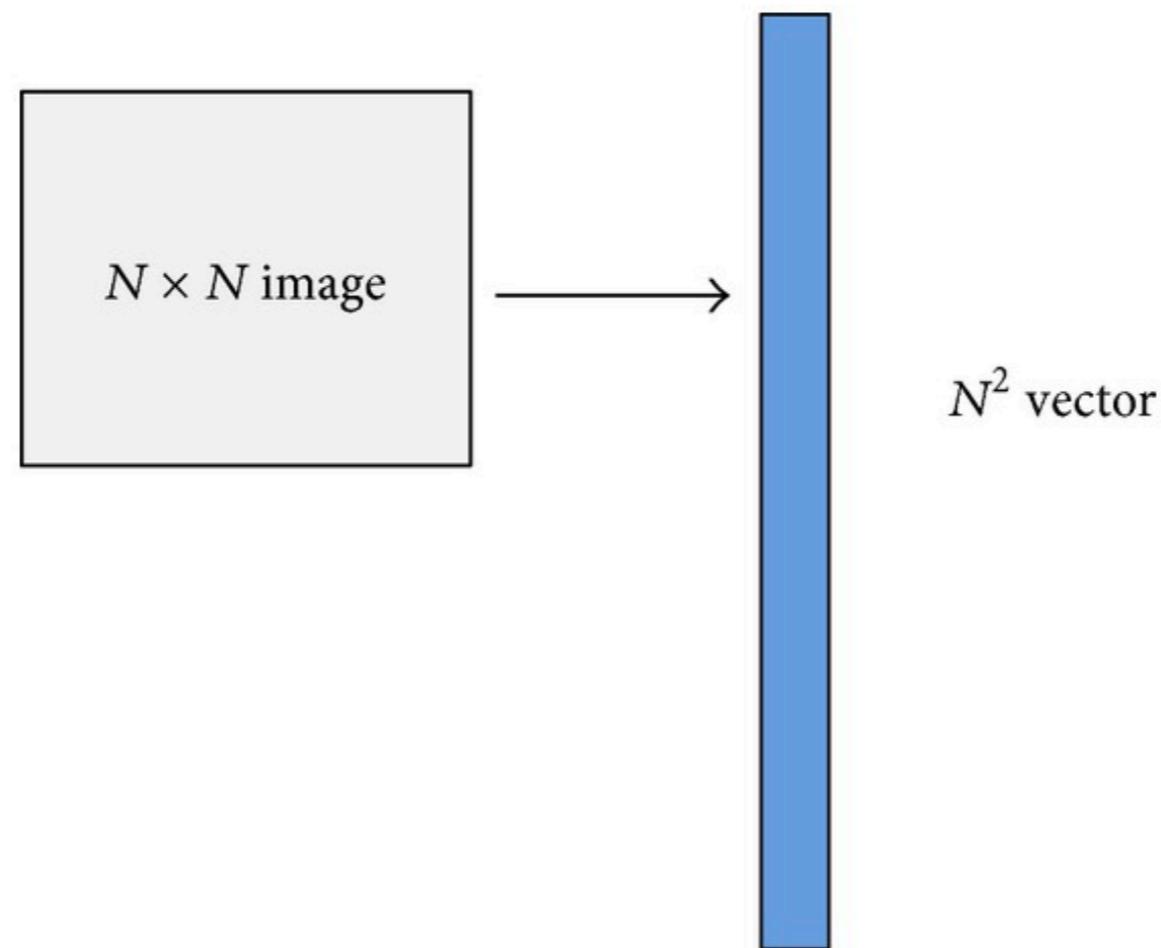
# Representing Images



The easiest way to represent images as matrices is to interpret each pixel as a real number value



# Representing Images



To get a vector, just concatenate rows or columns



# Representing Images



Red



Green



Blue

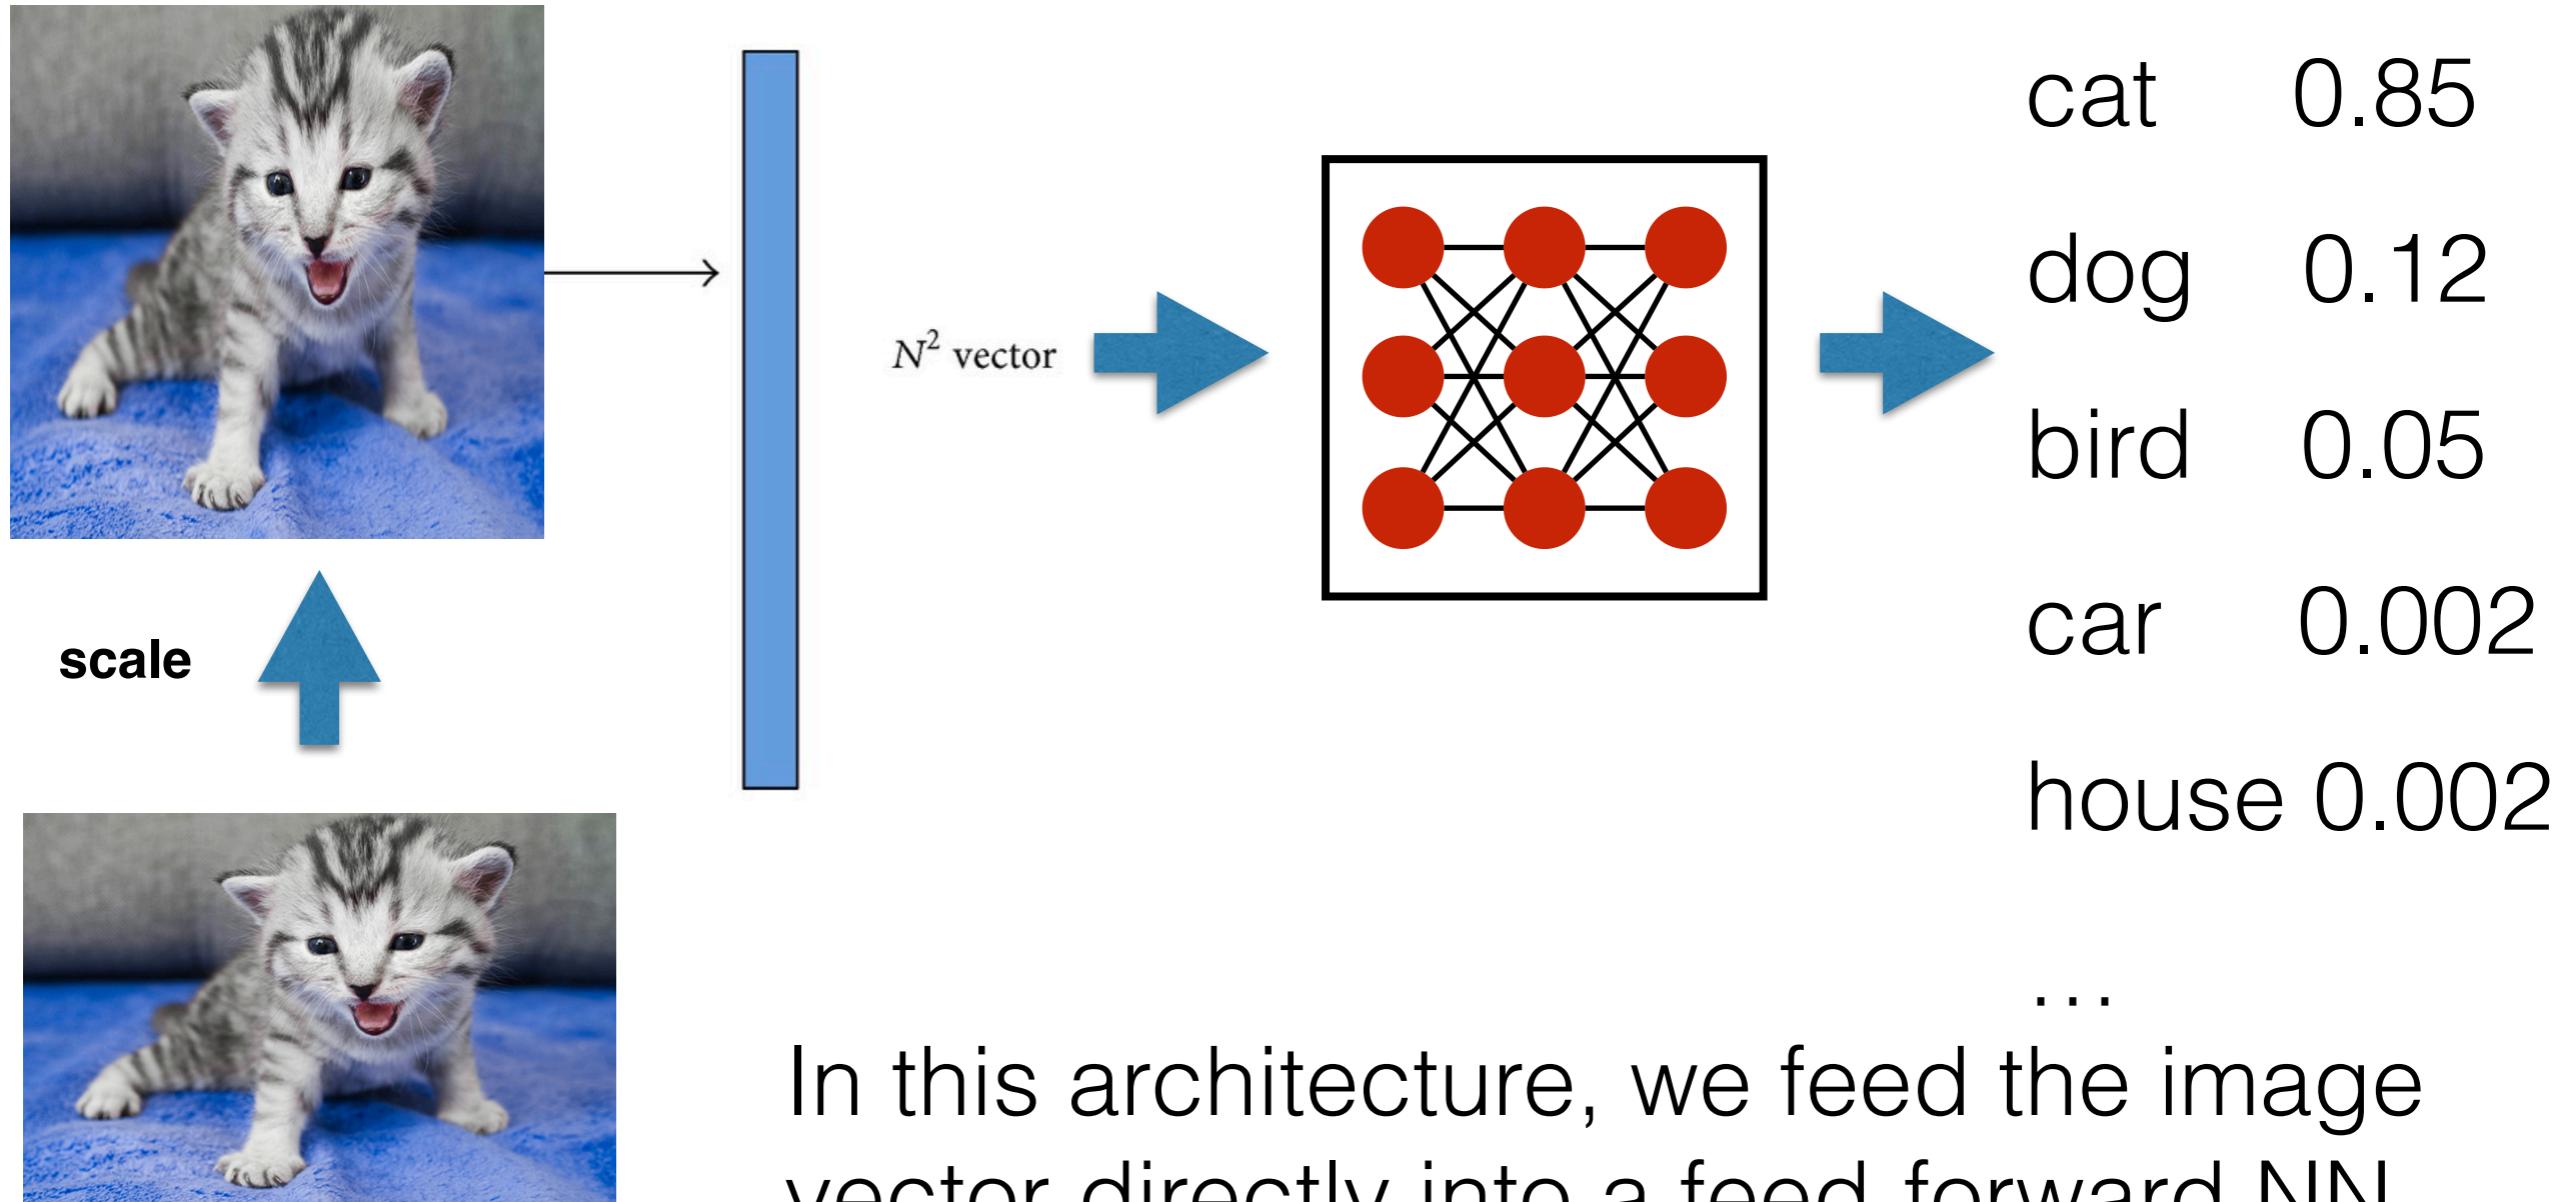


RGB color

We can split an image into a red, green and blue channel which can be transformed into vectors and concatenated.



# Image Classification using Feed-Forward Networks



In this architecture, we feed the image vector directly into a feed-forward NN trained for classification.



# Problems





# Problems

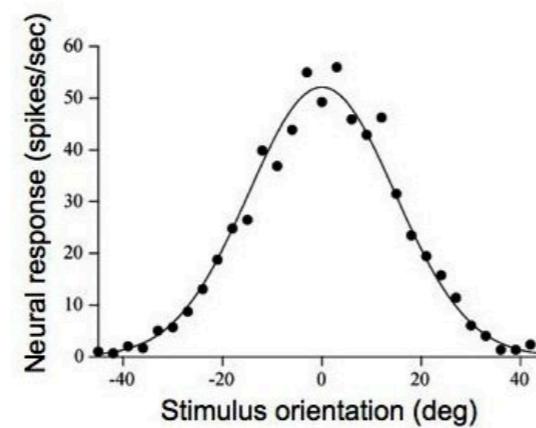
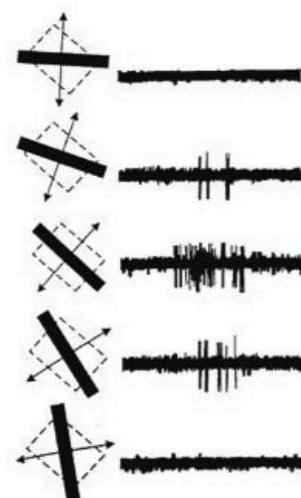
Learning to recognize a circle in the left upper corner doesn't (directly) help in recognizing a circle in the right lower corner.



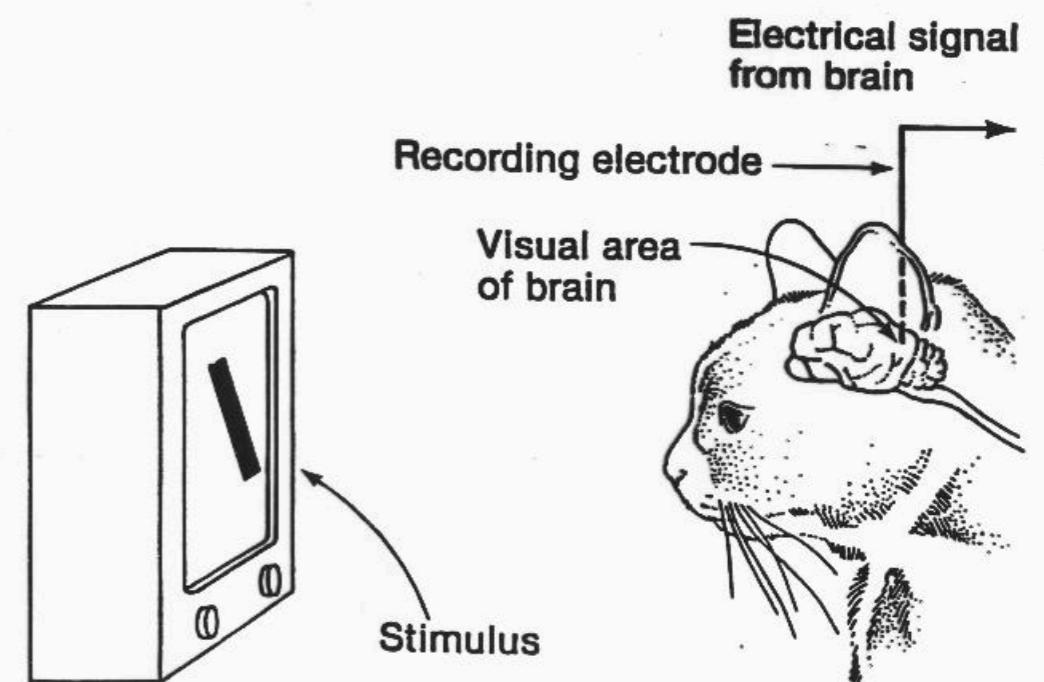
The main problem in the FFNN architecture is that there is no parameter sharing across the spatial dimension.



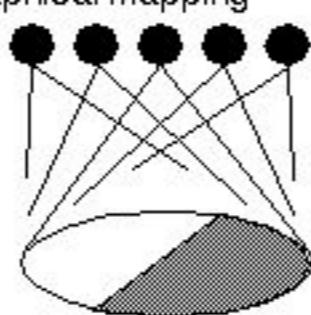
# Biological Inspiration



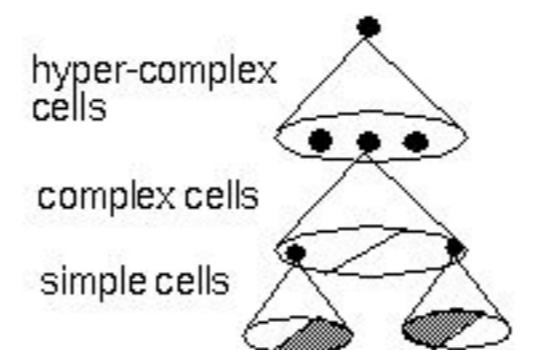
Hubel & Wiesel, 1968



Hubel & Weisel  
topographical mapping



featural hierarchy



- high level
- mid level
- ▨ low level

Hubel & Wiesel 1958

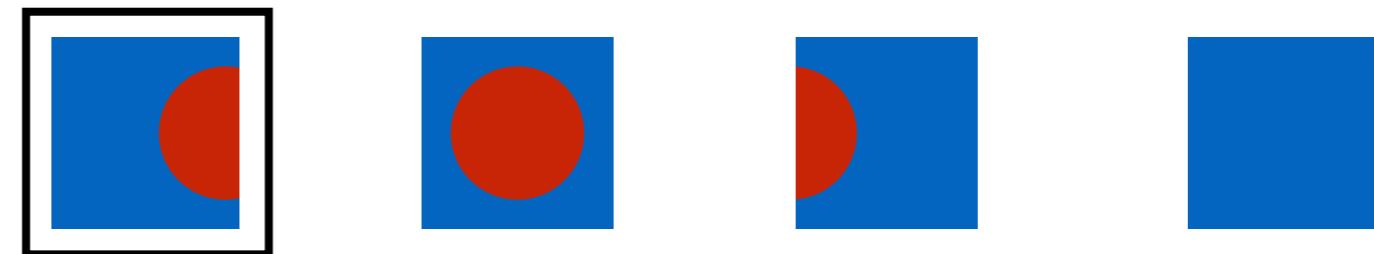


# Biological Inspiration

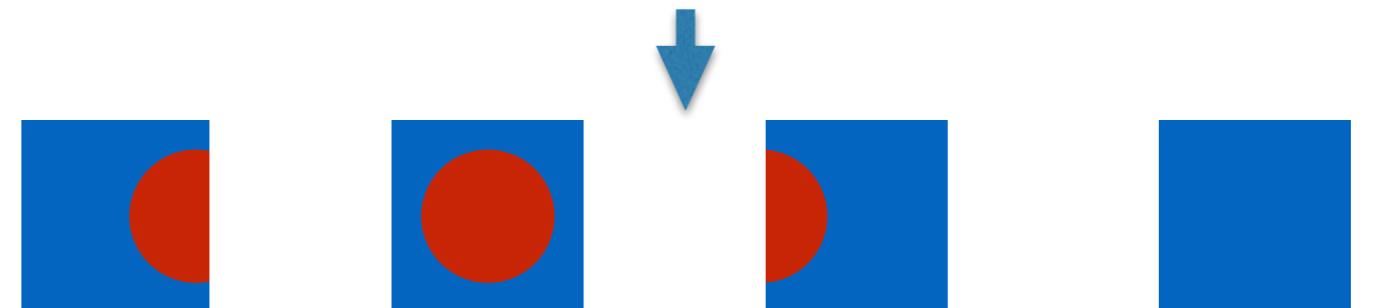




# Convolutional filters: Intuition



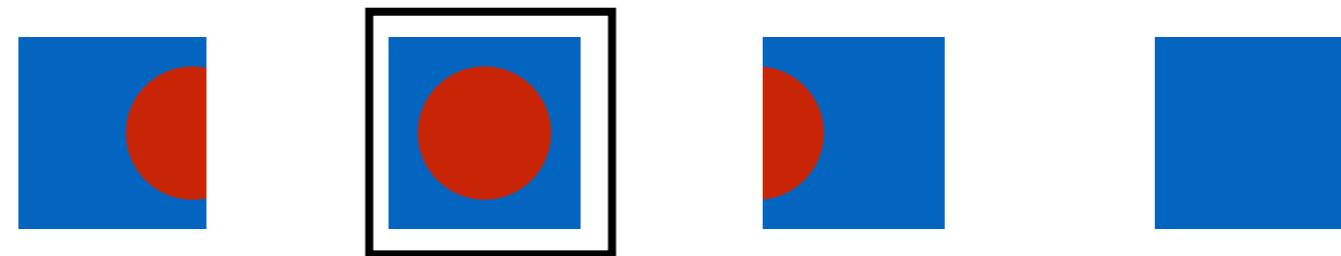
The images are composed of pieces like this



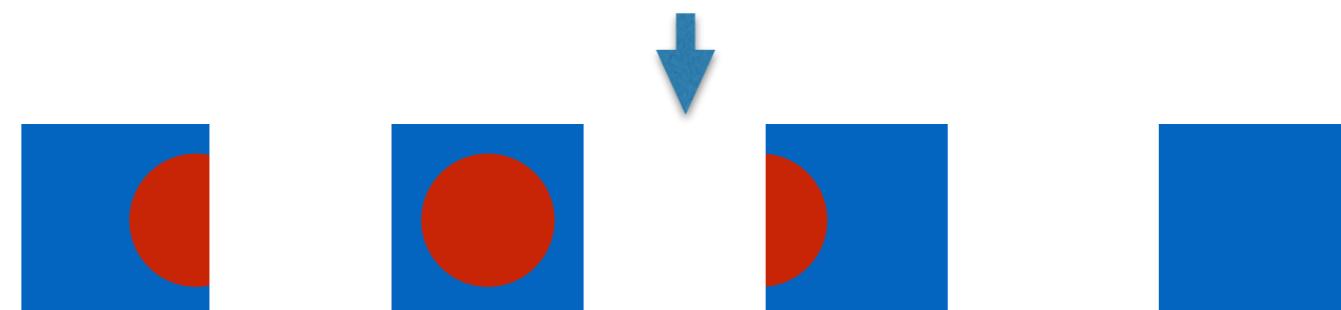
Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



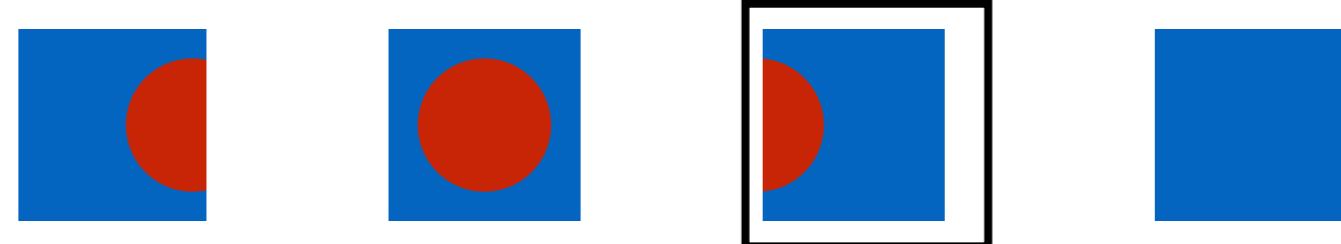
The images are composed of pieces like this



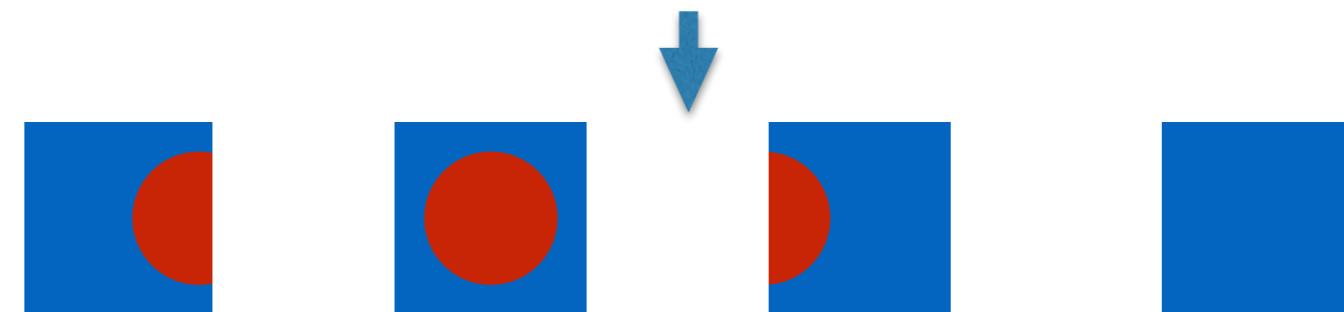
Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



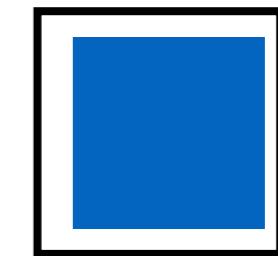
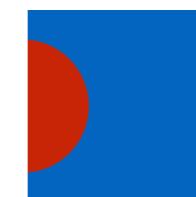
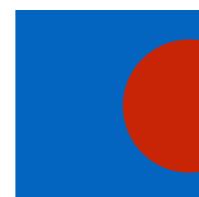
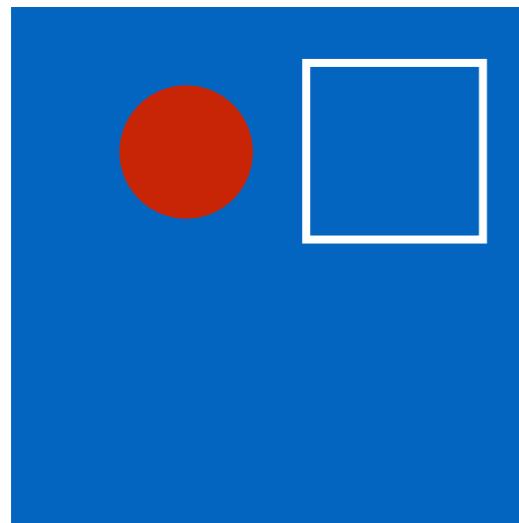
The images are composed of pieces like this



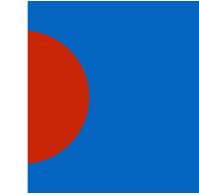
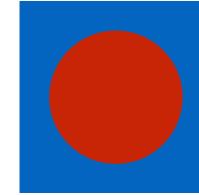
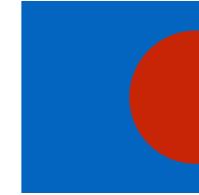
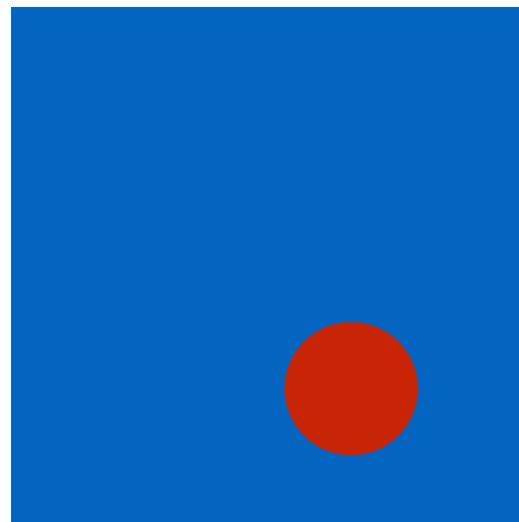
Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



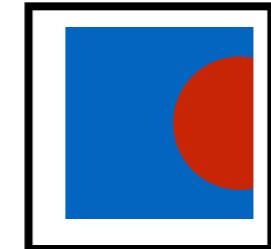
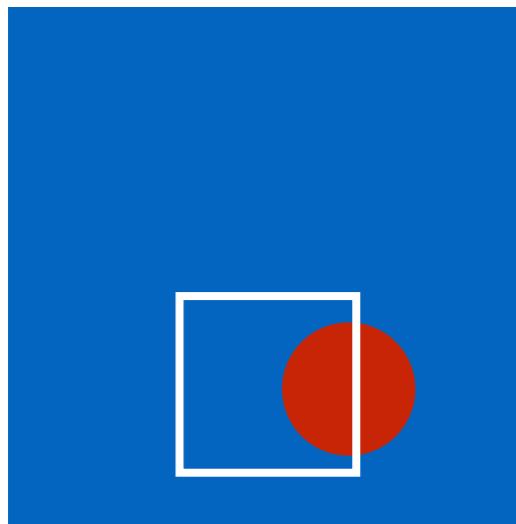
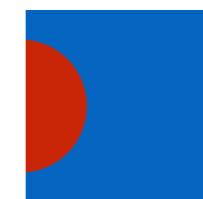
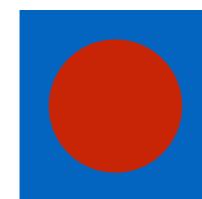
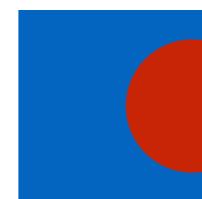
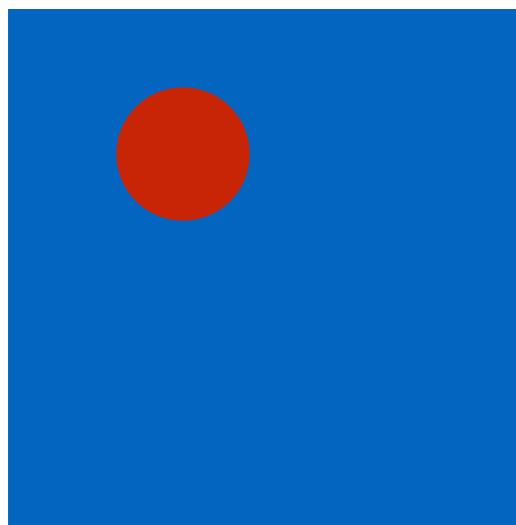
The images are composed of pieces like this



Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



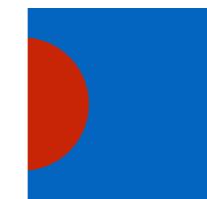
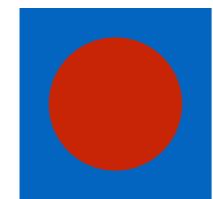
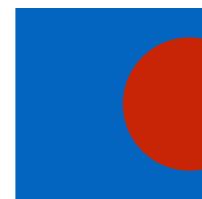
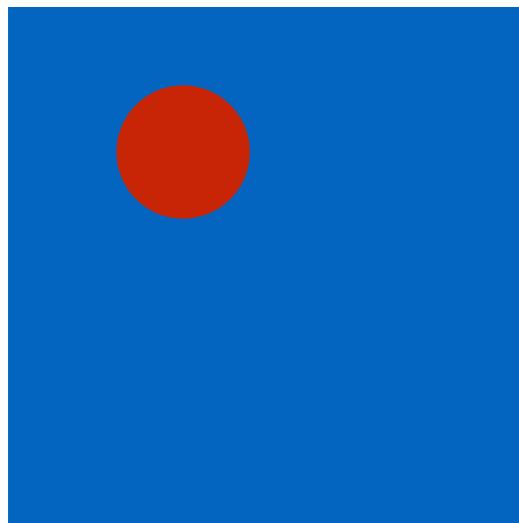
The images are composed of pieces like this



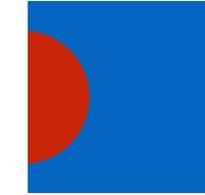
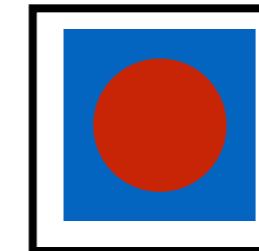
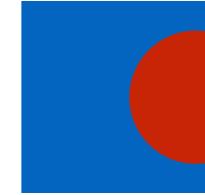
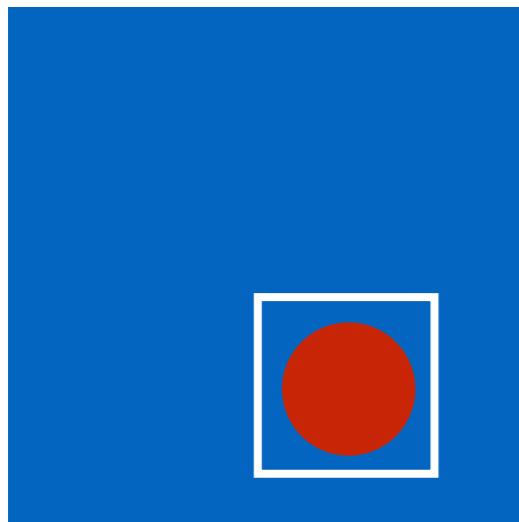
Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



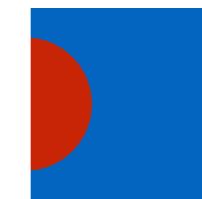
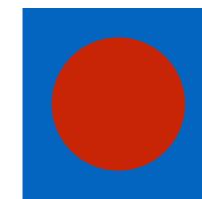
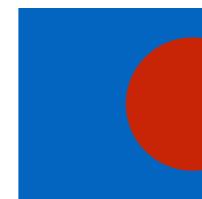
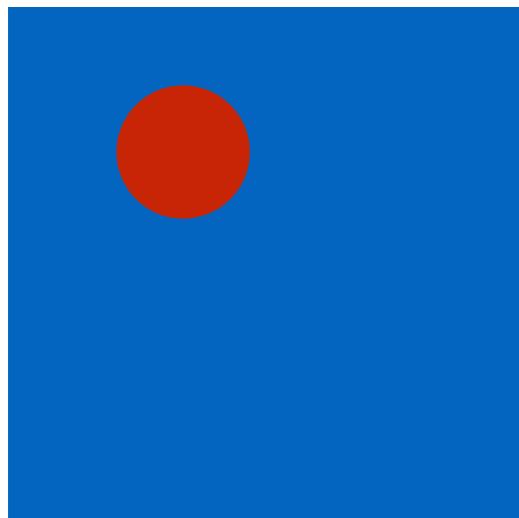
The images are composed of pieces like this



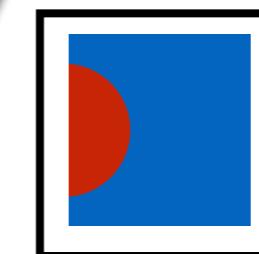
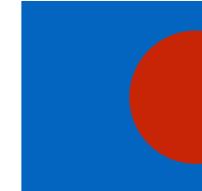
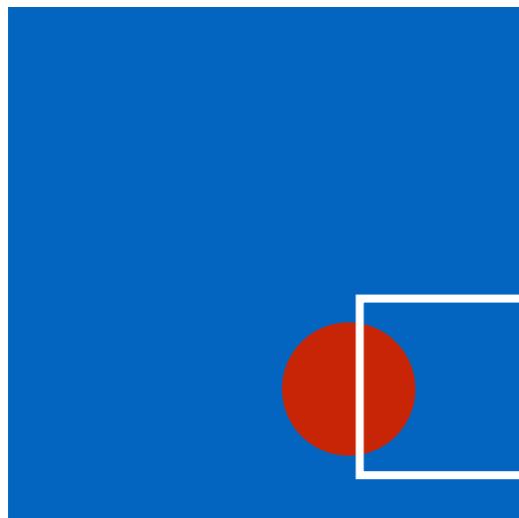
Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional filters: Intuition



The images are composed of pieces like this



Using convolutional filters, we can look at a small section of the image at a time.



# Convolutional Filter

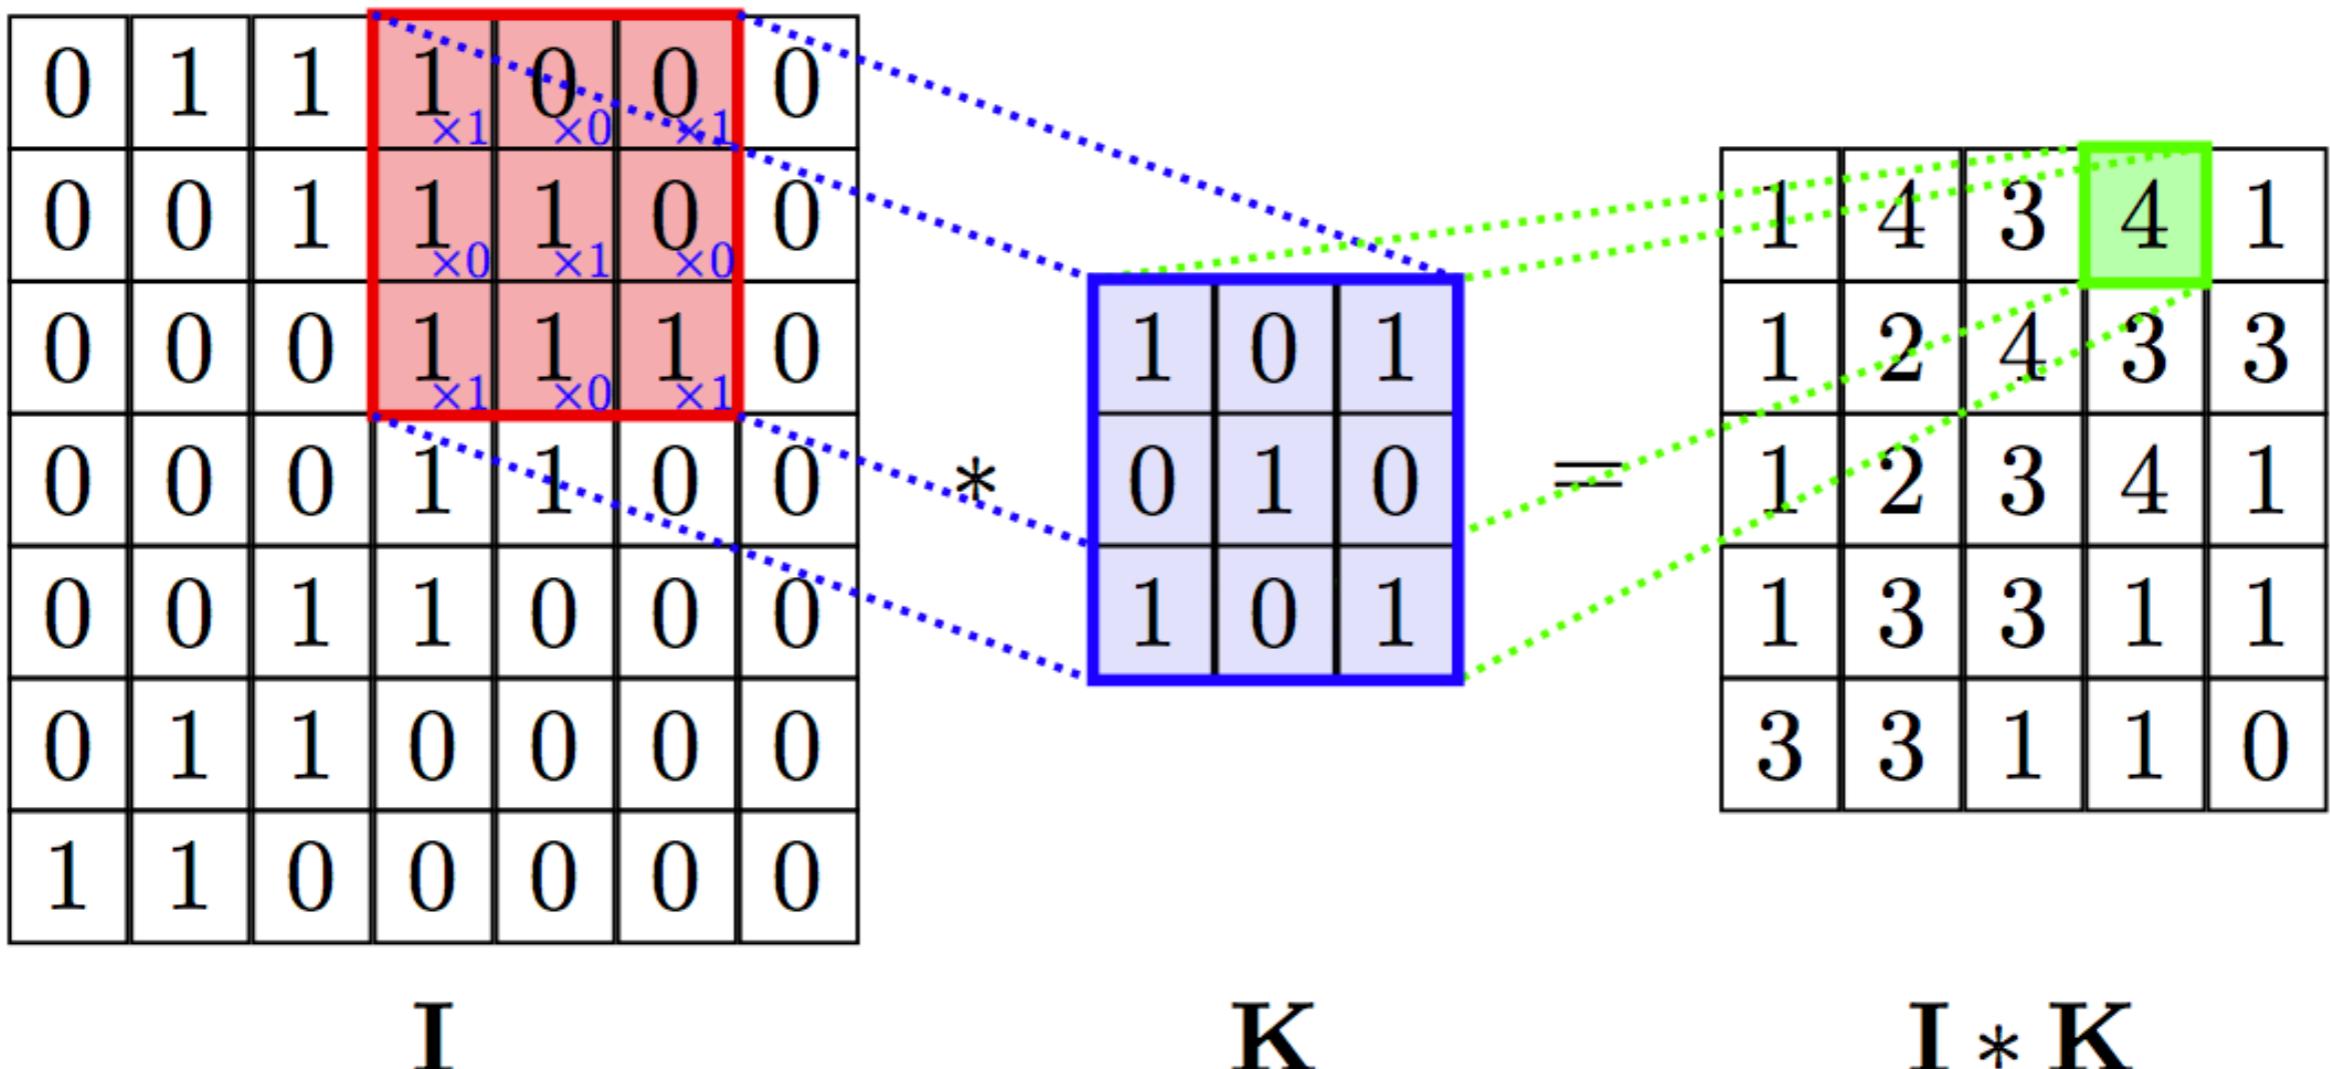


Image  $K$  is convolved with filter  $I$ . Each element in  $I * K$  is a dot product of  $I$  and  $K$  interpreted as vectors.



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	1	1	1
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Filter:

-1 1 -1  
-1 1 -1  
-1 1 -1



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	1	1	1
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Filter:

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix}$$

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} \cdot \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} = 0$$



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	1	1
0	0	0	0	0	0	1	1
0	0	0	0	0	0	0	0

Filter:

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix}$$

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} \cdot \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} = 0$$



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filter:

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix}$$

$$\begin{array}{ccc|ccc|c} -1 & 1 & -1 & 0 & 0 & 0 & 0 \\ -1 & 1 & -1 & 0 & 0 & 0 & = 0 \\ -1 & 1 & -1 & 0 & 0 & 0 & \\ \hline -1 & 1 & -1 & 0 & 1 & 0 & \\ -1 & 1 & -1 & 0 & 1 & 0 & = 3 \\ -1 & 1 & -1 & 0 & 1 & 0 & \end{array}$$



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	1	1	0	0
0	0	0	0	0	0	0	0	0	0

Filter:

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix}$$

$$\begin{array}{ccc|ccc|c} -1 & 1 & -1 & 0 & 0 & 0 & 0 \\ -1 & 1 & -1 & 0 & 0 & 0 & 0 \\ -1 & 1 & -1 & 0 & 0 & 0 & 0 \\ \hline & & & = & & & 0 \end{array}$$
$$\begin{array}{ccc|ccc|c} -1 & 1 & -1 & 0 & 1 & 0 & 0 \\ -1 & 1 & -1 & 0 & 1 & 0 & 0 \\ -1 & 1 & -1 & 0 & 1 & 0 & 0 \\ \hline & & & = & & & 3 \end{array}$$

Result:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	-1	1	-1	0	0	0	0	0
0	0	-2	2	-2	0	0	0	0	0
0	0	-3	3	-3	0	0	0	0	0
0	0	-2	2	-2	0	0	0	0	0
0	0	-1	1	-1	-1	0	-1	0	0
0	0	0	0	0	-1	0	-1	0	0
0	0	0	0	0	-1	0	-1	0	0
0	0	0	0	0	-1	0	-1	0	0

Filter  
fires strongly  
here



# Convolutional Filter

Image:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	1	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filter:

$$\begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix}$$

$$\begin{array}{r} \begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} \cdot \begin{matrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{matrix} = 0 \\ \begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} \cdot \begin{matrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \end{matrix} = 3 \\ \begin{matrix} -1 & 1 & -1 \\ -1 & 1 & -1 \\ -1 & 1 & -1 \end{matrix} \cdot \begin{matrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 0 & 1 & 0 \end{matrix} = 3 \end{array}$$

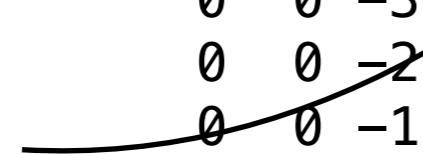
Result:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	-1	1	-1	0	0	0	0	0
0	0	-2	2	-2	0	0	0	0	0
0	0	-3	3	-3	0	0	0	0	0
0	0	-2	2	-2	0	0	0	0	0
0	0	-1	1	-1	0	-1	0	0	0
0	0	0	0	0	-1	0	-1	0	0
0	0	0	0	0	-1	0	-1	0	0
0	0	0	0	0	-1	0	-1	0	0

After ReLU:

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0
0	0	0	0	2	0	0	0	0	0
0	0	0	0	3	0	0	0	0	0
0	0	0	0	2	0	0	0	0	0
0	0	0	0	2	0	0	0	0	0
0	0	0	0	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filter  
fires strongly  
here

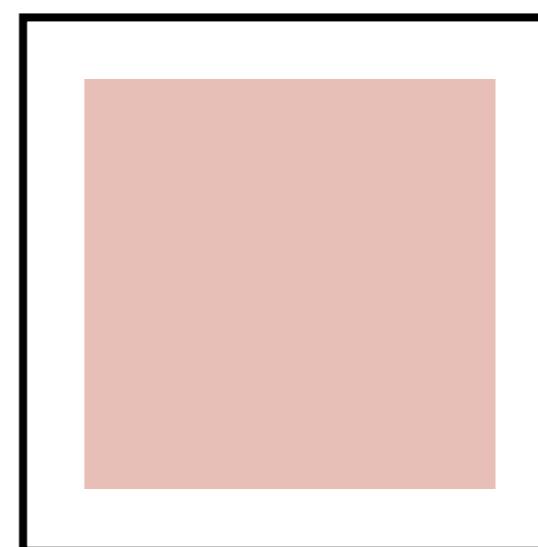




# Example: Detecting Edges

**Image:**

```
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0  
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```



+5

0

-5

**Applying to monochr.  
region:**

$$\begin{array}{cccc} -1 & -1 & -1 & 0 0 0 \\ -1 & 8 & -1 & \cdot 0 0 0 = 0 \\ -1 & -1 & -1 & 0 0 0 \end{array}$$
  
$$\begin{array}{cccc} -1 & -1 & -1 & 1 1 1 \\ -1 & 8 & -1 & \cdot 1 1 1 = 0 \\ -1 & -1 & -1 & 1 1 1 \end{array}$$

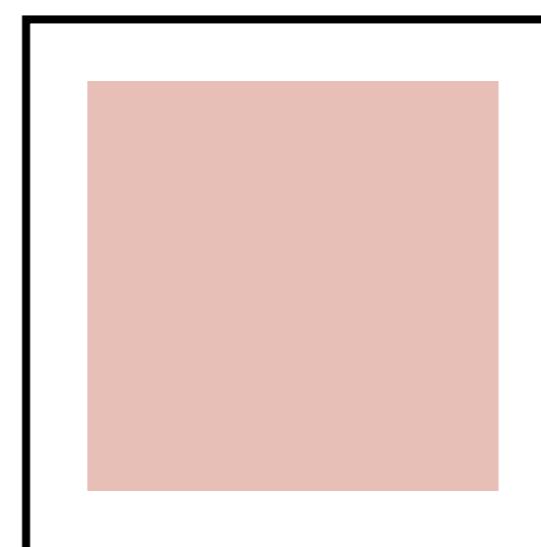
**Filter:**

```
-1 -1 -1  
-1 8 -1  
-1 -1 -1
```



# Example: Detecting Edges

## Image:



1

1

# Applying to monochr. region:

$$\begin{array}{ccc|ccc} -1 & -1 & -1 & 0 & 0 & 0 \\ -1 & 8 & -1 & \cdot & 0 & 0 & 0 \\ -1 & -1 & -1 & & 0 & 0 & 0 \end{array} = 0$$

## Filter:

$$\begin{pmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{pmatrix}$$

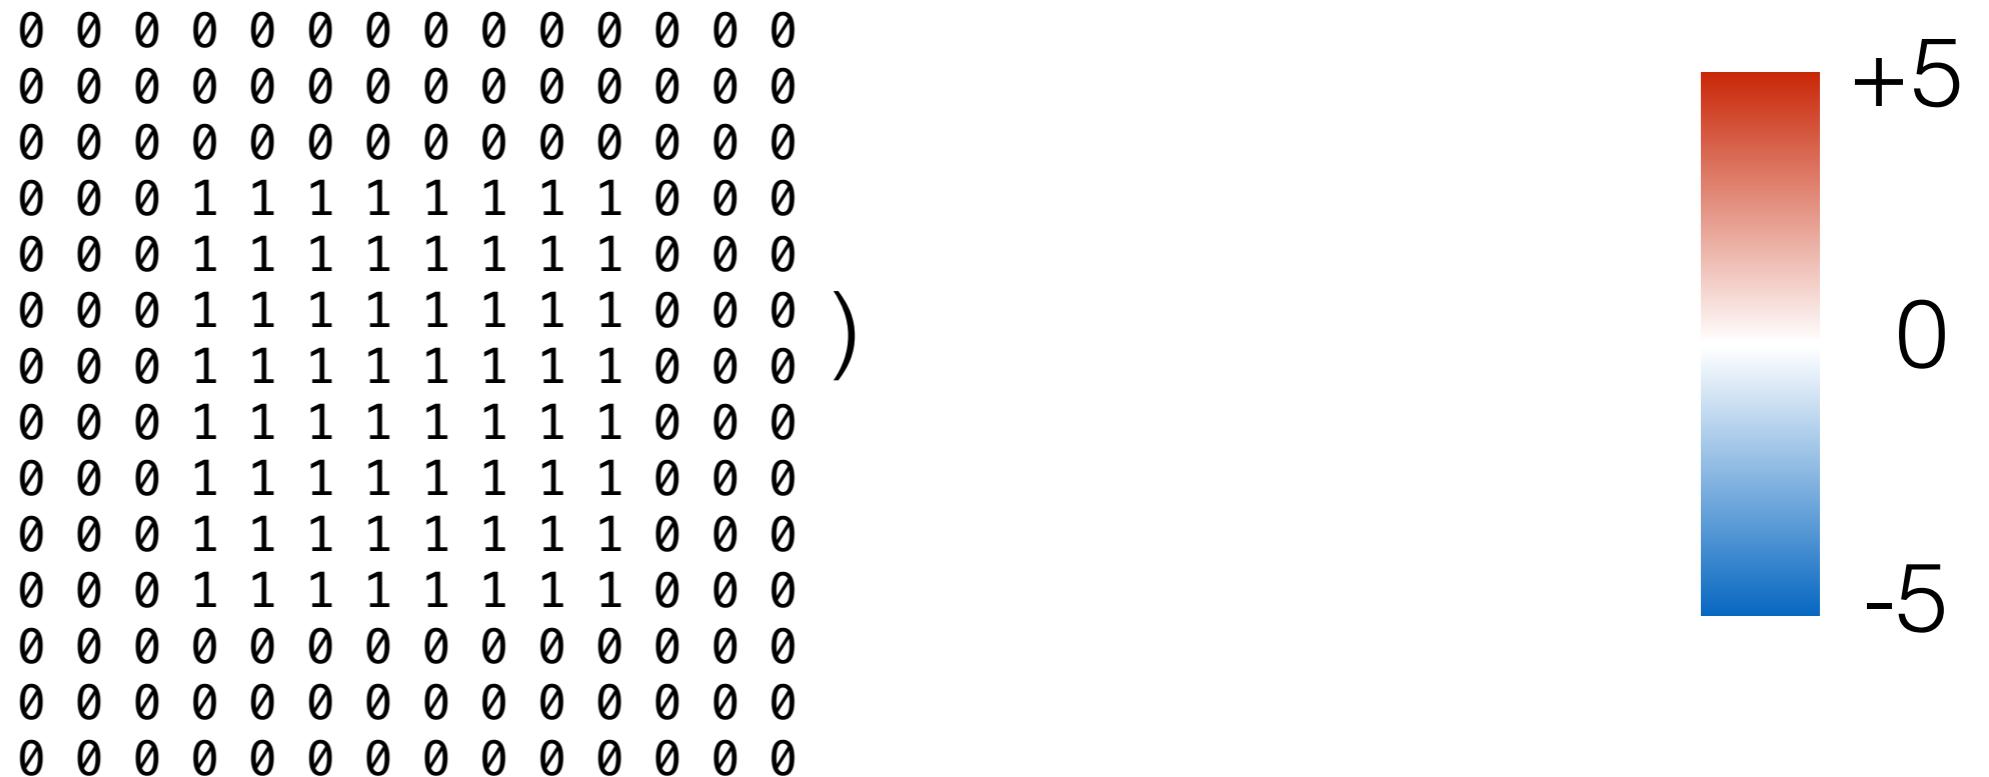
## Applying near an edge:

$$\begin{array}{ccc|ccc} -1 & -1 & -1 & 0 & 0 & 0 \\ -1 & 8 & -1 & 0 & 1 & 1 \\ -1 & -1 & -1 & 0 & 1 & 1 \end{array} = 5$$



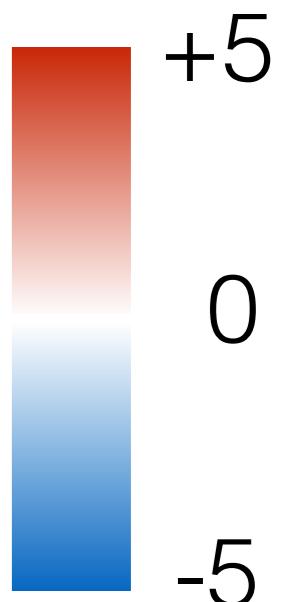
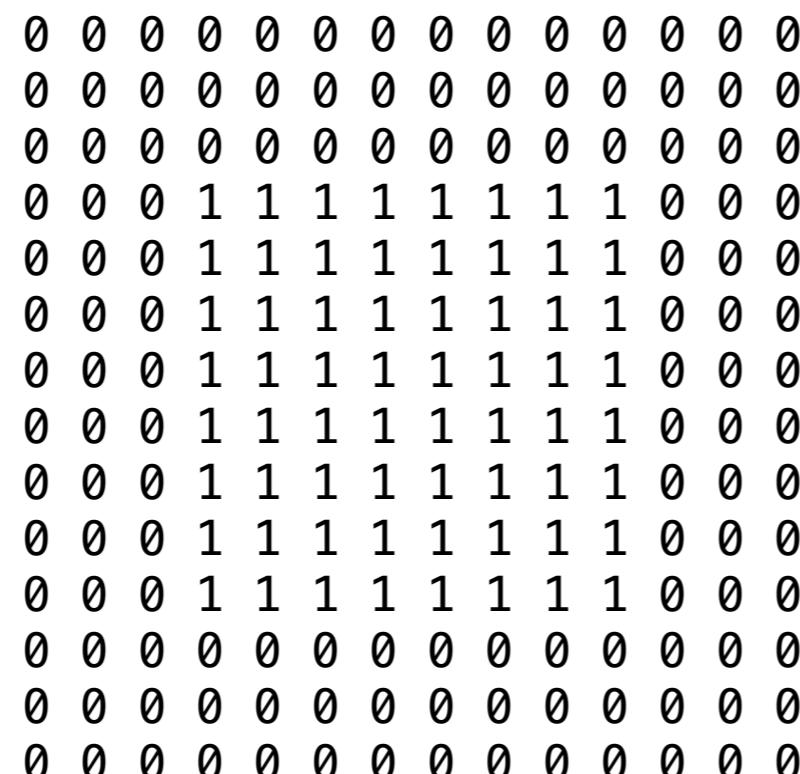
# Example: Detecting Edges

```
conv2d( [ -1 -1 -1 0 0 0 1 1 1 1 1 1 1 1 1 0 0 0  
         -1 8 -1 0 0 0 1 1 1 1 1 1 1 1 1 1 0 0 0  
         -1 -1 -1 0 0 0 1 1 1 1 1 1 1 1 1 1 0 0 0 ] )
```





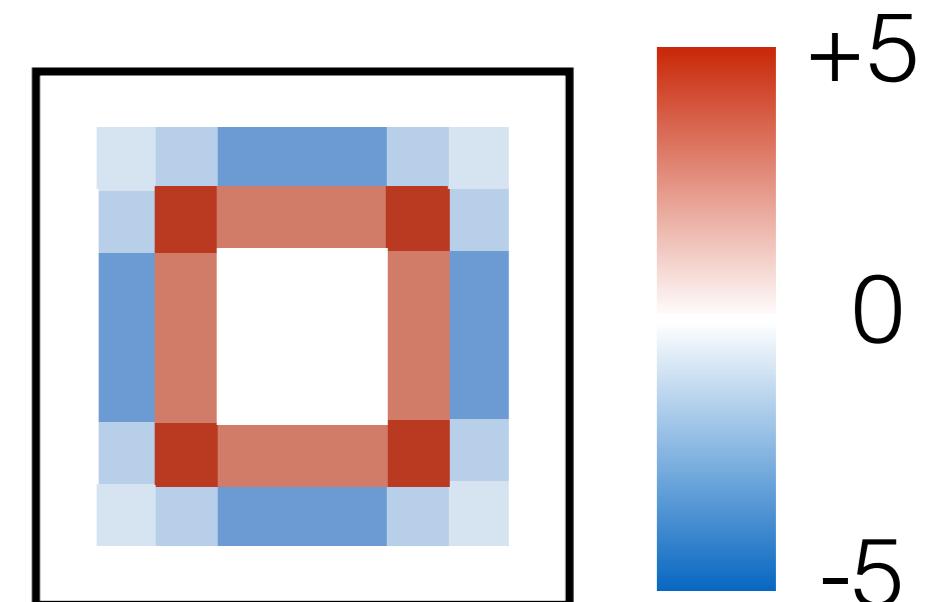
# Example: Detecting Edges



1



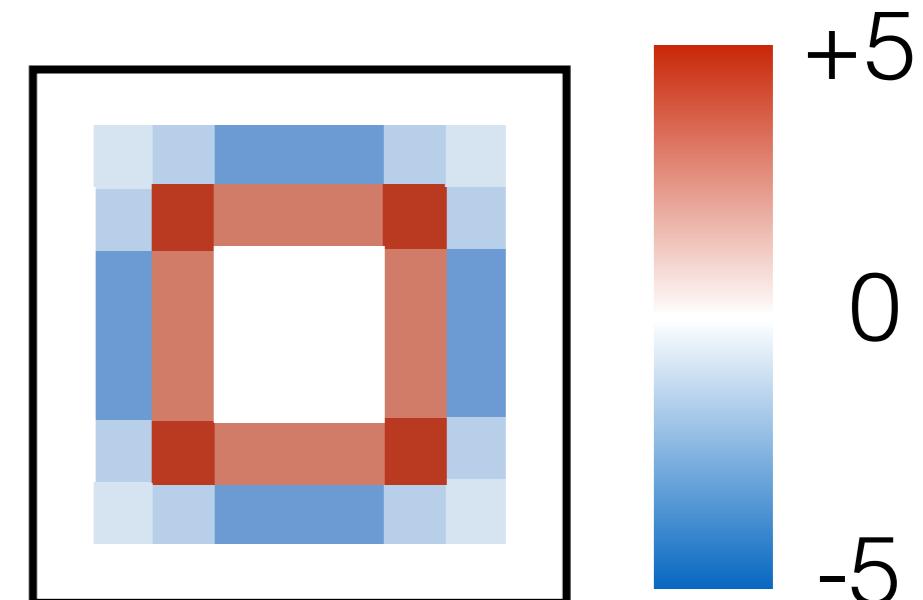
# Example: Detecting Edges



2



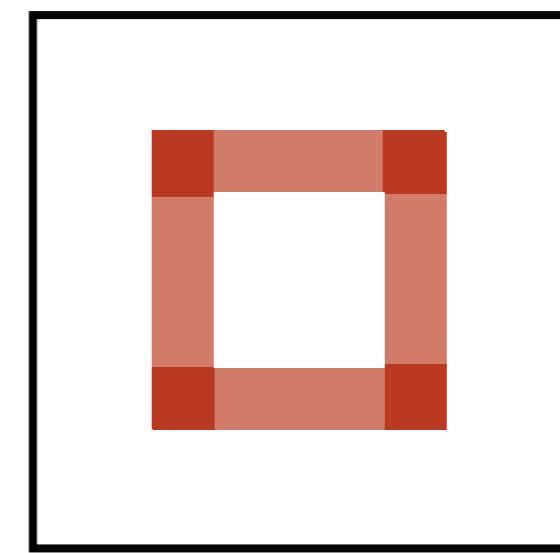
# Example: Detecting Edges



# Add a ReLU:

$$\text{ReLU}(\text{conv2d}(F, I)) =$$

1





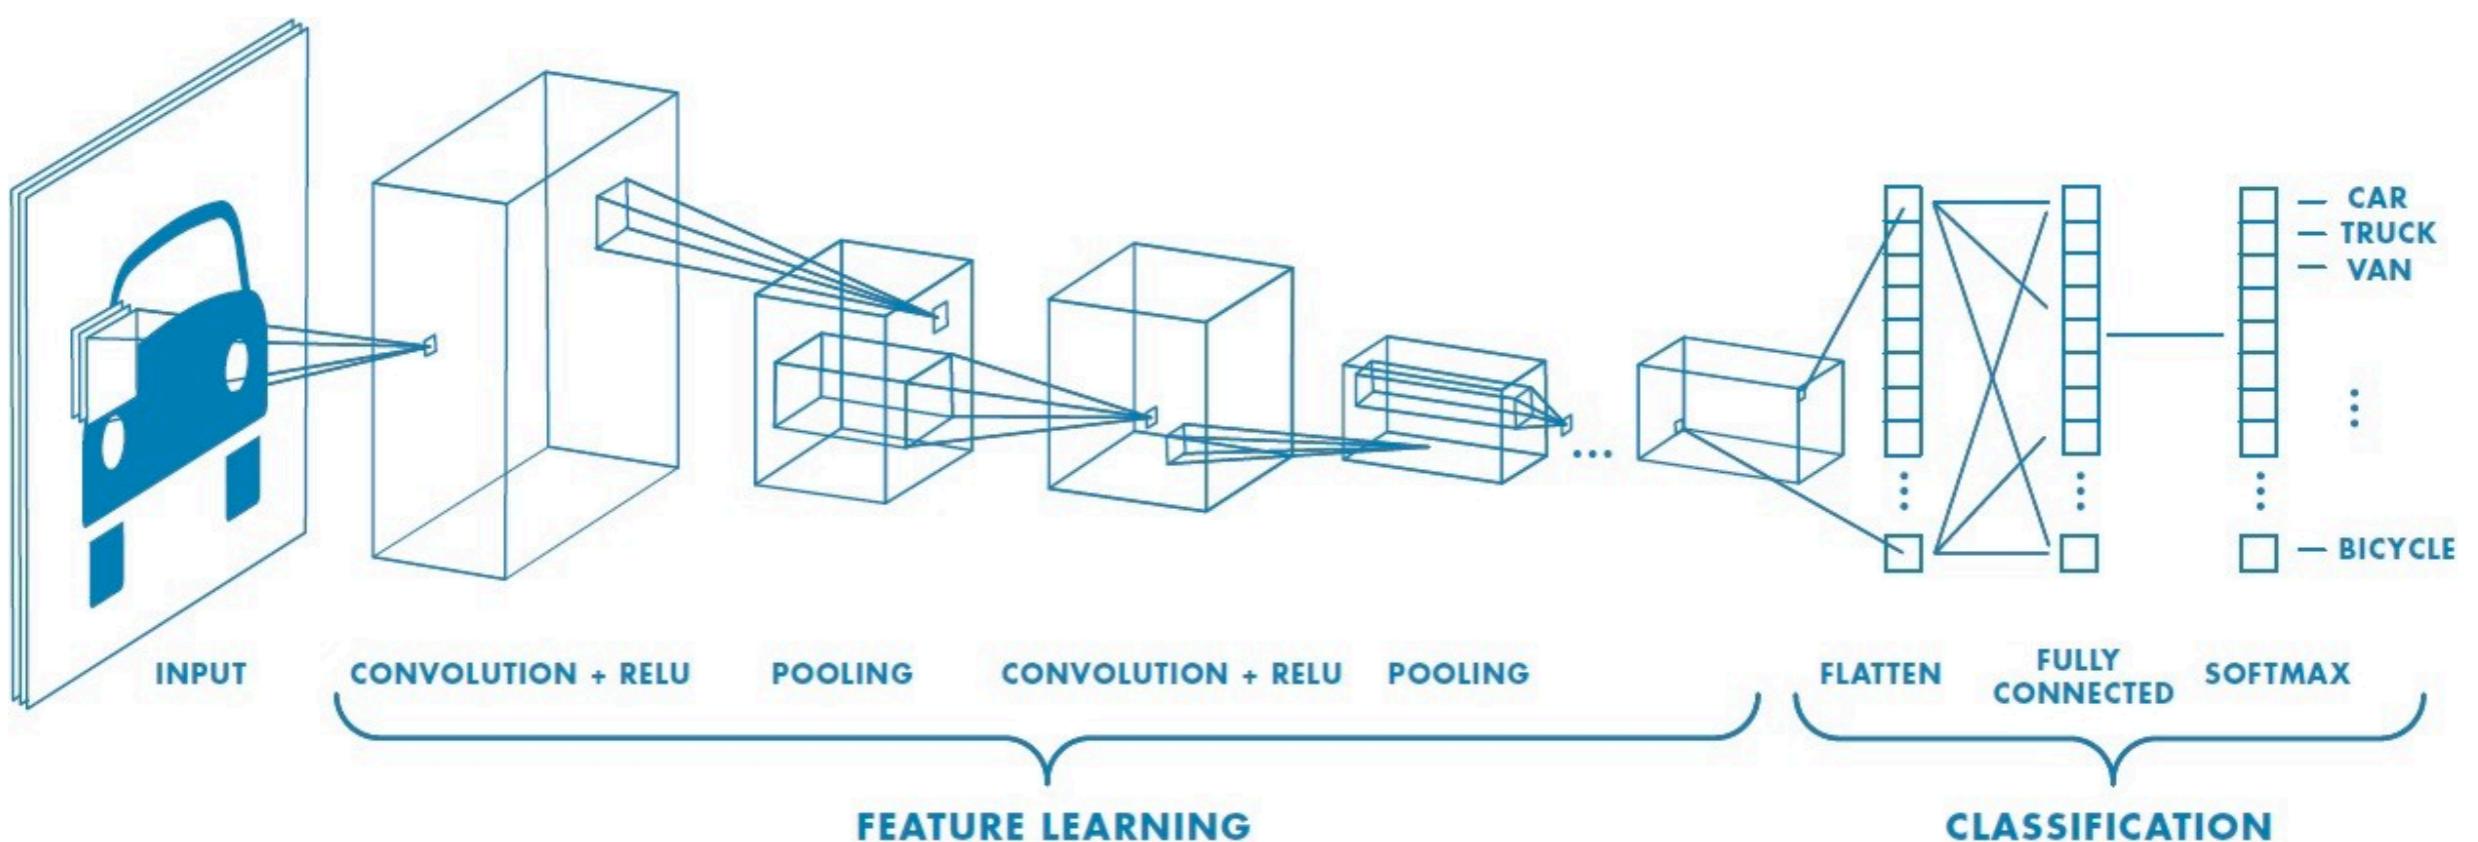
# Example: Detecting Edges



Monotonous regions vanish. Only edges between regions are preserved.

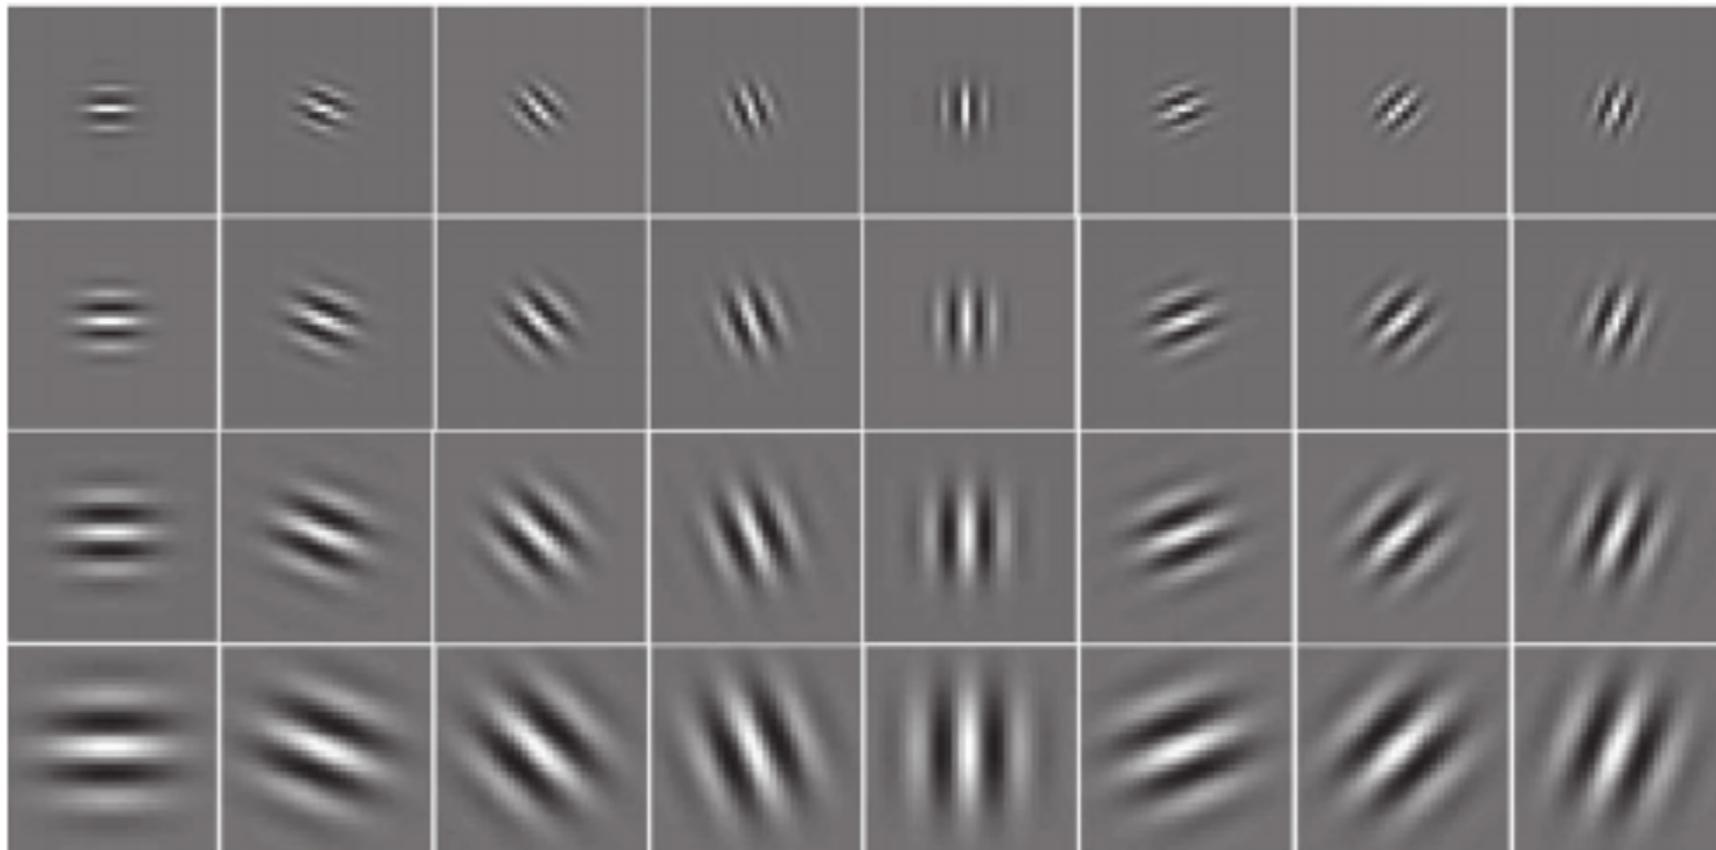


# General CNN architecture





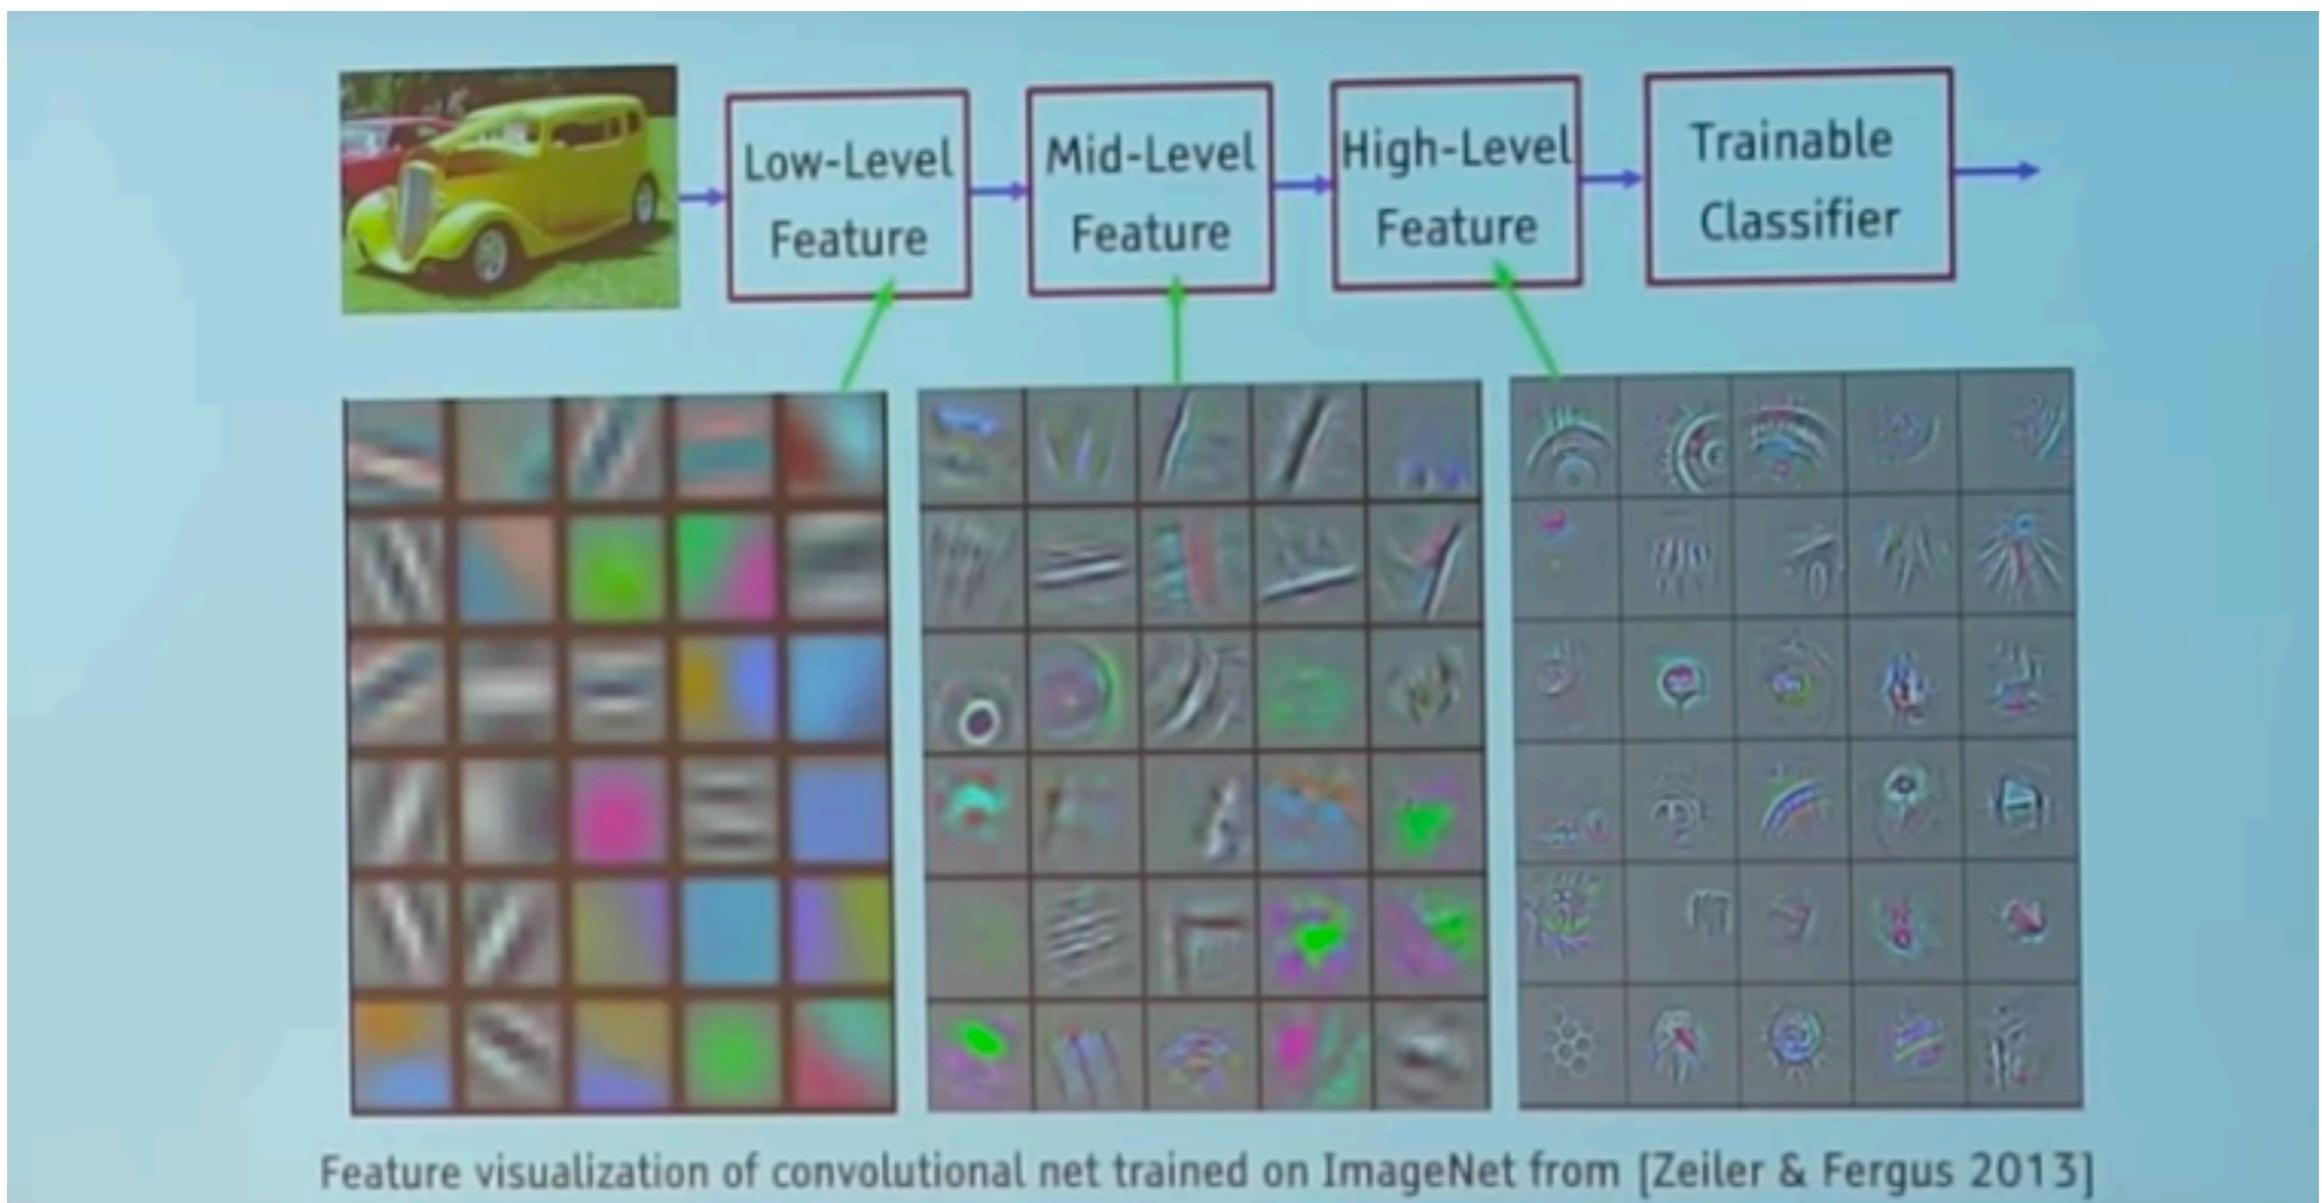
# Gabor Filters



Convolutional neural networks trained for image recognition tend to learn filters which can detect edges.

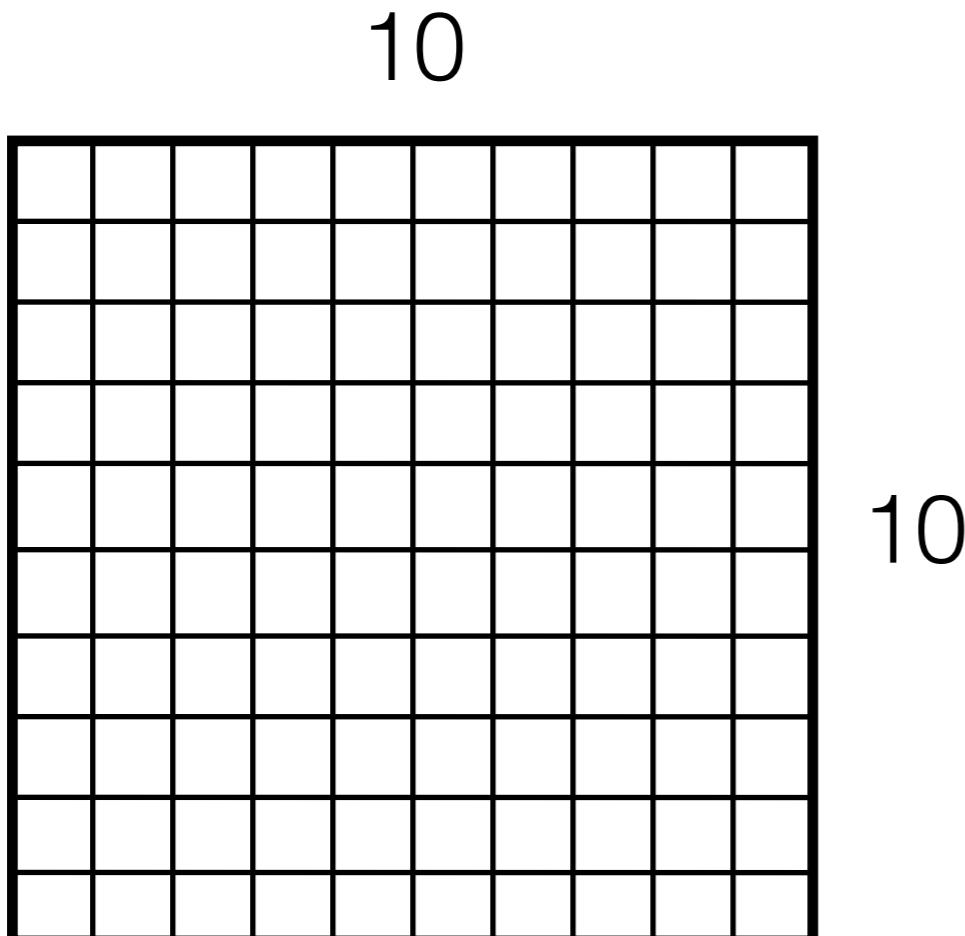


# Filter learned by a CNN



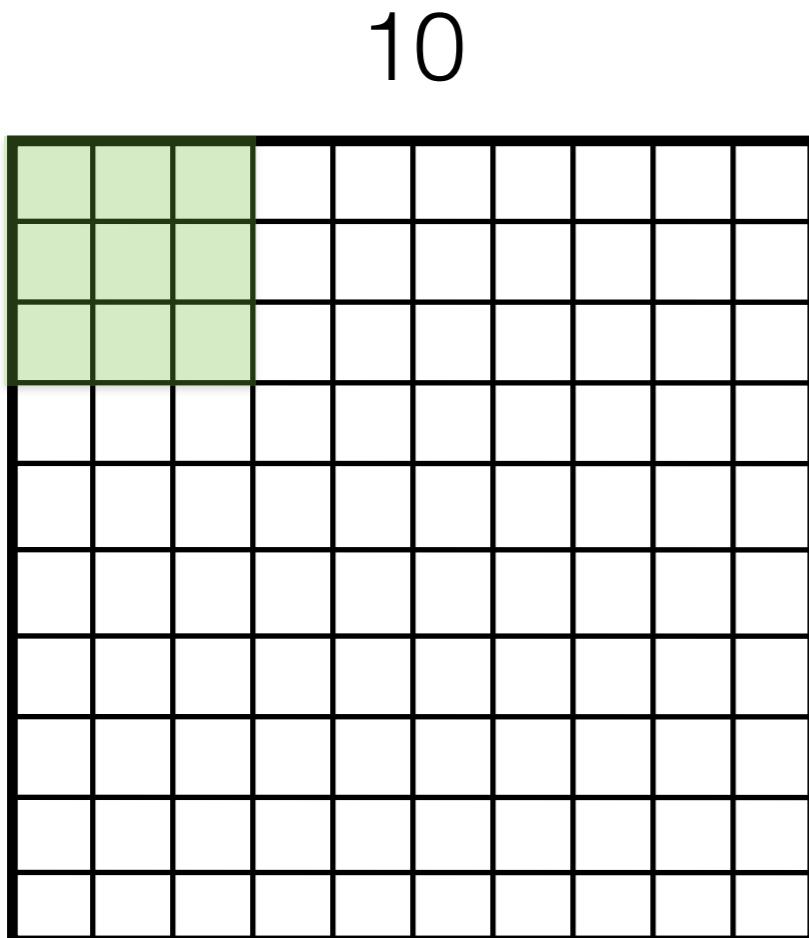


# Convolution Changes the Dimension





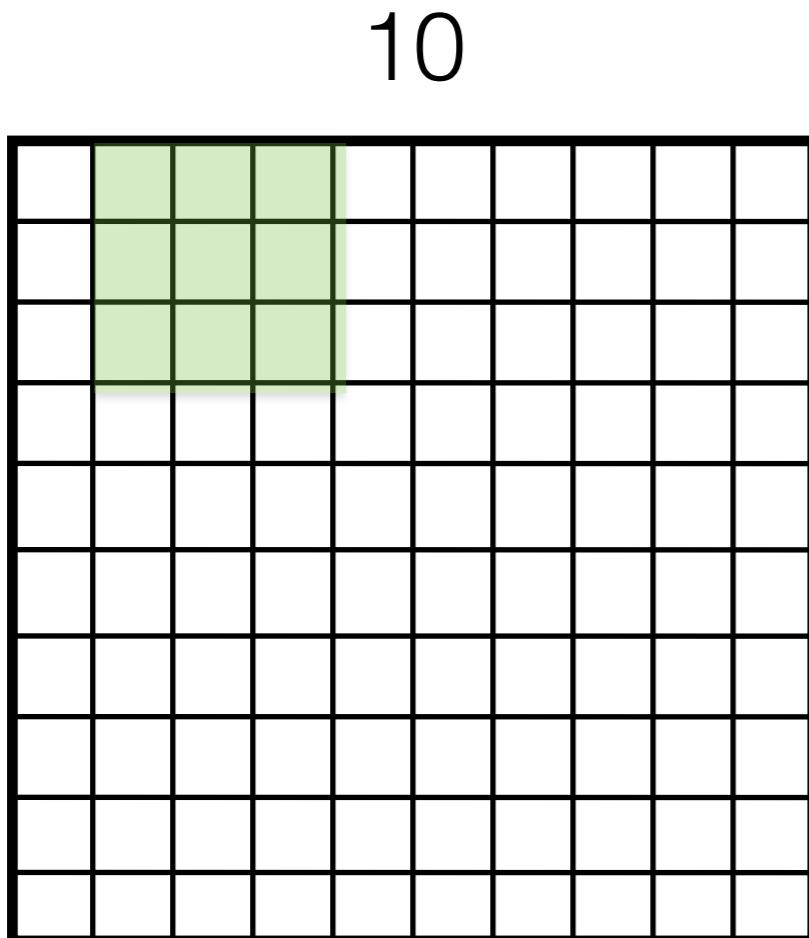
# Convolution Changes the Dimension



2.1  
10



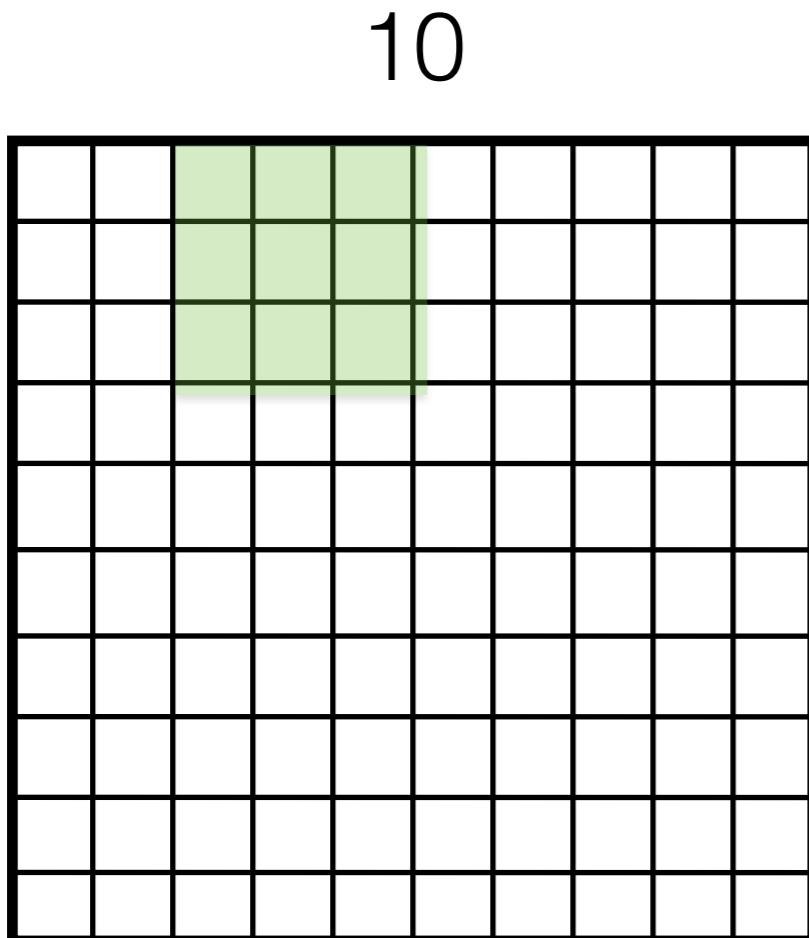
# Convolution Changes the Dimension



2.1 1.2  
10



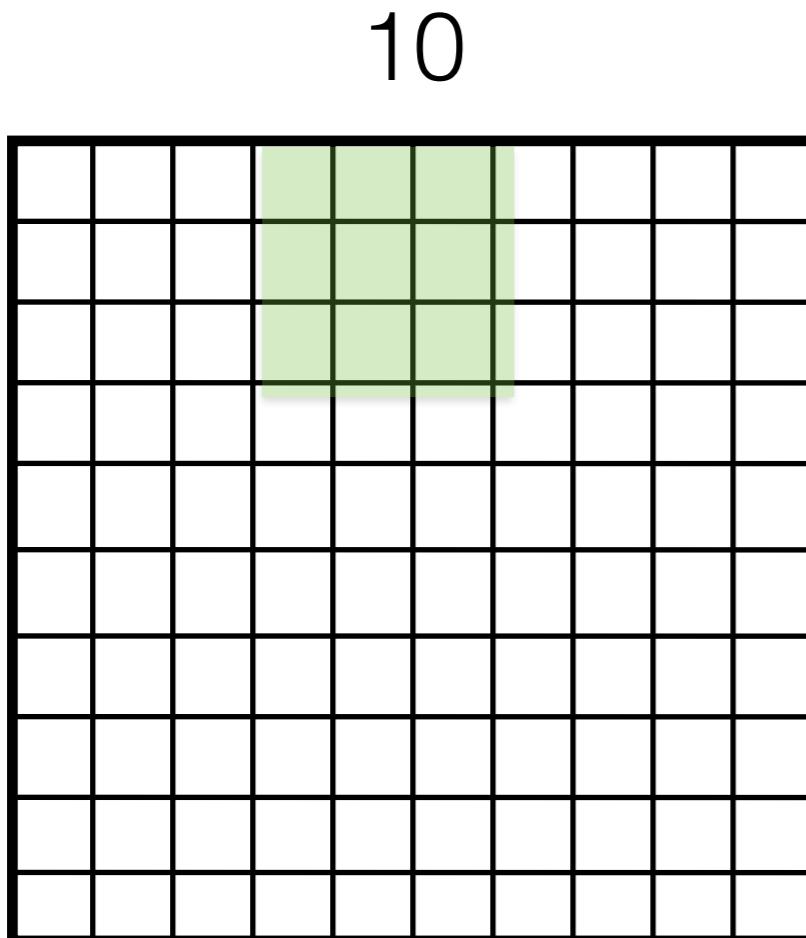
# Convolution Changes the Dimension



2.1 1.2 -2.0  
10



# Convolution Changes the Dimension

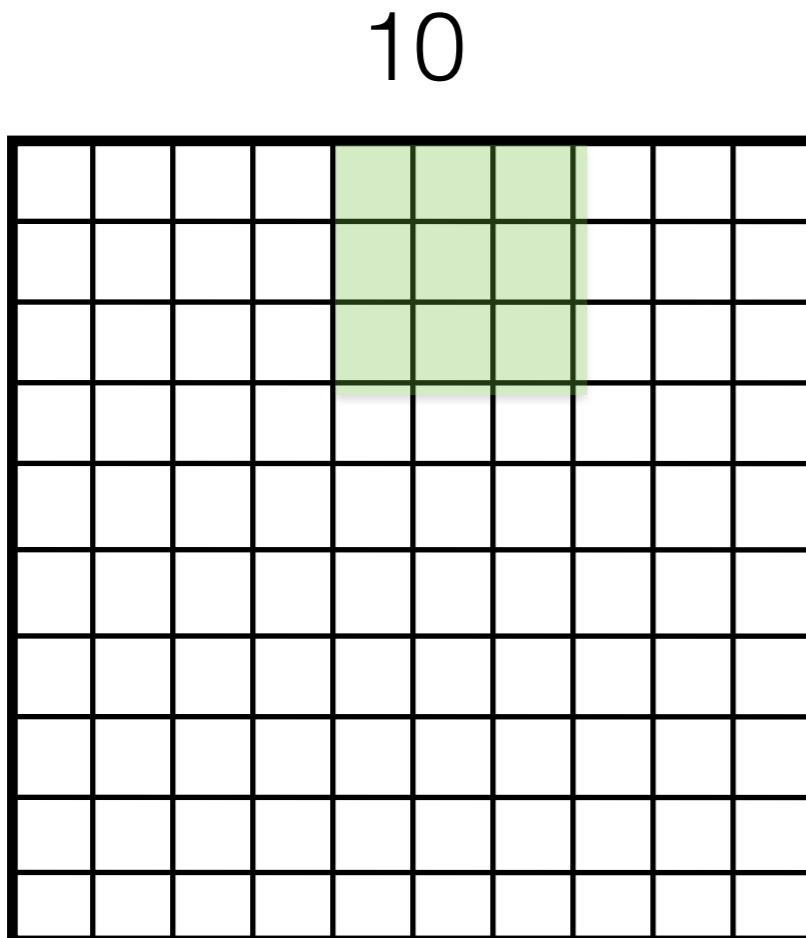


2.1 1.2 -2.0 -2.2

10



# Convolution Changes the Dimension

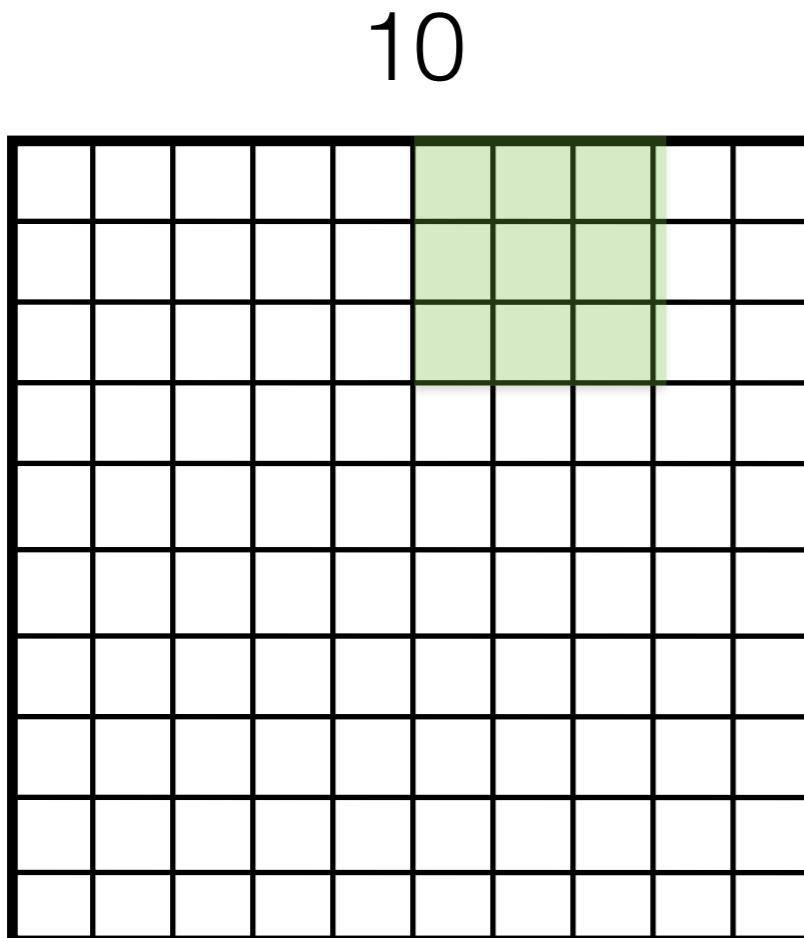


2.1 1.2 -2.0 -2.2 0.7

10



# Convolution Changes the Dimension

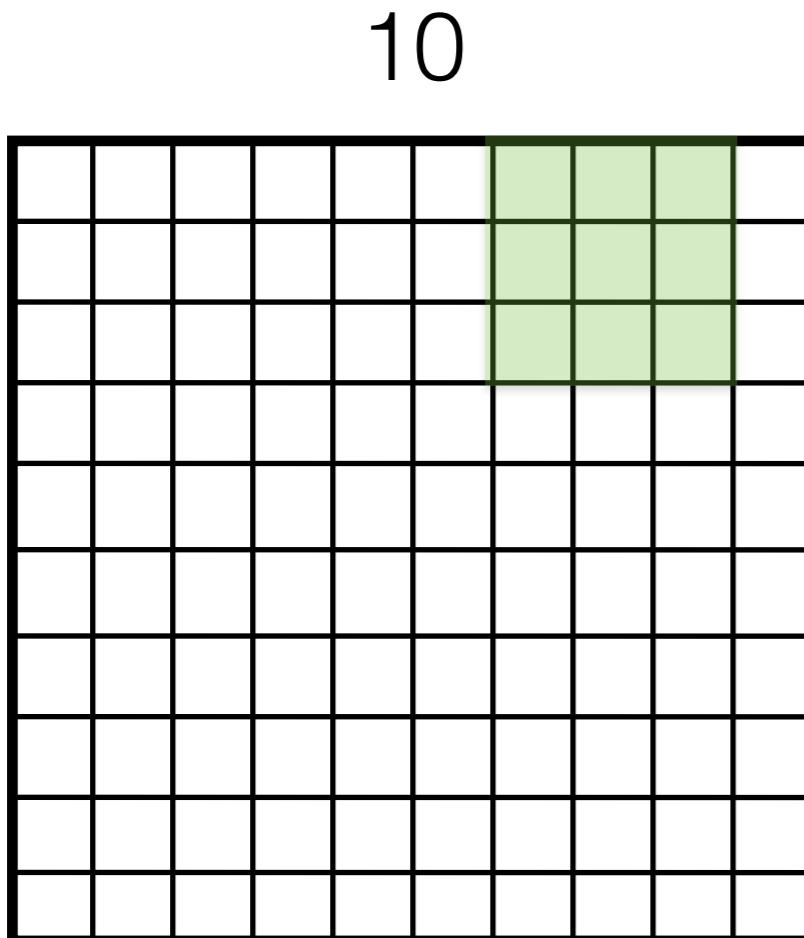


2.1 1.2 -2.0 -2.2 0.7 0.3

10



# Convolution Changes the Dimension

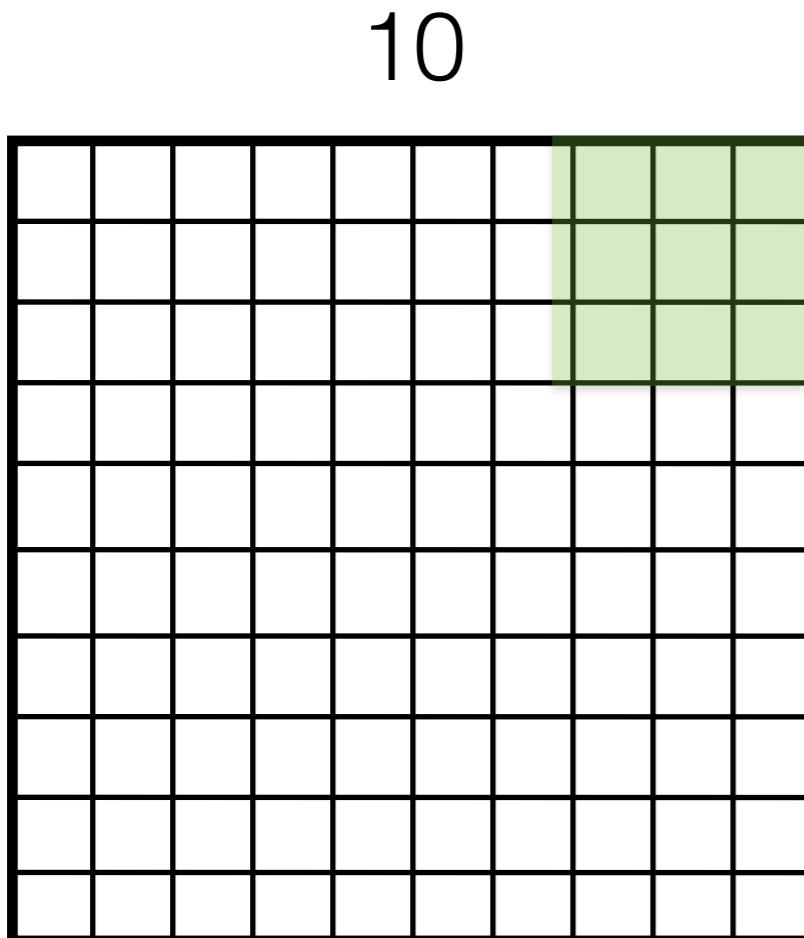


2.1 1.2 -2.0 -2.2 0.7 0.3 -1.0

10



# Convolution Changes the Dimension

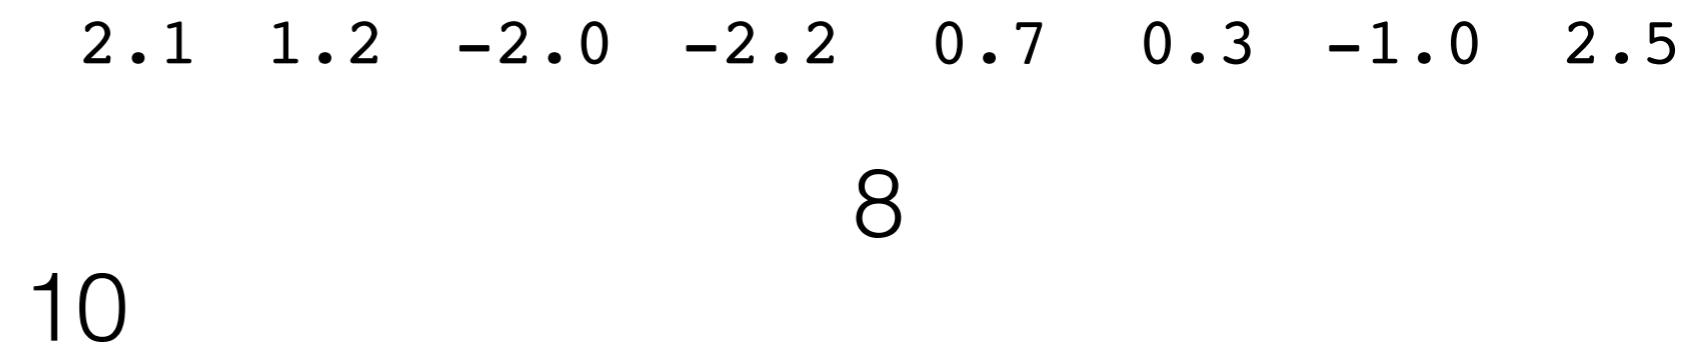
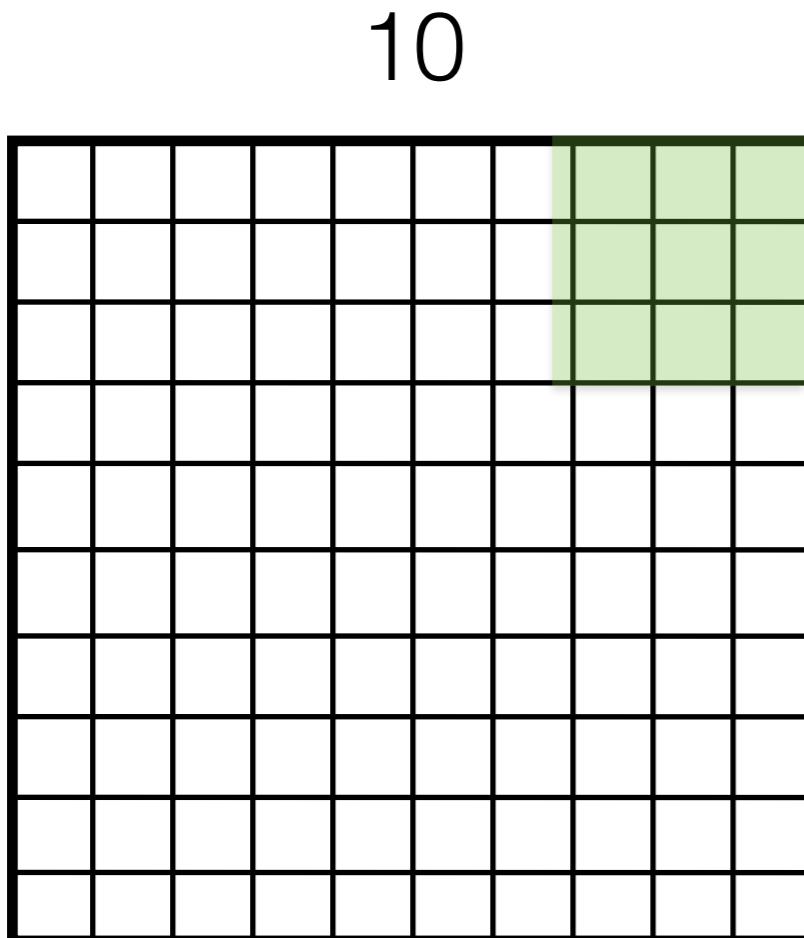


2.1 1.2 -2.0 -2.2 0.7 0.3 -1.0 2.5

10

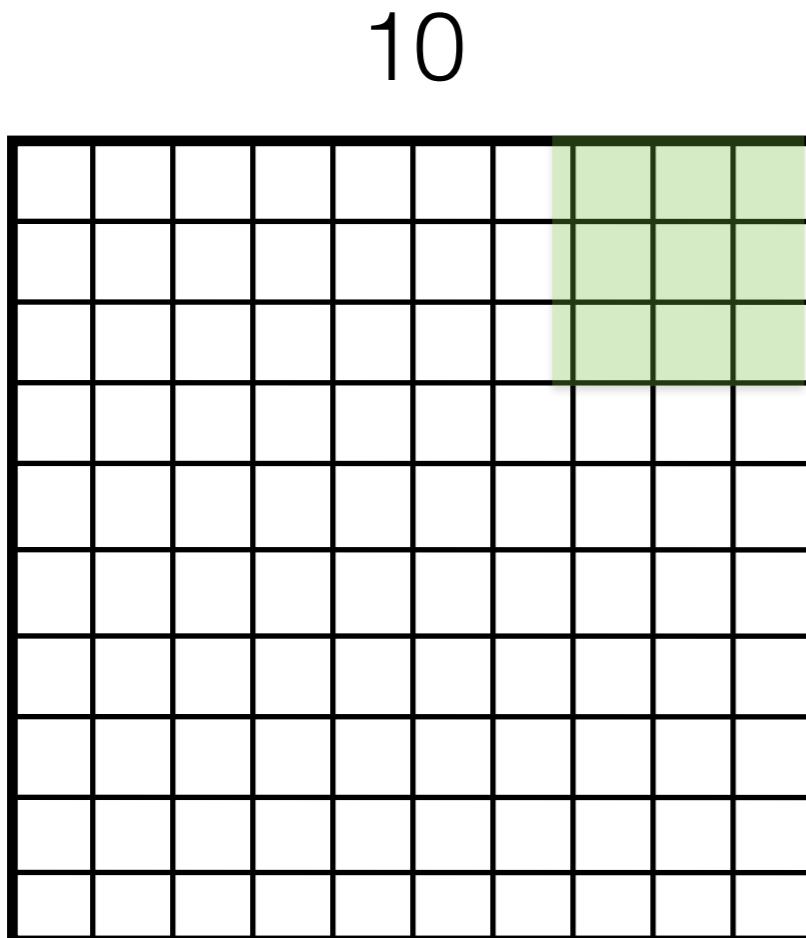


# Convolution Changes the Dimension





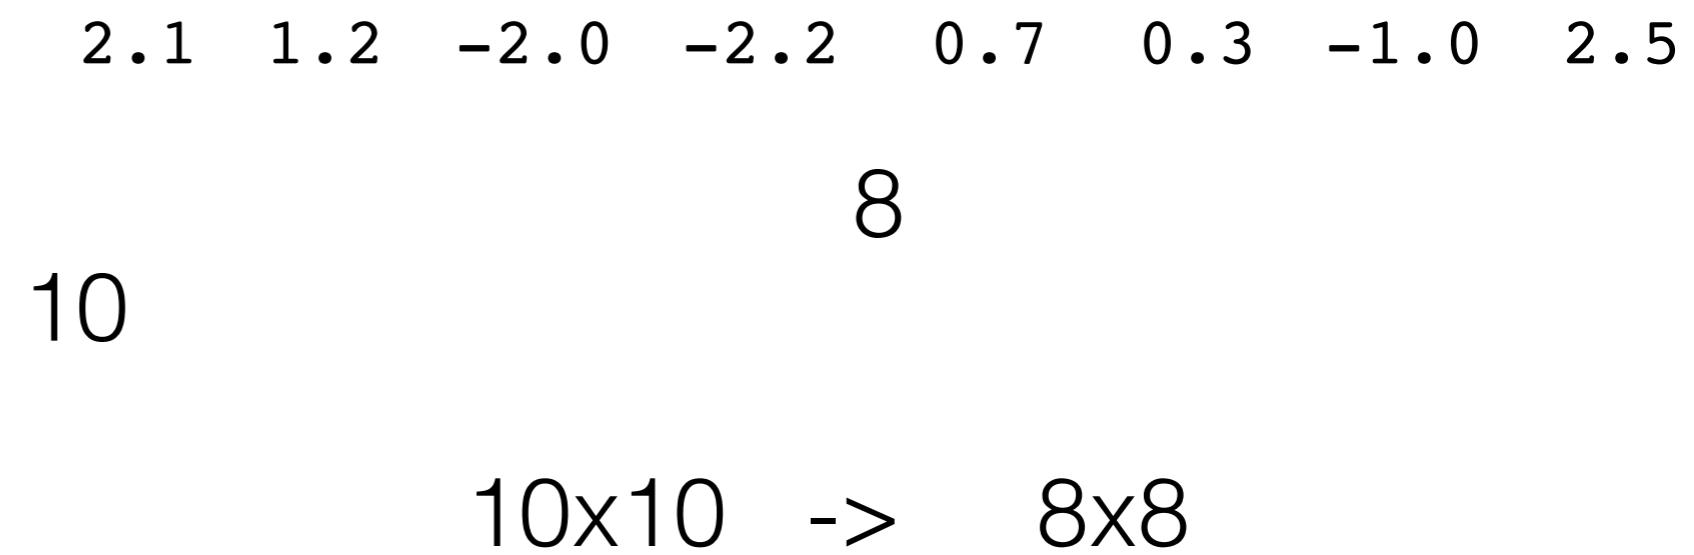
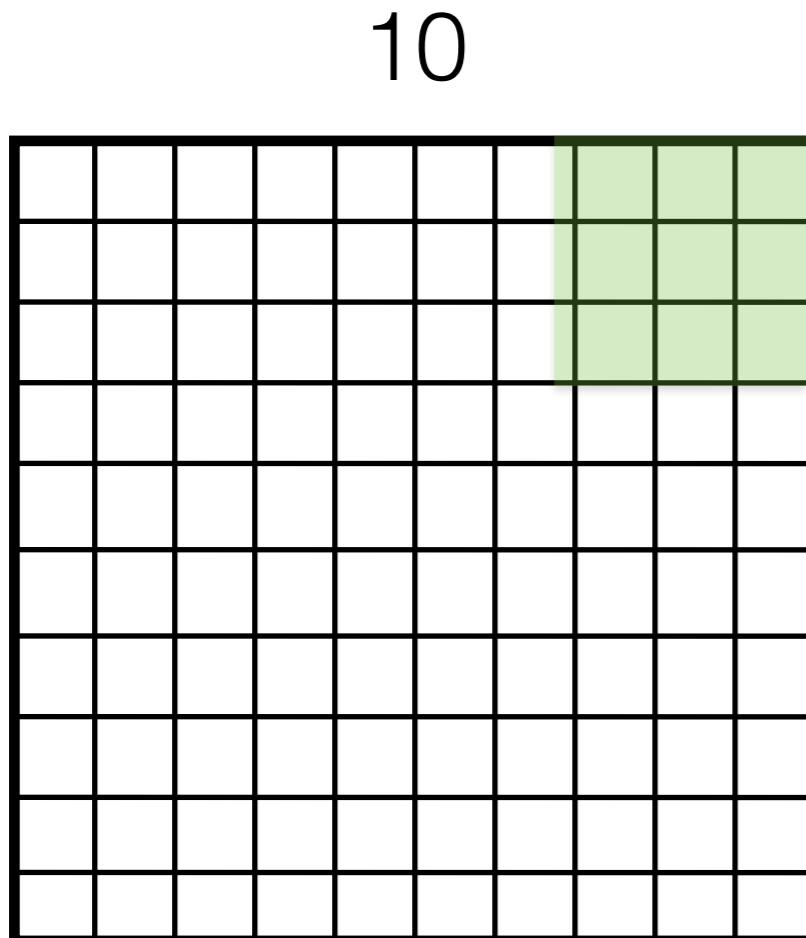
# Convolution Changes the Dimension



2.1 1.2 -2.0 -2.2 0.7 0.3 -1.0 2.5  
8  
10x10 -> 8x8



# Convolution Changes the Dimension



`output_dim = input_dim - filter_size + 1`



# Padding

0	0	0	0	0	0	0	0	0	0	0	0
0											0
0											0
0											0
0											0
0											0
0											0
0											0
0											0
0											0
0	0	0	0	0	0	0	0	0	0	0	0



We add zeros at the boundaries

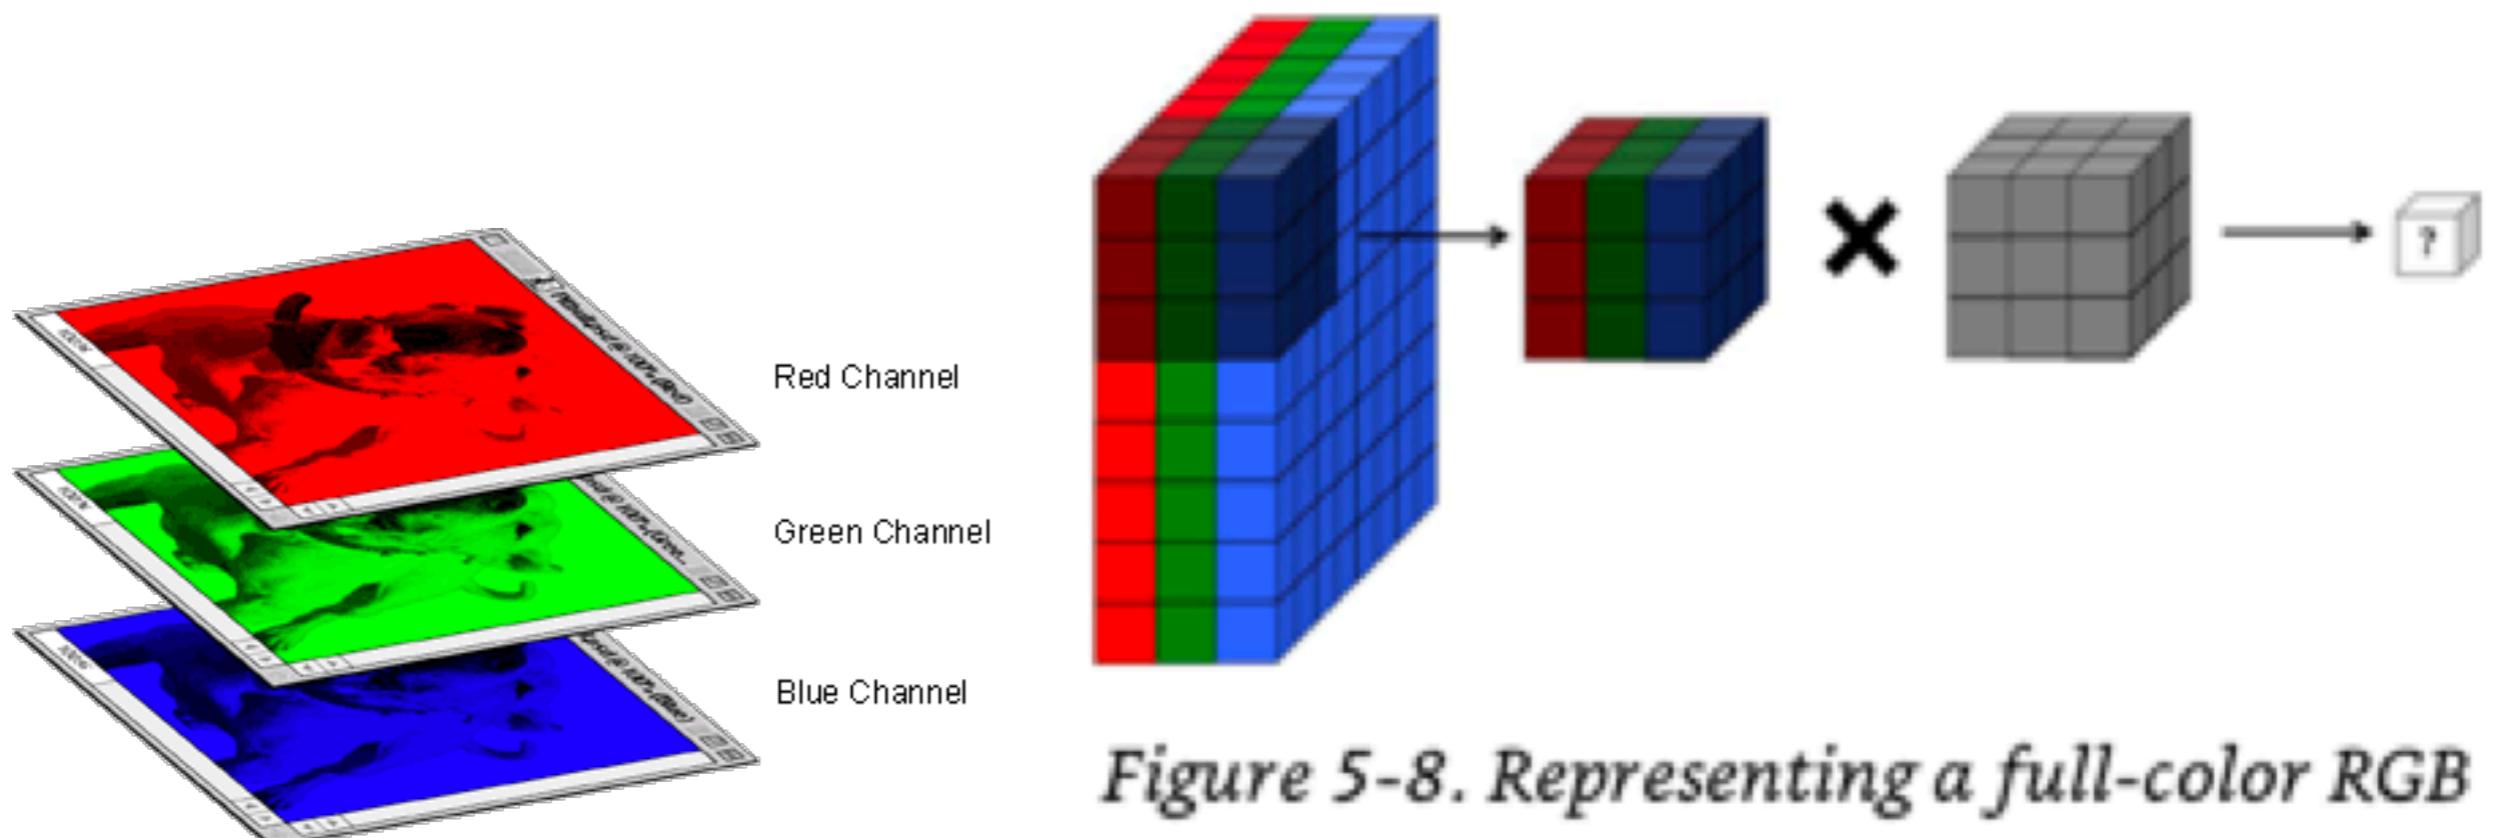
convolution

$10 \times 10 \rightarrow 12 \times 12 \rightarrow 10 \times 10$

padding



# Convolutional Filter with Colors

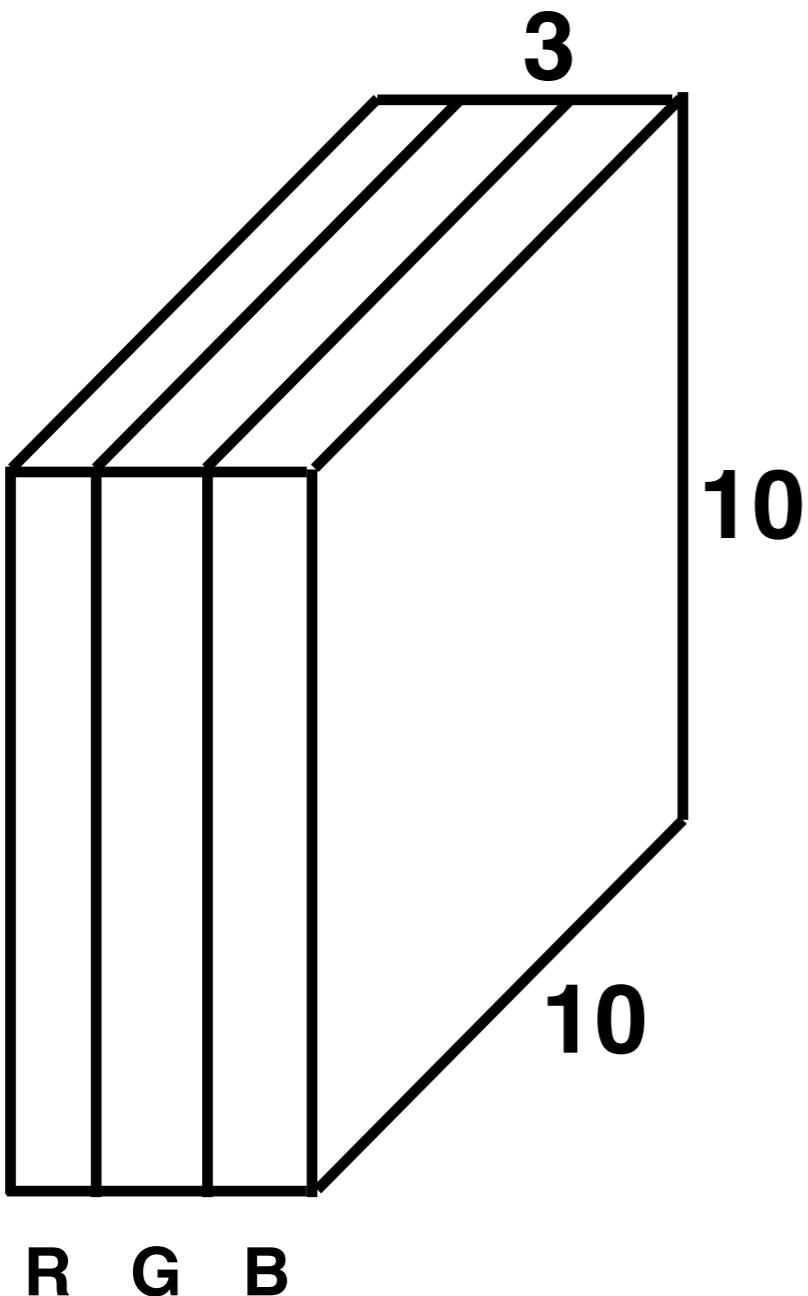


*Figure 5-8. Representing a full-color RGB image as a volume and applying a volumetric convolutional filter*

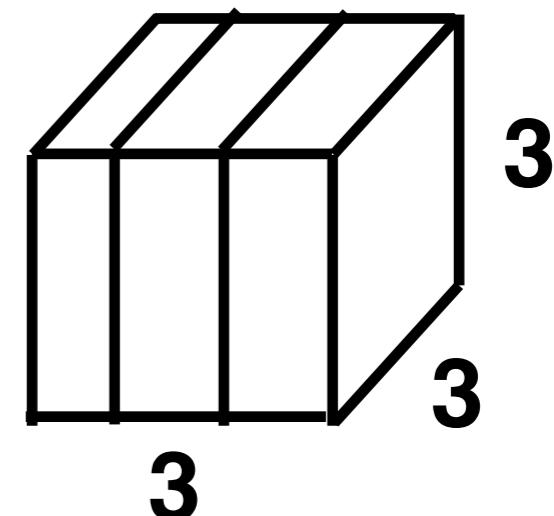


# Example of Convolutional Layer

Input: 12x12x3



Individual filter: 3x3x3

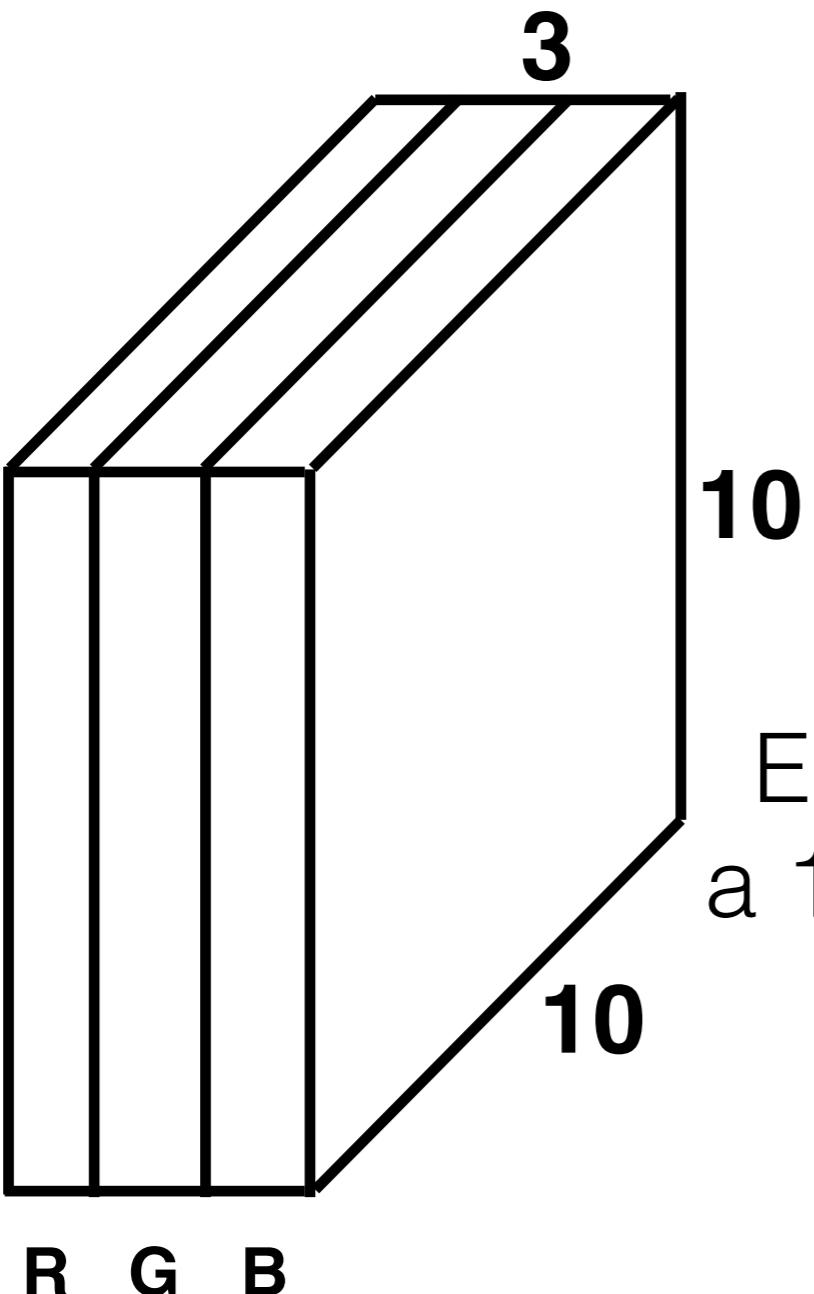


We have 5 filters altogether.

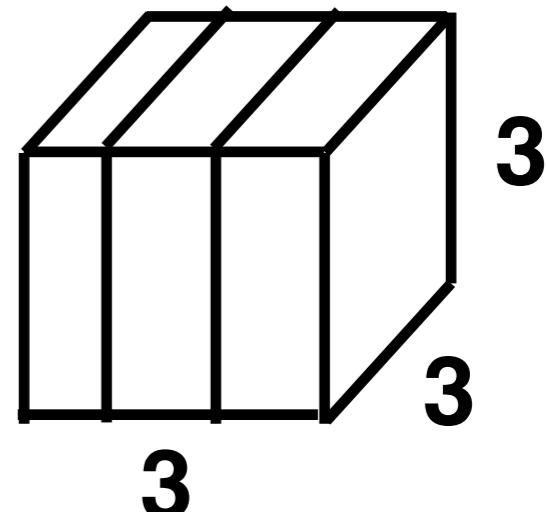


# Example of Convolutional Layer

Input: 12x12x3



Individual filter: 3x3x3



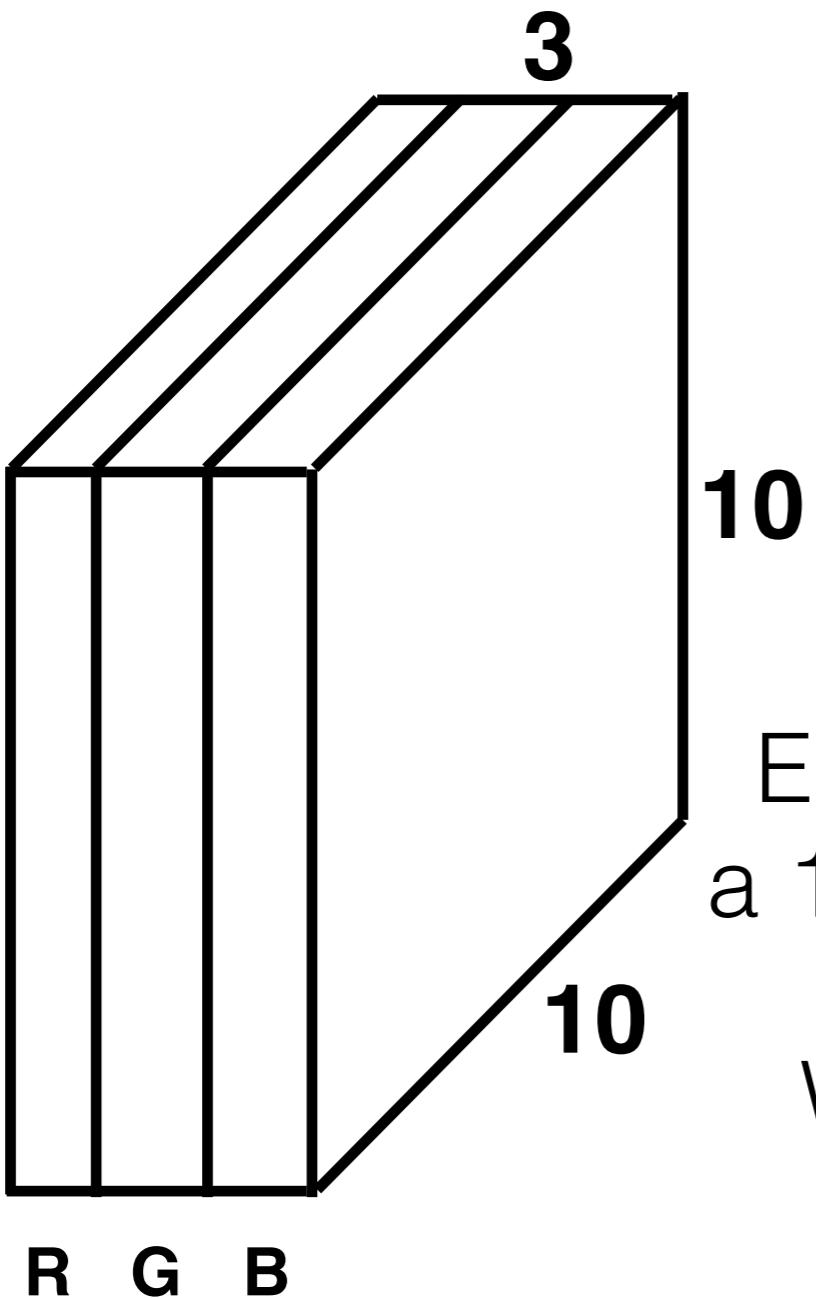
We have 5 filters altogether.

Each filter maps the 10x10x3 input into a 10x10 output because we use padding

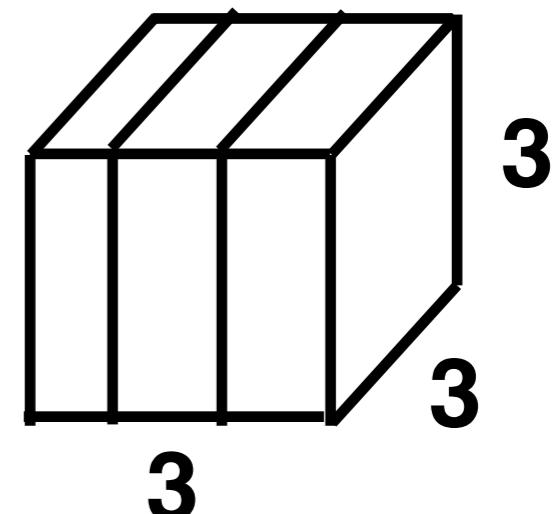


# Example of Convolutional Layer

Input: 12x12x3



Individual filter: 3x3x3



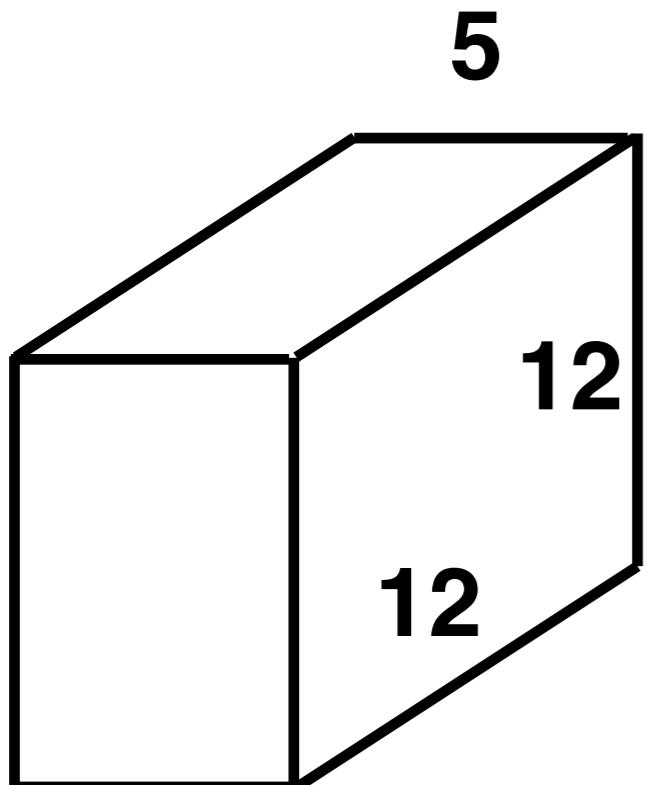
We have 5 filters altogether.

Each filter maps the 10x10x3 input into a 10x10 output because we use padding

We stack the individual outputs to a 10x10x5 output for the entire layer.



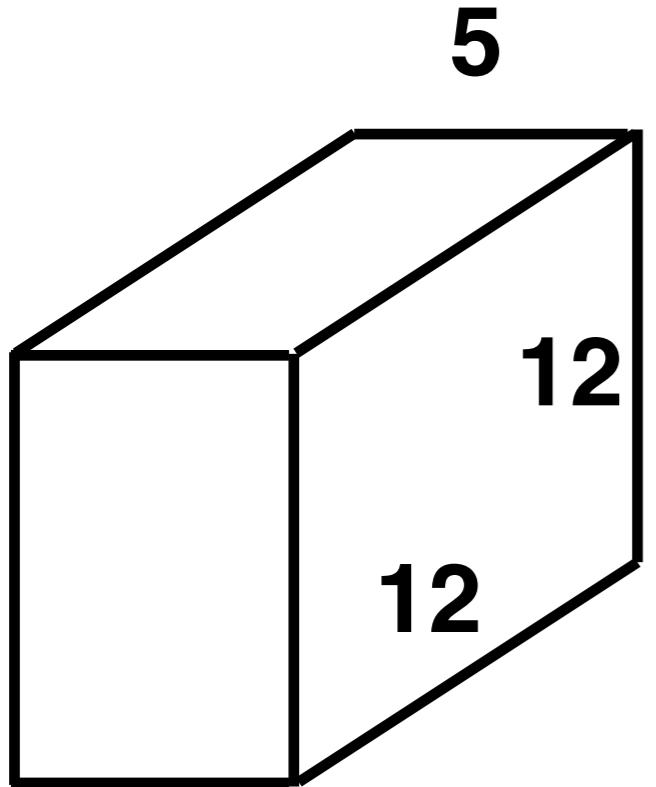
# Applying ReLU



-0.2	0.7	0.9	0.3	0.7	1.1	-0.5	-0.4	0.4	0.6	0.3	1.3
0.7	0.5	-0.2	1.3	0.3	0.8	0.3	-0.1	0.7	1.2	1.2	0.1
0.8	-0.	0.3	0.2	-0.2	0.	1.2	1.3	1.5	0.4	-0.1	0.5
0.7	0.4	0.6	0.7	0.2	1.0	0.2	1.3	0.9	1.0	-0.3	1.0
-0.2	-0.	-0.3	1.4	1.4	0.1	0.5	1.3	1.1	1.	0.1	0.5
0.9	1.1	1.3	0.5	1.2	1.1	0.3	0.4	-0.5	0.6	1.1	-0.4
0.7	-0.1	-0.	0.9	-0.5	0.1	-0.	0.8	0.1	1.2	-0.2	-0.
0.5	0.7	0.5	1.0	0.1	-0.2	0.8	0.8	0.5	-0.	0.5	-0.4
-0.5	0.3	0.4	0.2	-0.2	-0.4	-0.	1.3	-0.3	1.3	1.4	1.5
-0.4	1.2	0.4	-0.4	0.2	1.4	0.	-0.3	0.6	1.2	1.0	0.1
1.0	1.2	0.6	1.1	-0.5	-0.1	0.	0.2	1.0	0.8	1.3	-0.0
0.9	1.1	1.5	-0.1	0.7	-0.2	1.4	1.1	1.2	1.0	-0.1	0.3



# Applying ReLU



-0.2	0.7	0.9	0.3	0.7	1.1	-0.5	-0.4	0.4	0.6	0.3	1.3
0.7	0.5	-0.2	1.3	0.3	0.8	0.3	-0.1	0.7	1.2	1.2	0.1
0.8	-0.	0.3	0.2	-0.2	0.	1.2	1.3	1.5	0.4	-0.1	0.5
0.7	0.4	0.6	0.7	0.2	1.0	0.2	1.3	0.9	1.0	-0.3	1.0
-0.2	-0.	-0.3	1.4	1.4	0.1	0.5	1.3	1.1	1.	0.1	0.5
0.9	1.1	1.3	0.5	1.2	1.1	0.3	0.4	-0.5	0.6	1.1	-0.4
0.7	-0.1	-0.	0.9	-0.5	0.1	-0.	0.8	0.1	1.2	-0.2	-0.
0.5	0.7	0.5	1.0	0.1	-0.2	0.8	0.8	0.5	-0.	0.5	-0.4
-0.5	0.3	0.4	0.2	-0.2	-0.4	-0.	1.3	-0.3	1.3	1.4	1.5
-0.4	1.2	0.4	-0.4	0.2	1.4	0.	-0.3	0.6	1.2	1.0	0.1
1.0	1.2	0.6	1.1	-0.5	-0.1	0.	0.2	1.0	0.8	1.3	-0.0
0.9	1.1	1.5	-0.1	0.7	-0.2	1.4	1.1	1.2	1.0	-0.1	0.3

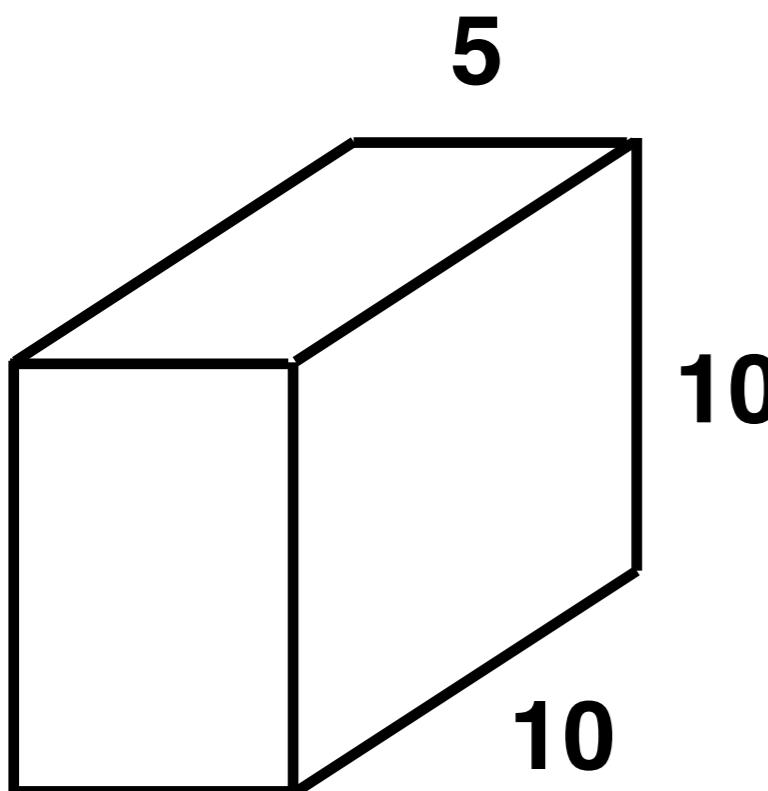


0	0.7	0.9	0.3	0.7	1.1	0	0	0.4	0.6	0.3	1.3
0.7	0.5	0	1.3	0.3	0.8	0.3	0	0.7	1.2	1.2	0.1
0.8	0	0.3	0.2	0	0	1.2	1.3	1.5	0.4	0	0.5
0.7	0.4	0.6	0.7	0.2	1.0	0.2	1.3	0.9	1.0	0	1.0
0	0	0	1.4	1.4	0.1	0.5	1.3	1.1	1.0	0.1	0.5
0.9	1.1	1.3	0.5	1.2	1.1	0.3	0.4	0	0.6	1.1	0
0.7	0	0	0.9	0	0.1	0	0.8	0.1	1.2	0	0
0.5	0.7	0.5	1.	0.1	0	0.8	0.8	0.5	0	0.5	0
0	0.3	0.4	0.2	0	0	0	1.3	0	1.3	1.4	1.5
0	1.2	0.4	0	0.2	1.4	0.	0	0.6	1.2	1.	0.1
1.	1.2	0.6	1.1	0	0	0.	0.2	1.	0.8	1.3	0
0.9	1.1	1.5	0	0.7	0	1.4	1.1	1.2	1.	0	0.3



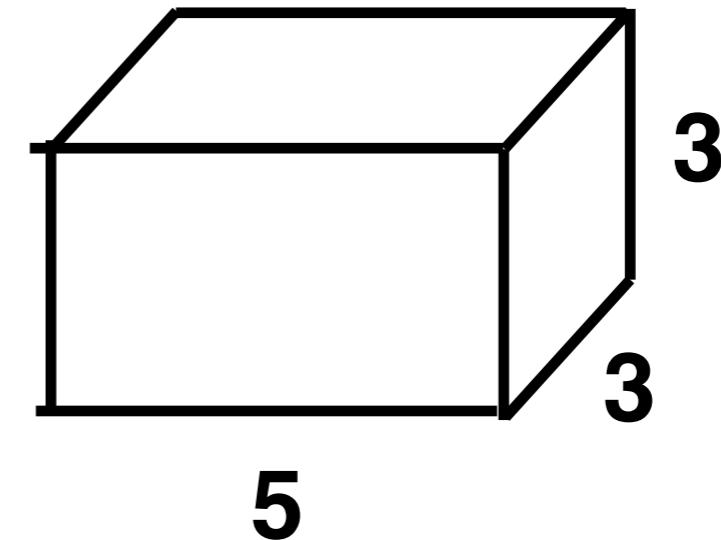
# Additional Layers

Input:



- Since the input has depth 5, we need filters of depth 5.
- The number of filters here again determines the depth of the output.

Filter:





# Stride

0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	1	1	1	1	1	1	1	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0

Stride 1

Stride 2

Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

# Stride 2

## Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

# Stride 1

## Stride 2

## Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

## Stride 2

# Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

## Stride 2

# Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

## Stride 2

# Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

# Stride 2

# Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

# Stride 2

# Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

# Stride 2

## Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Stride

## Stride 1

# Stride 2

## Stride 3

Greater stride can improve efficiency because we compute fewer dot products.

Greater stride also decreases resolution



# Fully specifying a Convolutional Layer

1. We need to specify the input dimension,
2. filter size,
3. number of filters,
4. padding size and
5. stride.



# Fully specifying a Convolutional Layer

1. Input size:  $12 \times 12 \times 3$ ,
2. filter size:  $3 \times 3 \times 3$ ,
3. number of filters: 5
4. padding: 1
5. stride: 1

Input:

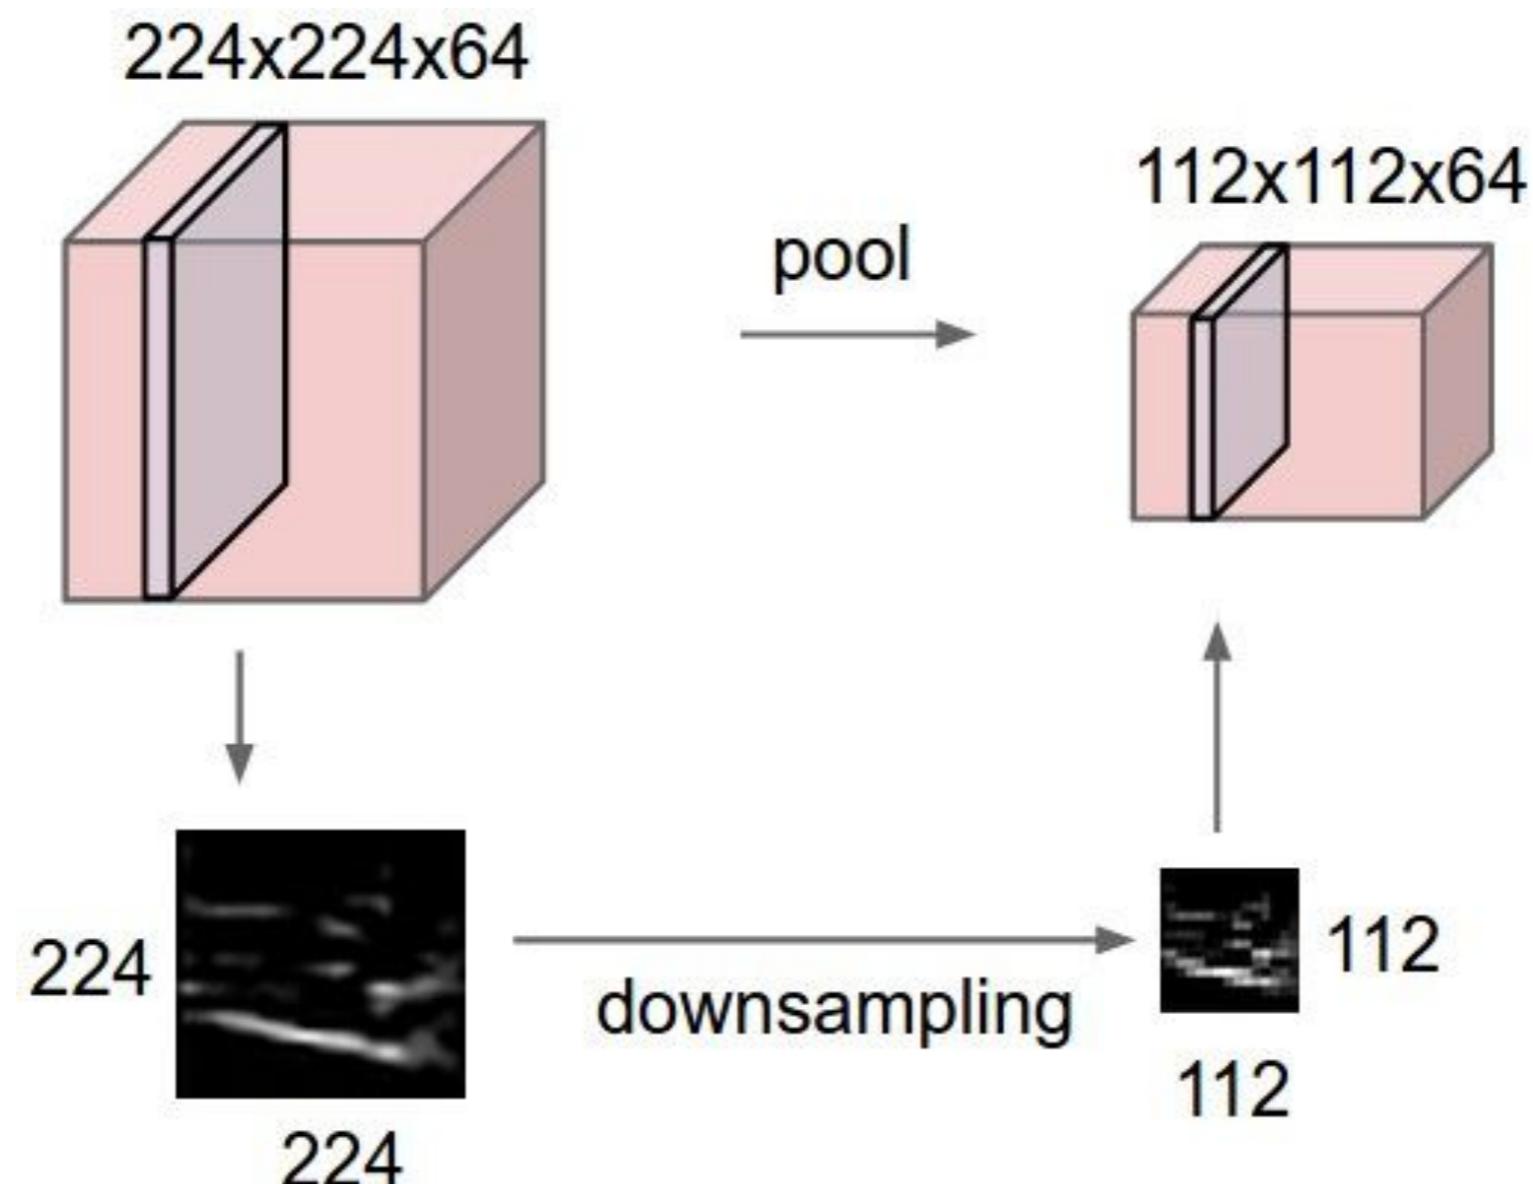
$12 \times 12 \times 3$

Output:

$12 \times 12 \times 5$

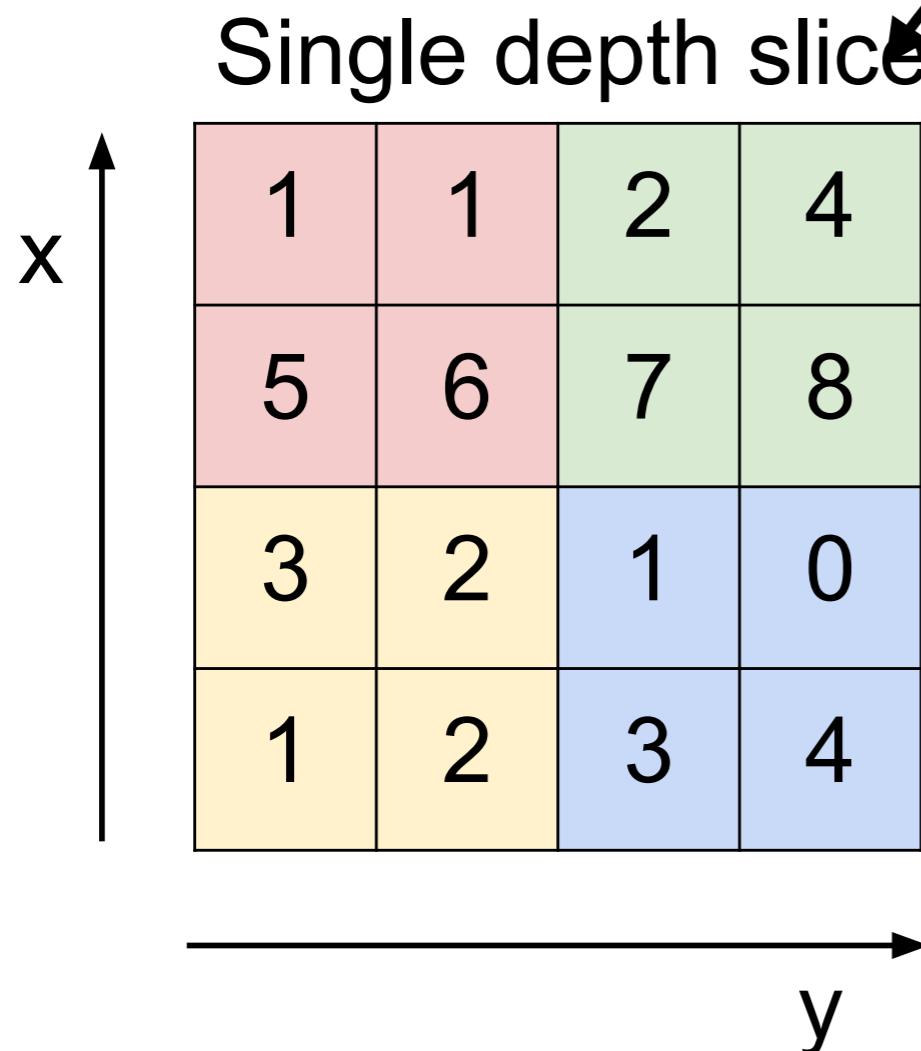


# Pooling Layers

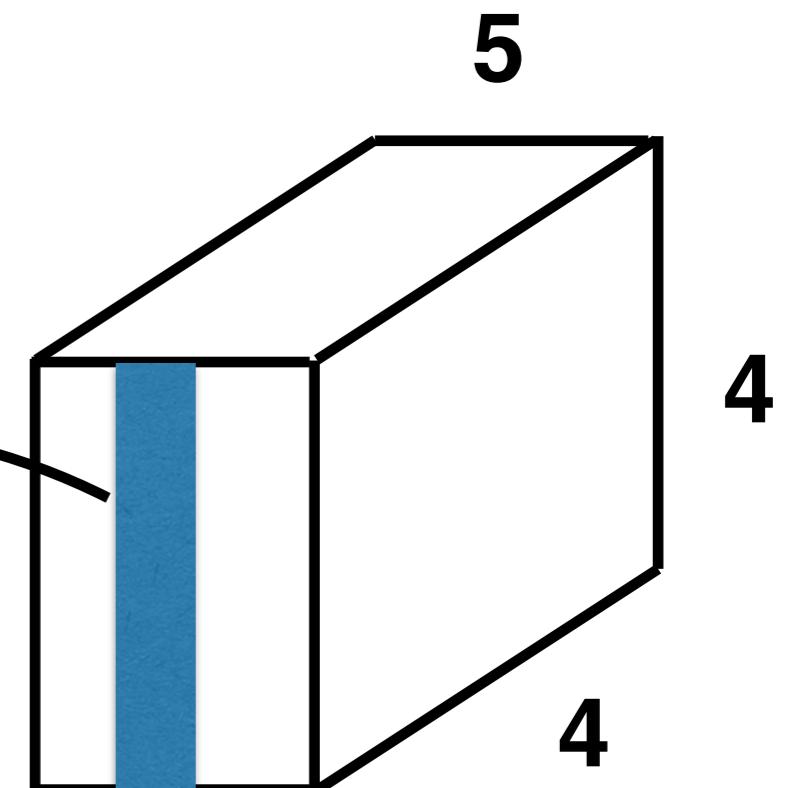




# Max Pooling

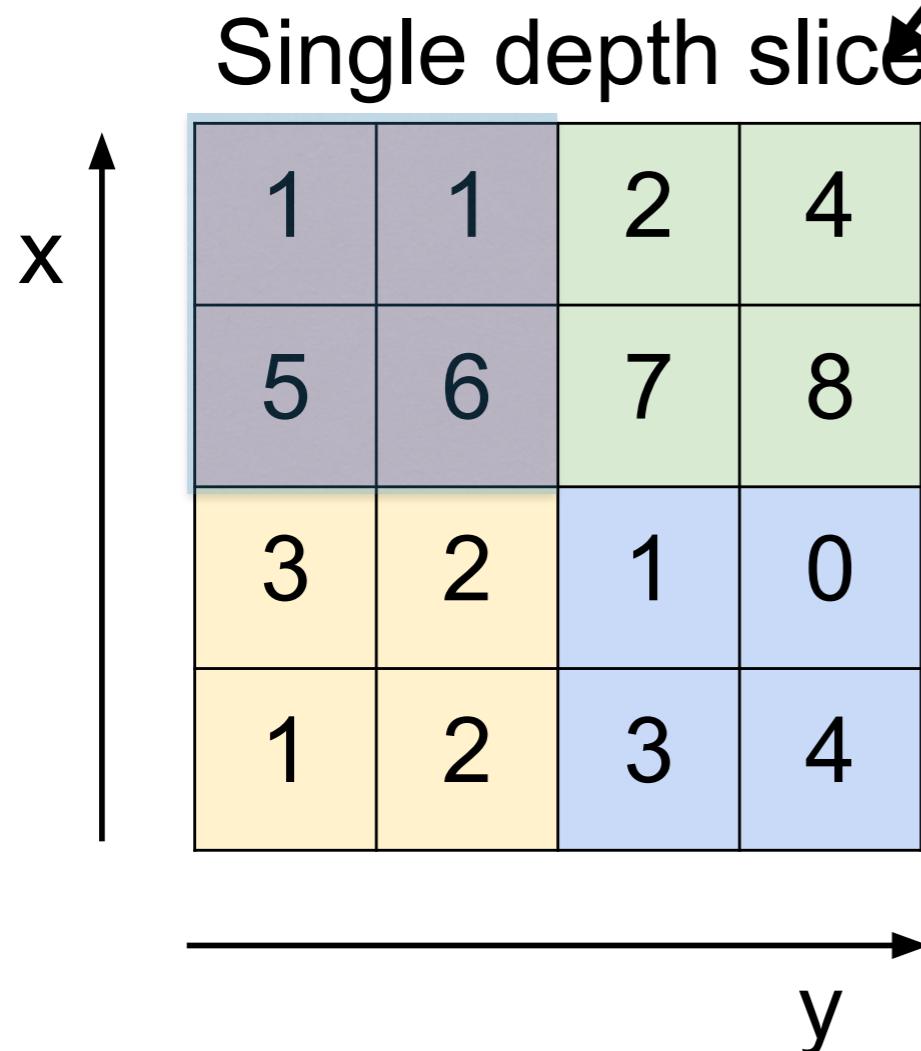


max pool with 2x2 filters  
and stride 2





# Max Pooling



max pool with 2x2 filters  
and stride 2

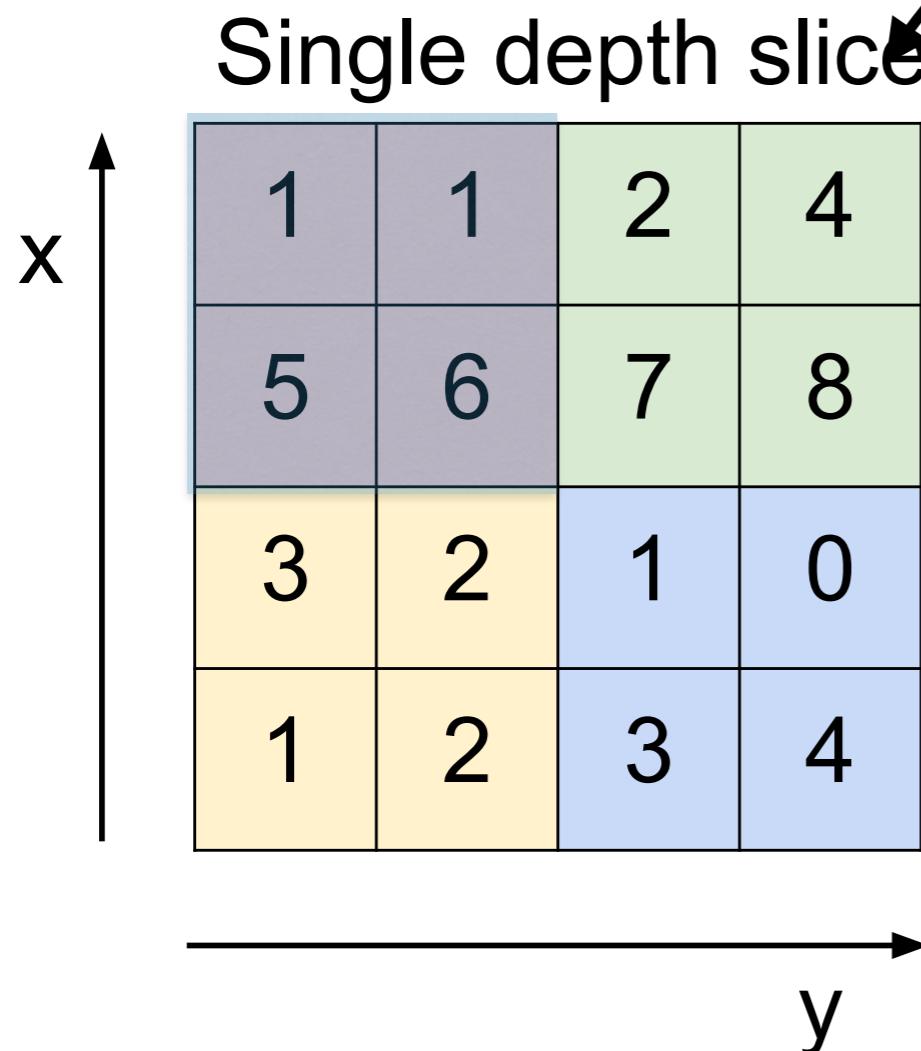
5

4

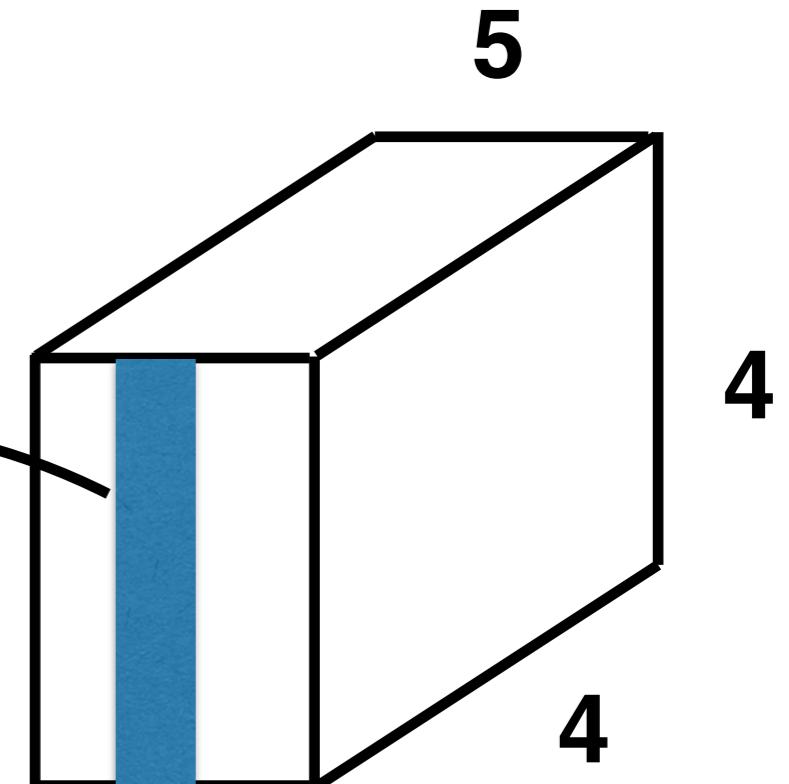
4



# Max Pooling

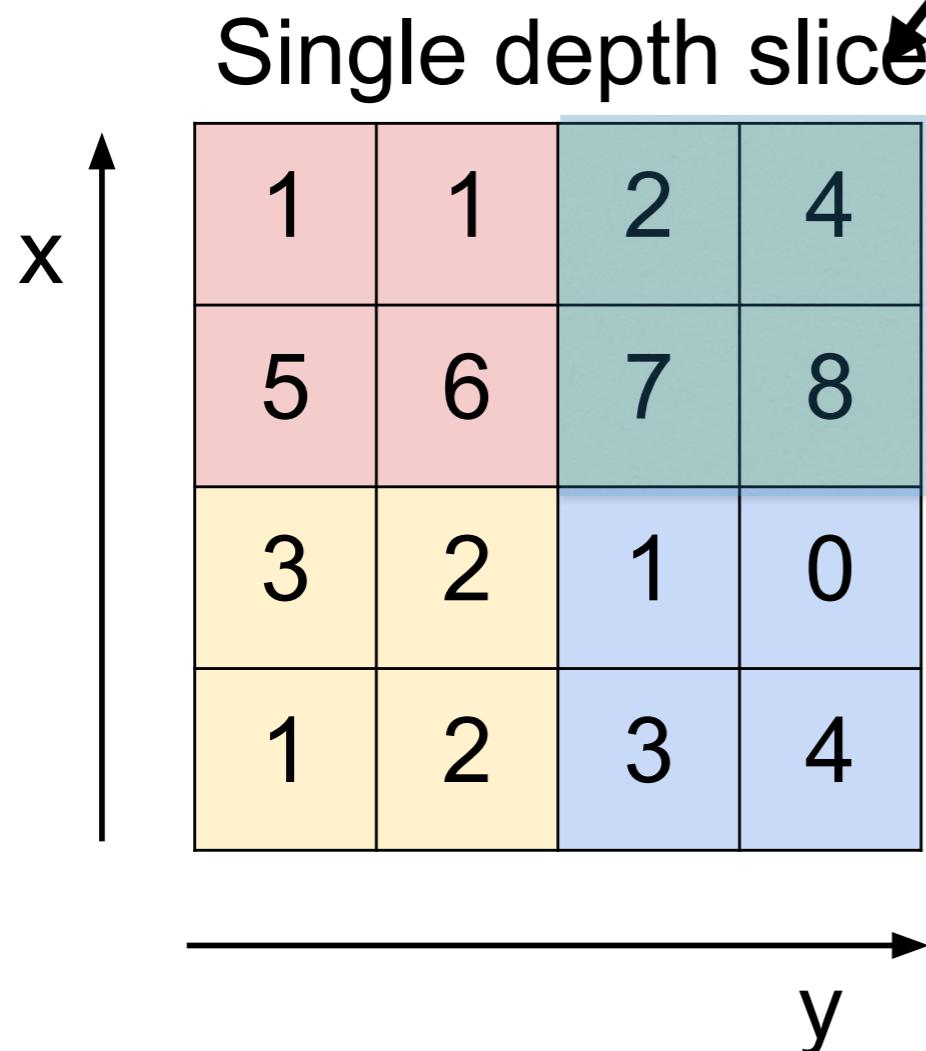


max pool with 2x2 filters  
and stride 2

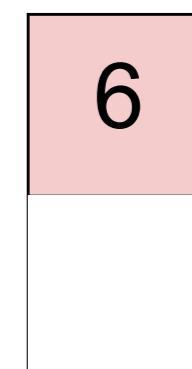




# Max Pooling

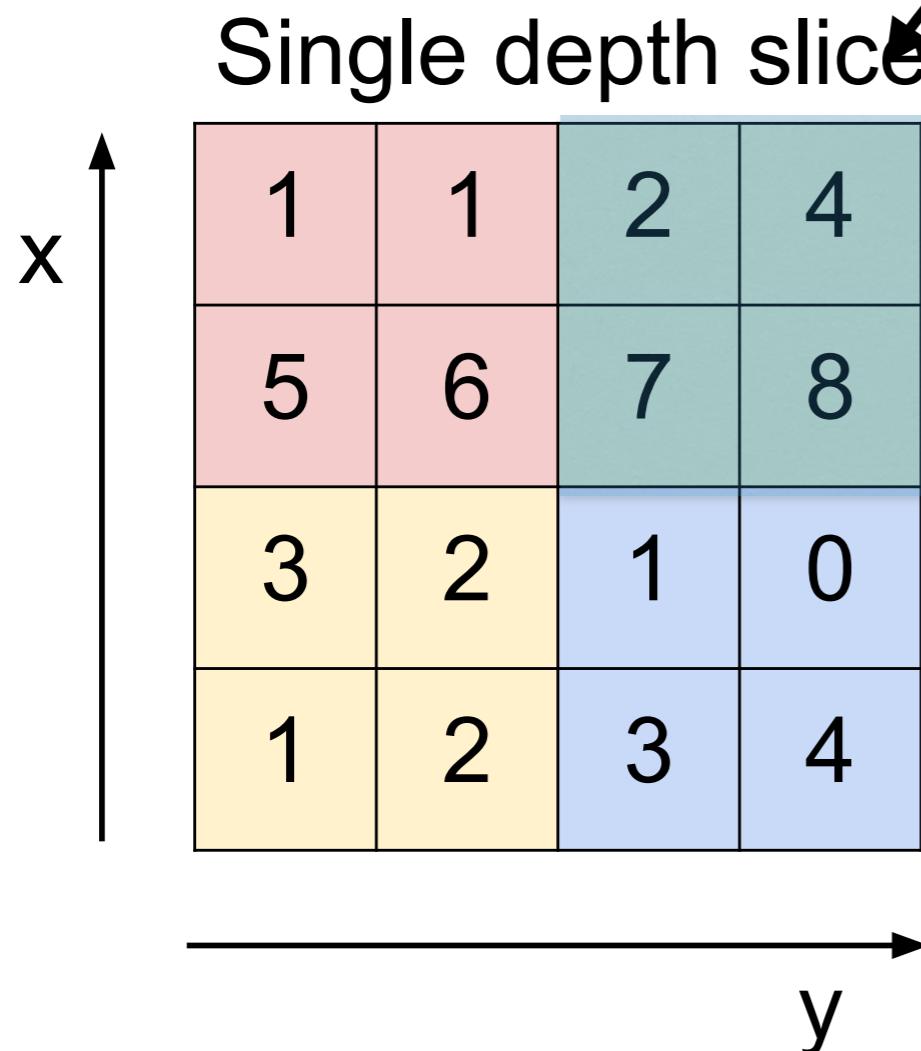


max pool with 2x2 filters  
and stride 2

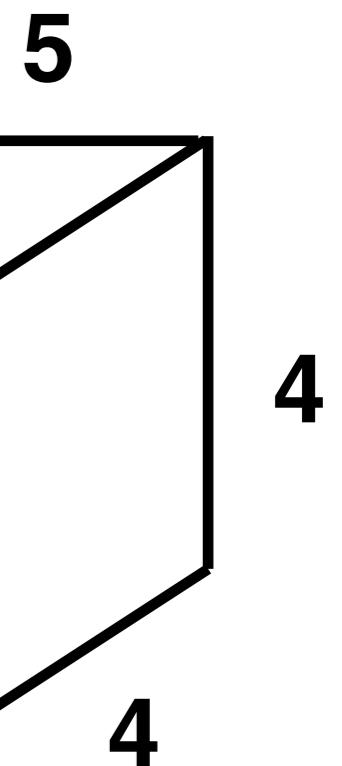
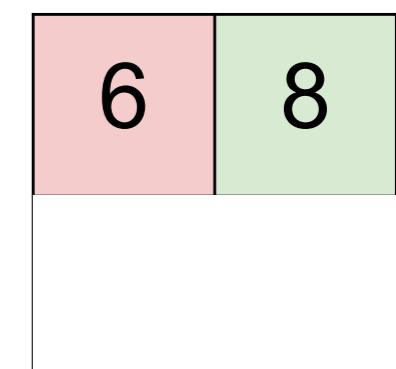




# Max Pooling

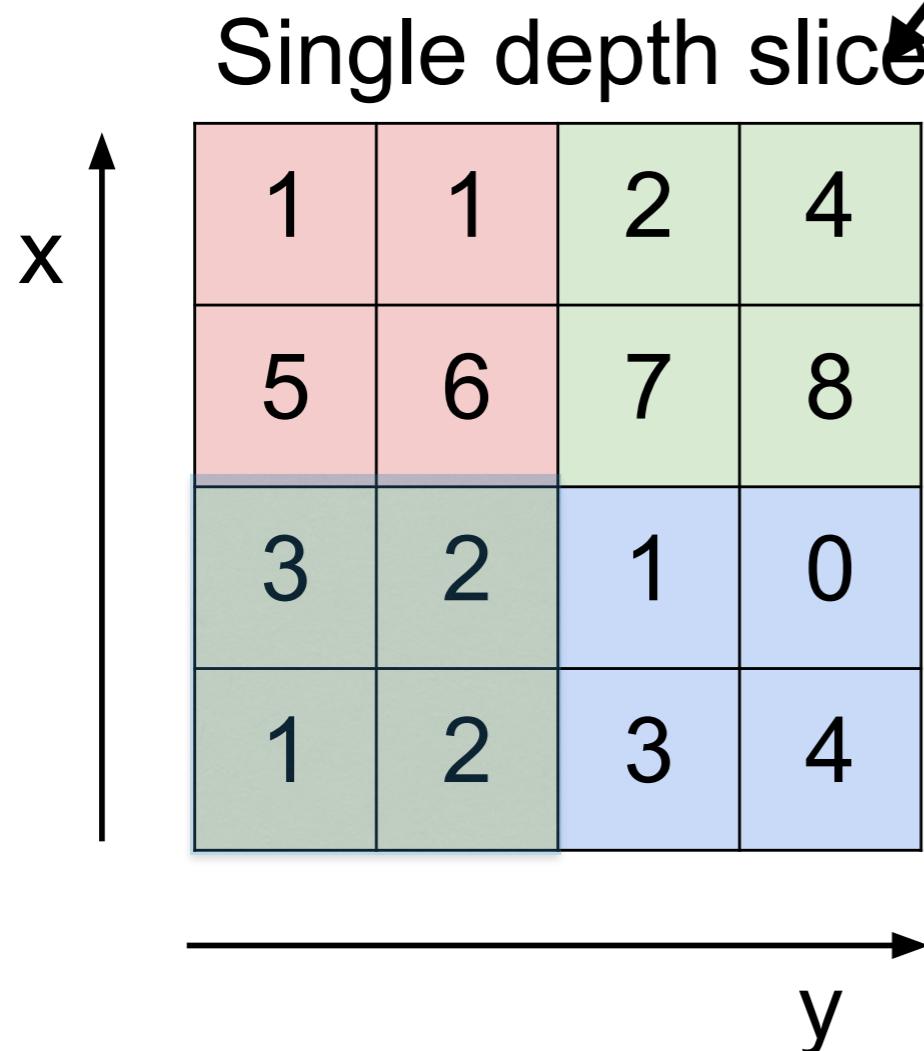


max pool with 2x2 filters  
and stride 2



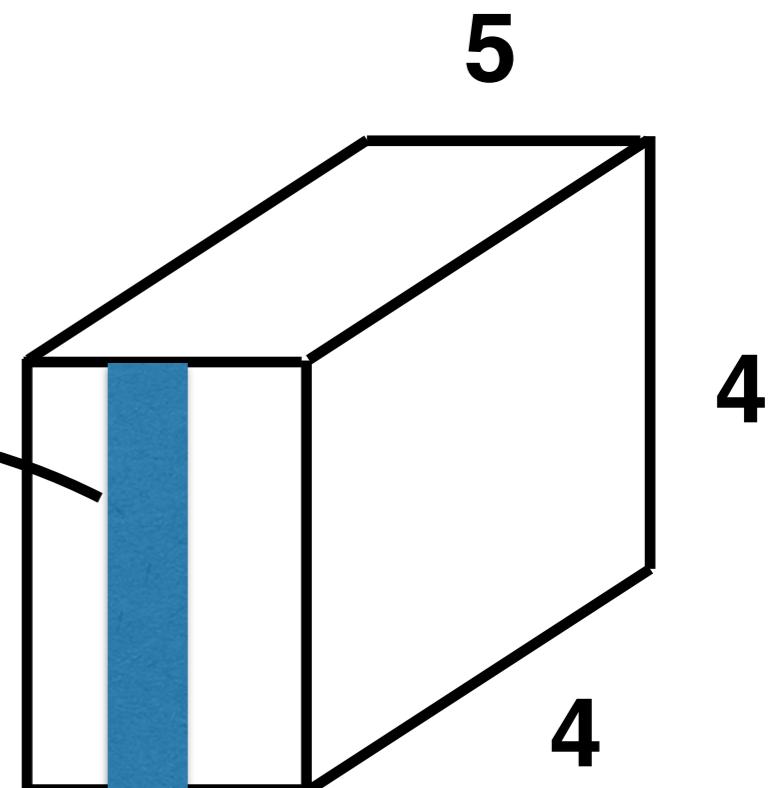


# Max Pooling



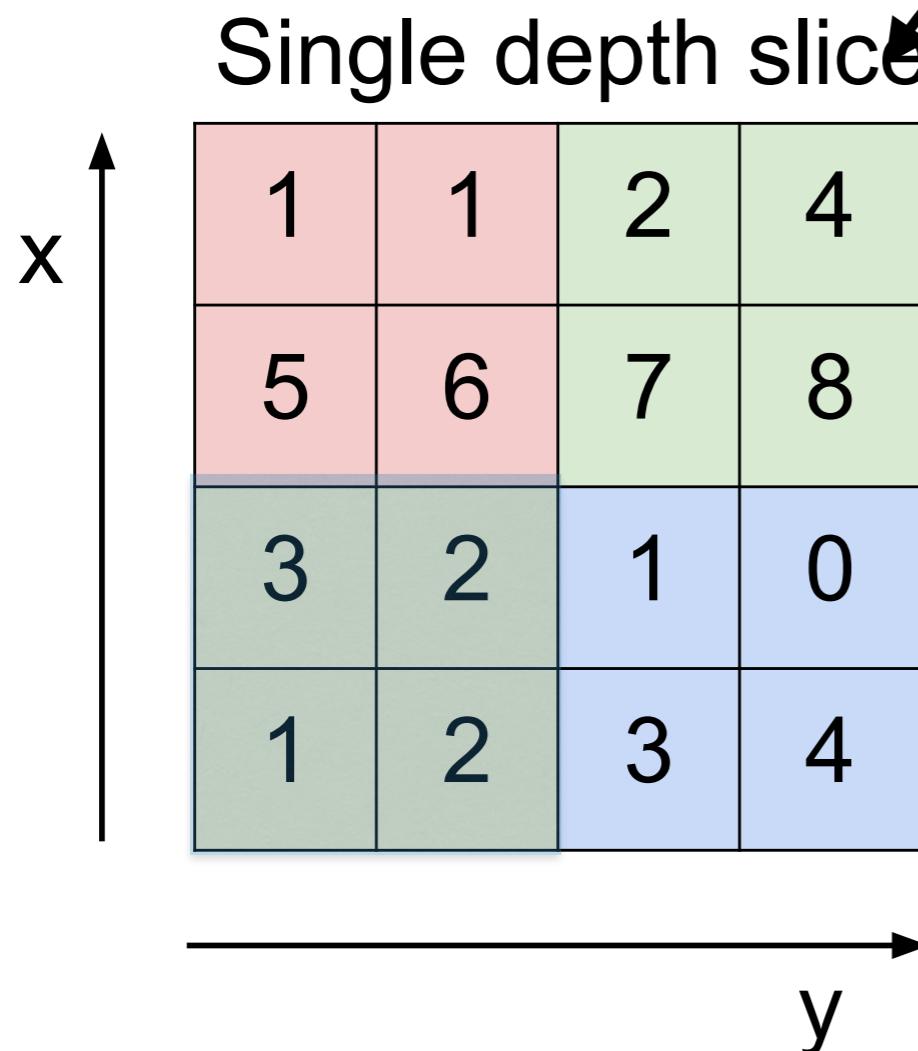
max pool with 2x2 filters  
and stride 2

6 8

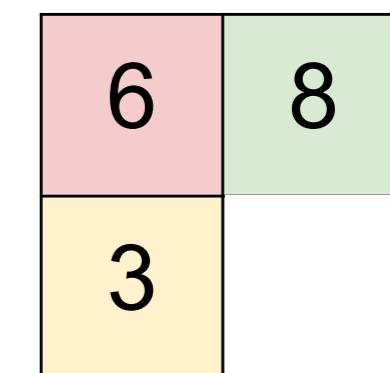




# Max Pooling

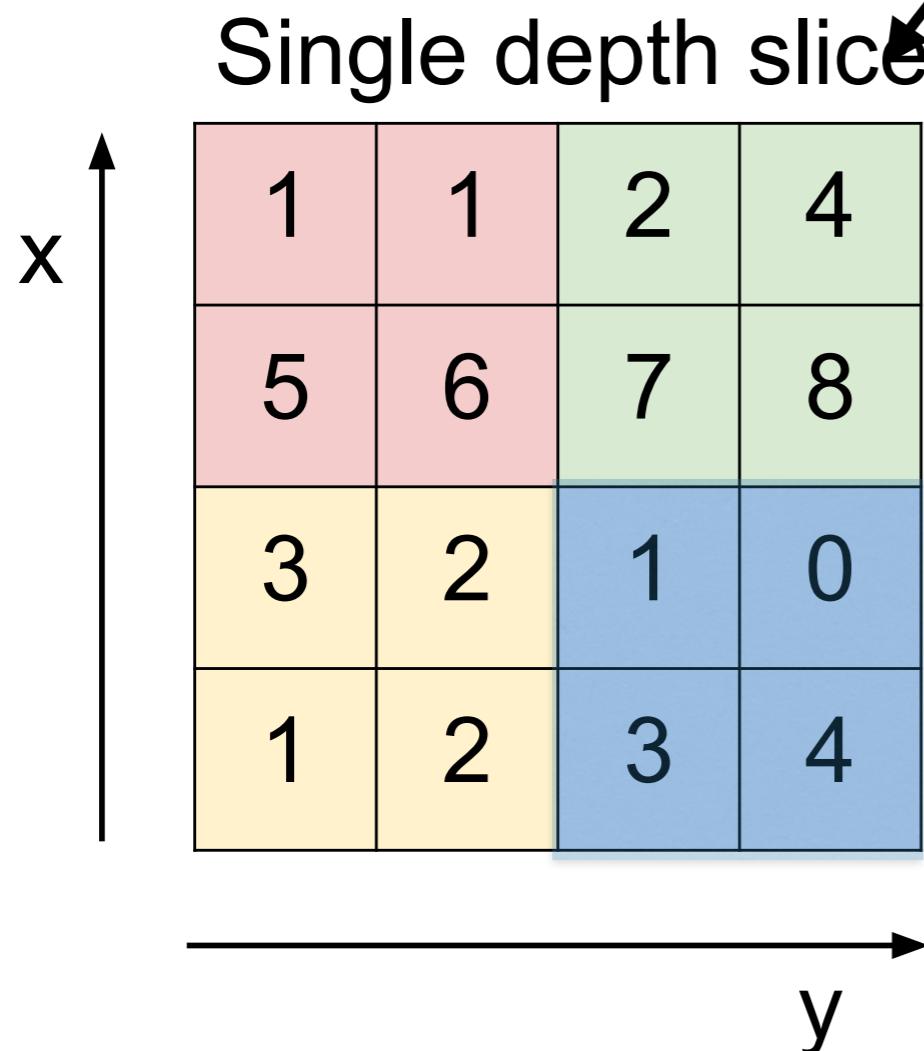


max pool with 2x2 filters  
and stride 2

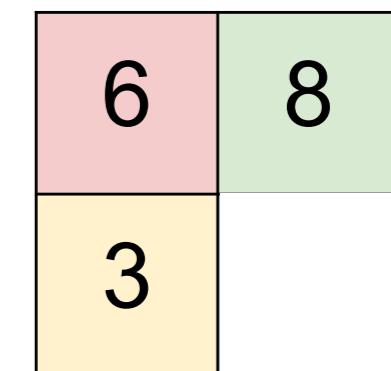




# Max Pooling

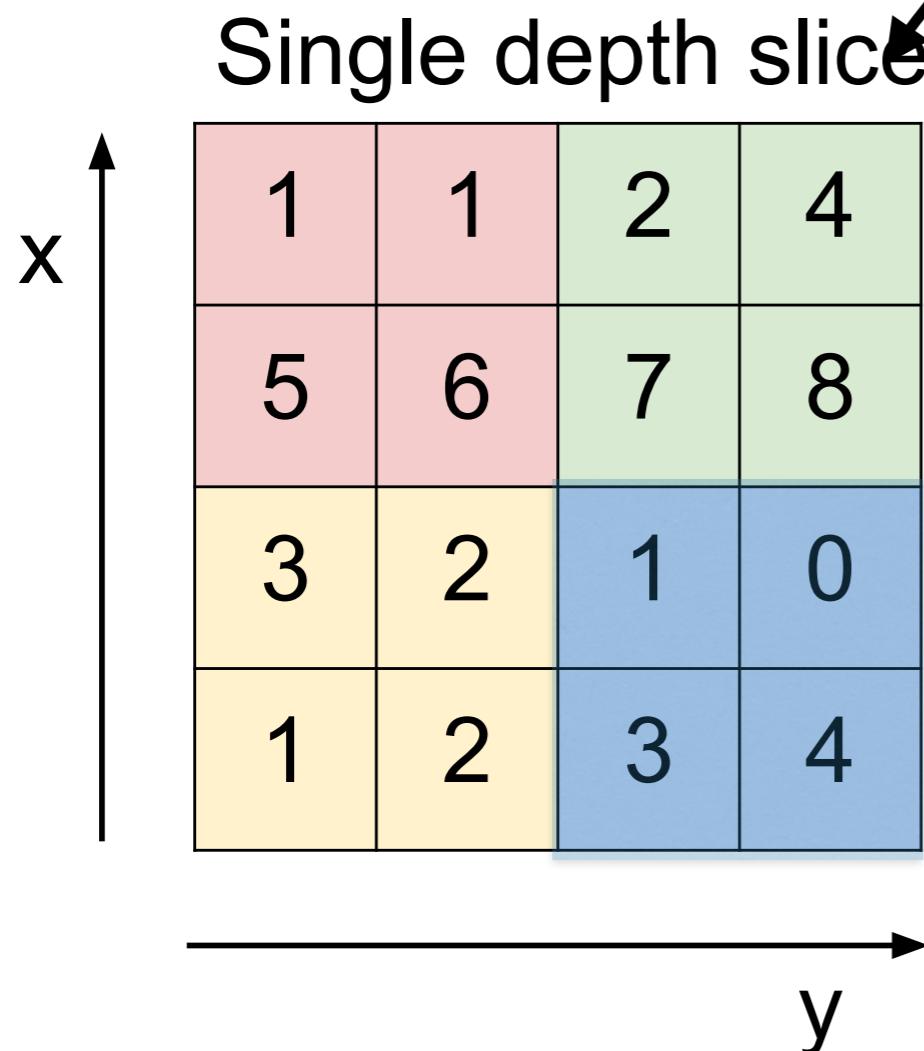


max pool with 2x2 filters  
and stride 2



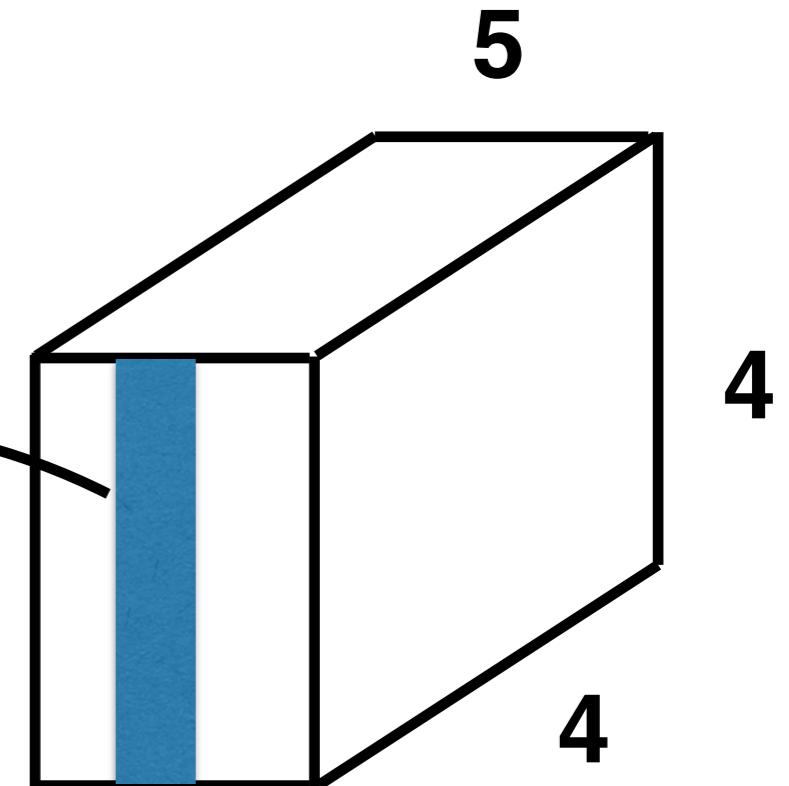


# Max Pooling



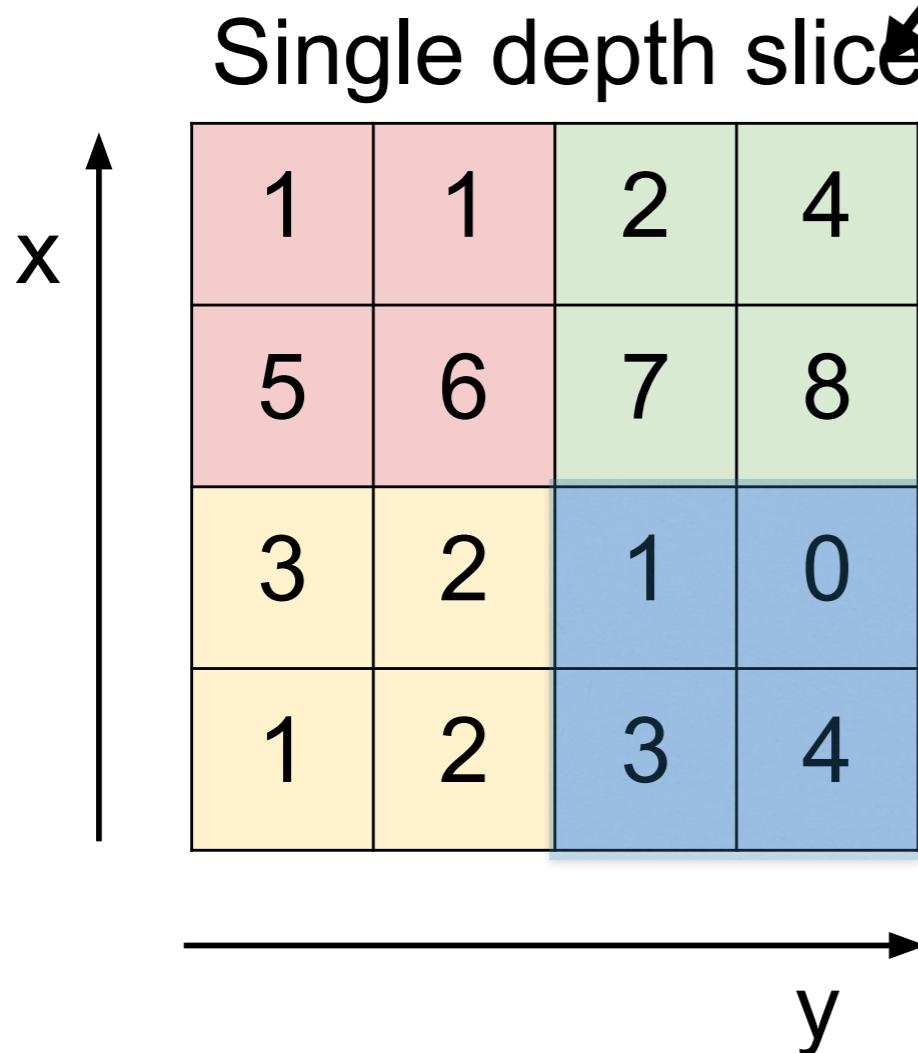
max pool with 2x2 filters  
and stride 2

6	8
3	4



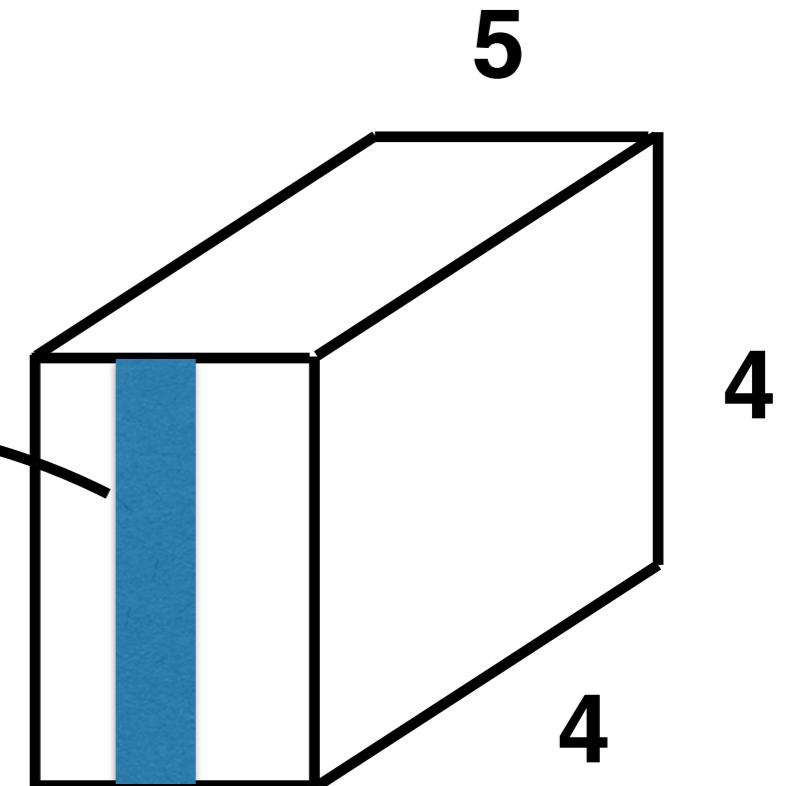


# Max Pooling



max pool with 2x2 filters  
and stride 2

Input: 4x4x5   Output: 2x2x5



6	8
3	4



# Max Pooling

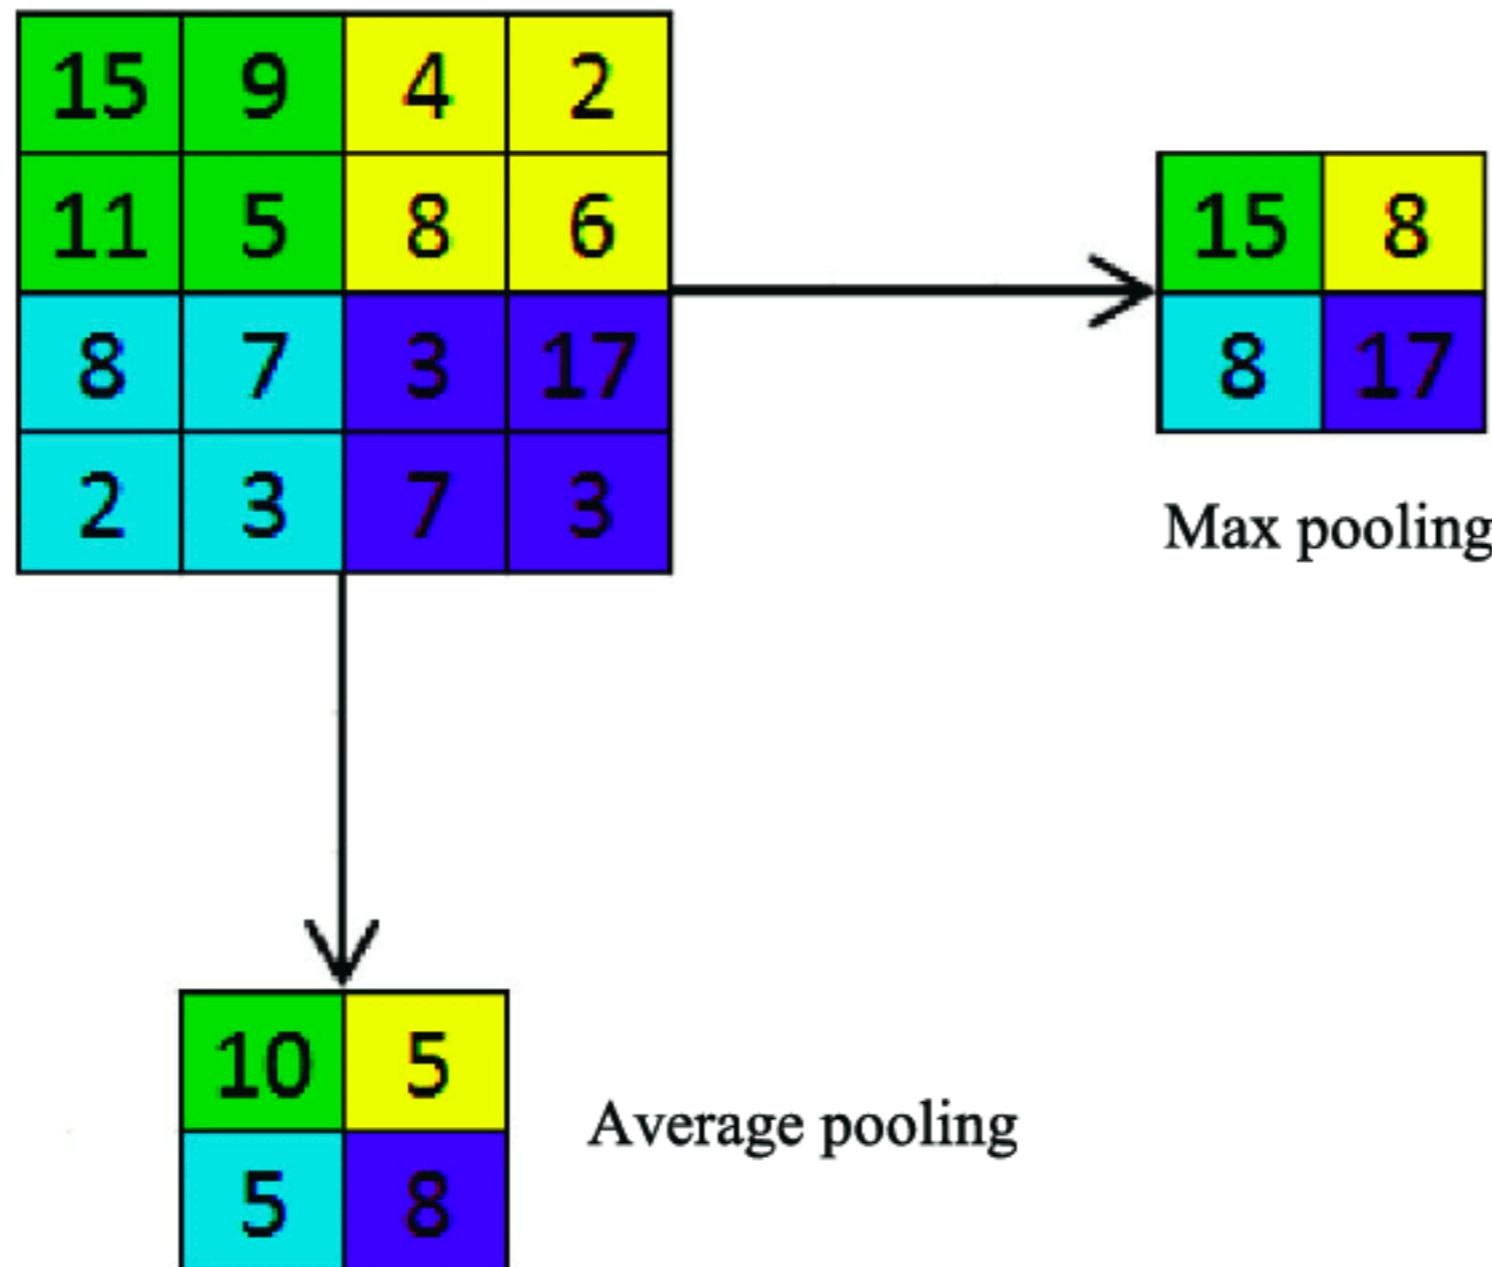
Idea: We trade some spacial resolution for efficiency.

For image recognition, it is often not necessary to know the exact position of an object.

In fact, pooling can improve results by increasing spatial tolerance.



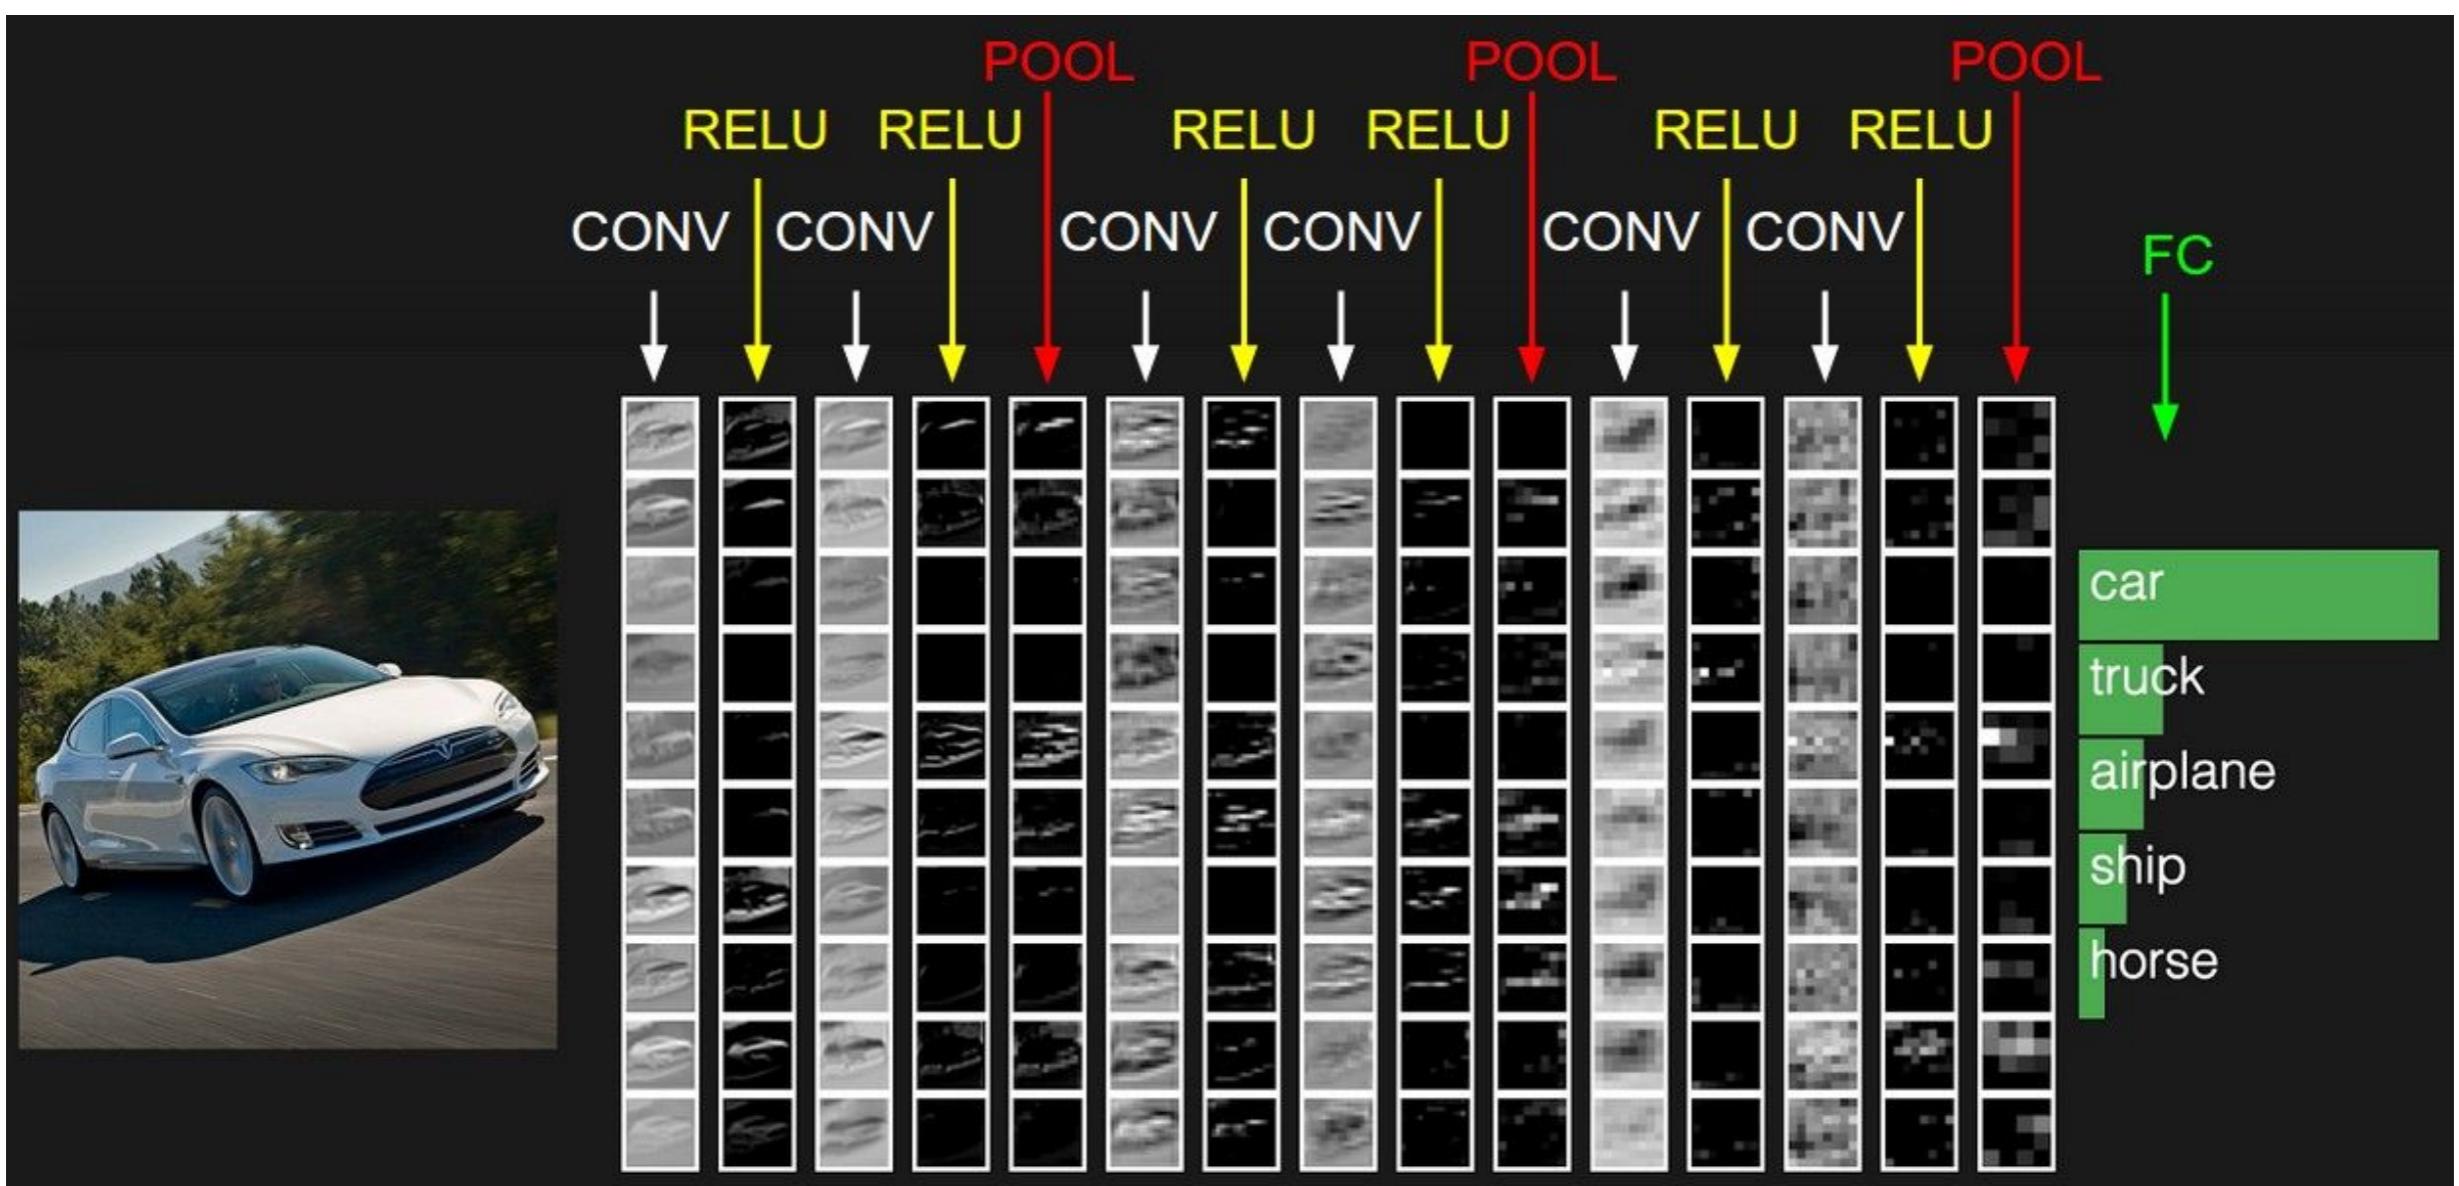
# Average Pooling



Average pooling seems to not work as well as max pooling.

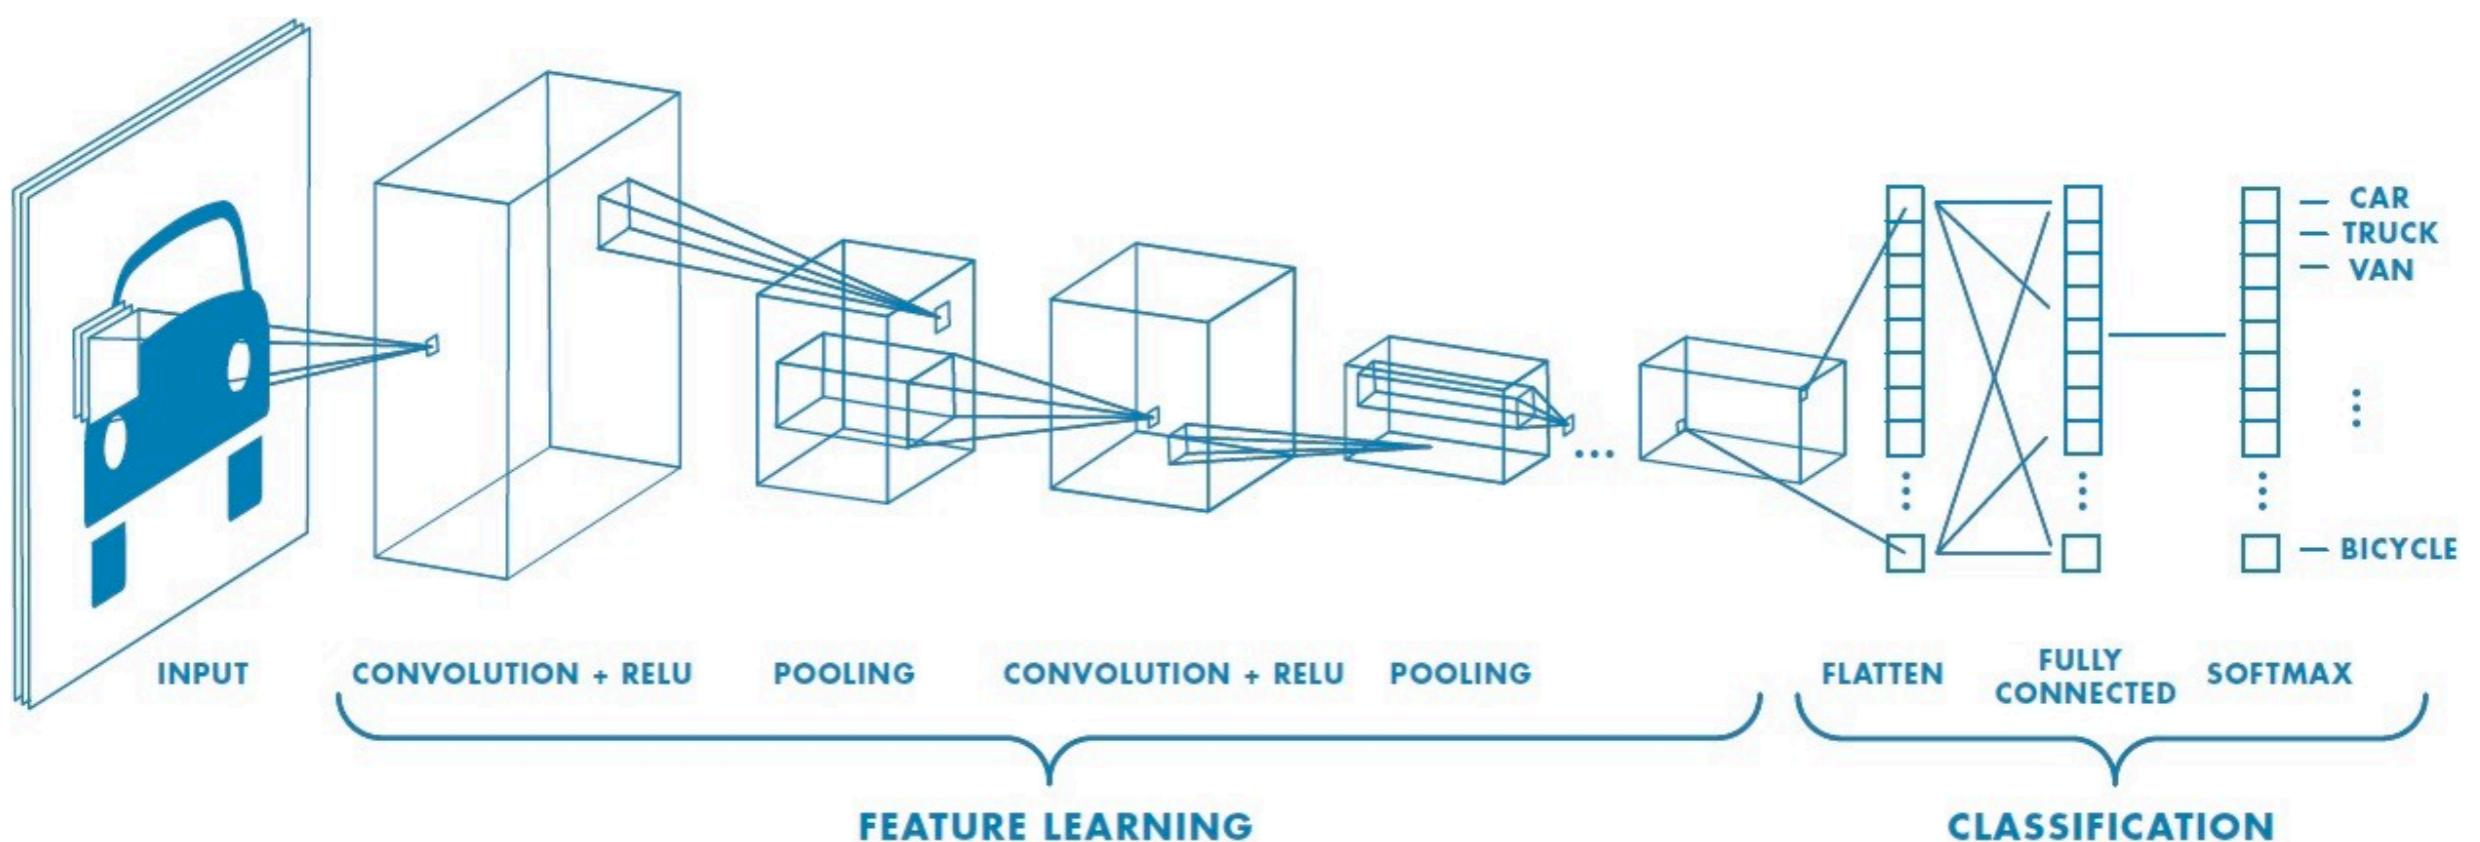


# Overall Architecture





# Final Classification





# Training a CNN

We train CNN's using backprop like any other NN.

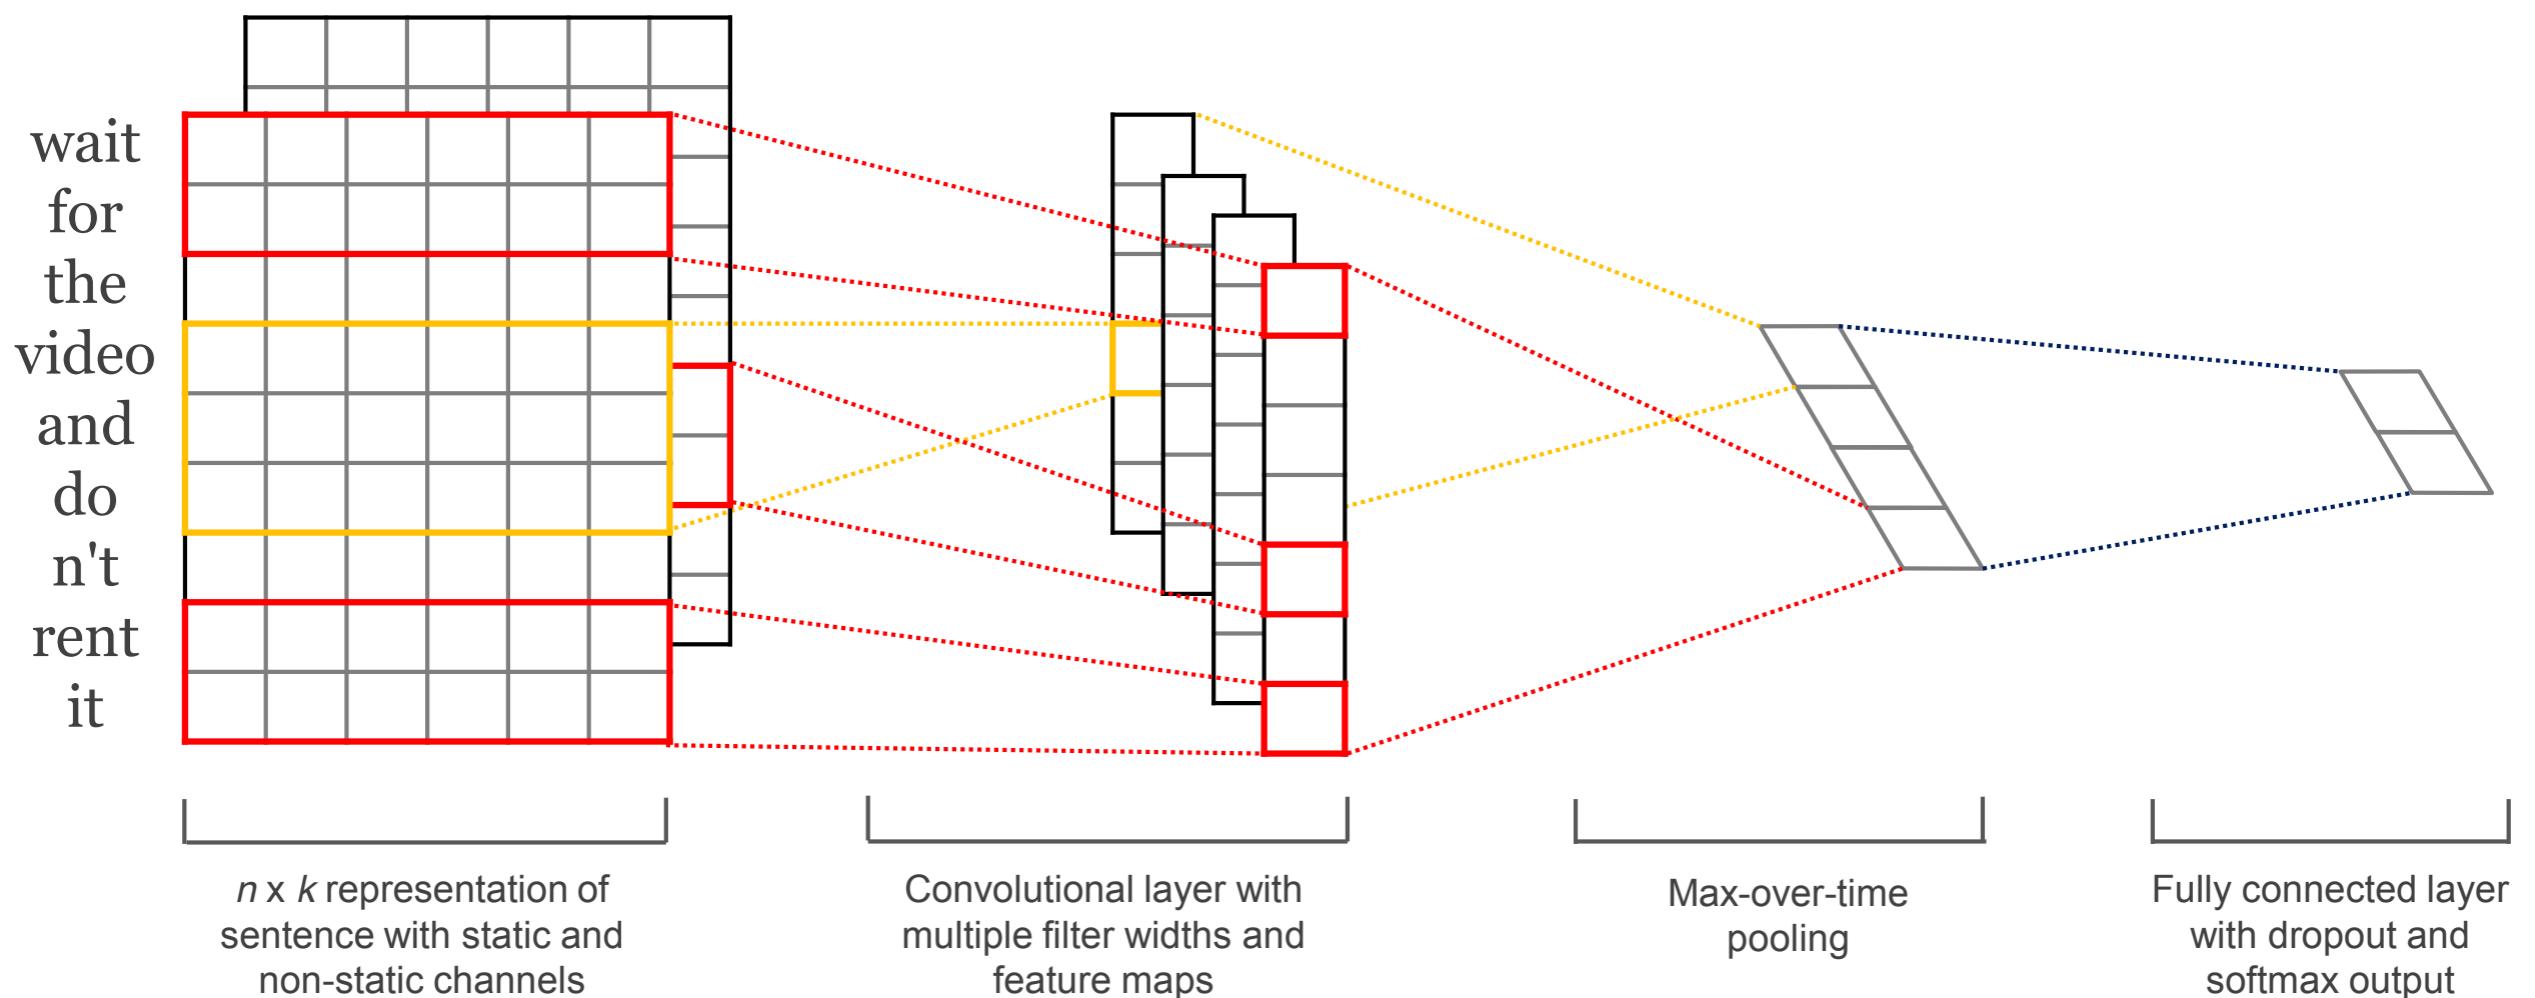
Convolutional filters are initialized randomly!

Because of the large number of dot products that needs to be computed, a GPU is typically used.

CNN's also have a fairly large memory footprint because the intermediate results of each convolution operation need to be stored for backprop.



# CNNs for NLP



**Yoon Kim: Convolutional Neural Networks for Sentence Classification EMNLP 2014**